

case_6.1

May 29, 2020

1 What factors are driving pay discrimination between men and women in your organization?

```
[3]: ### Load relevant packages
import pandas as pd
from scipy import stats
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.formula.api as sm
# import plotly.plotly as py
# https://community.plot.ly/t/solved-update-to-plotly-4-0-0-broke-application/
# 26526/2
import os

%matplotlib inline
plt.style.use('ggplot')
```

1.1 Case Introduction (10 mts)

Business Context. You are a data scientist in a large organization. Your company is undergoing an internal review of its hiring practices and employee compensation. In recent years, your firm has had low success in converting high-quality female candidates that it has wanted to hire. Management hypothesizes that this is due to possible pay discrimination and wants to figure out what is causing it.

Business Problem. As part of the internal review, the human resources department has approached you to specifically investigate the following question: “On balance, are men paid more than women in your organization? If so, what is driving this gap?”

Analytical Context. The human resources department has provided you with an employee database that contains information on various attributes such as performance, education, income, seniority, etc. We will use linear regression techniques on this dataset to solve the business problem described above. We will see how linear regression quantifies the correlation between the dependent variable (pay) and the independent variables (e.g. education, income, seniority, etc.)

The case is structured as follows: we will (1) perform exploratory data analysis to visually investigate the differences in pay; (2) use the observed insights to formally fit regression models; and finally (3) address the pay discrimination issue.

1.2 Data Exploration (30 mts)

```
[5]: Data = pd.read_csv('glassdoordata.csv')
```

```
[6]: Data.head()
```

```
[6]:
```

	jobtitle	gender	age	performance	education	department	\
0	Graphic Designer	Female	18	5	College	Operations	
1	Software Engineer	Male	21	5	College	Management	
2	Warehouse Associate	Female	19	4	PhD	Administration	
3	Software Engineer	Male	20	5	Masters	Sales	
4	Graphic Designer	Male	26	5	Masters	Engineering	

	seniority	income	bonus	
0	2	42363	9938	
1	5	108476	11128	
2	5	90208	9268	
3	4	108080	10154	
4	5	99464	9319	

The available features are: * **Job Title**: the title of the job (e.g. “Graphic Designer”, “Software Engineer”, etc); * **Gender**: male or female; * **Age**: age; * **Performance**: on a scale of 1 to 5, 1 being the lowest and 5 being the highest; * **Education**: different levels of education (e.g. “College”, “PhD”, “Masters”, “Highschool”); * **Department**: different departments in the organization (e.g. “Operations”, “Management”, etc); * **Seniority**: on a scale of 1 to 5, 1 being the lowest and 5 being the highest; * **Income, bonus**: both expressed in dollars

Since we are interested in total compensation, let's create a new column called pay:

```
[7]: Data['pay'] = Data['income'] + Data['bonus']
```

```
[8]: Data.head()
```

```
[8]:
```

	jobtitle	gender	age	performance	education	department	\
0	Graphic Designer	Female	18	5	College	Operations	
1	Software Engineer	Male	21	5	College	Management	
2	Warehouse Associate	Female	19	4	PhD	Administration	
3	Software Engineer	Male	20	5	Masters	Sales	
4	Graphic Designer	Male	26	5	Masters	Engineering	

	seniority	income	bonus	pay
0	2	42363	9938	52301
1	5	108476	11128	119604

```

2      5   90208   9268   99476
3      4   108080  10154  118234
4      5   99464   9319  108783

```

1.2.1 Exercise 1:

1.1 Construct a boxplot comparing pay between men and women. What can you conclude?

Answer. The median for men seems to be higher than the median for women.

```
[9]: sns.boxplot(x='gender', y = 'pay', data = Data)
#Data.boxplot(grid= False, column = ['pay'], by = ['gender'])
plt.title("Pay vs Gender");
```



1.2 Conduct a t - test on the difference in average pay between men and women. What can you conclude?

Answer. There is a statistically significant difference between the salary for men and women. The p - value is much smaller than 0.05.

```
[10]: ## A simple t test for difference of means
t2, p2 = stats.ttest_ind(Data.loc[Data['gender'] == 'Male', 'pay'], Data.
                         loc[Data['gender'] == 'Female', 'pay'])
```

```
print("t = " + str(t2))
print("p = " + str(p2))
```

```
t = 5.407461816876623
p = 8.000016978237565e-08
```

1.3 Construct a boxplot of pay, with respect to the following attributes: seniority, education, jobtitle and performance. Also do a scatter plot of pay vs. age. What patterns do you observe?

```
[11]: #Let us compute the average pay for men vs. women.
```

```
pay_by_gender = Data.groupby('gender')['pay']
pay_by_gender.mean()
```

```
[11]: gender
```

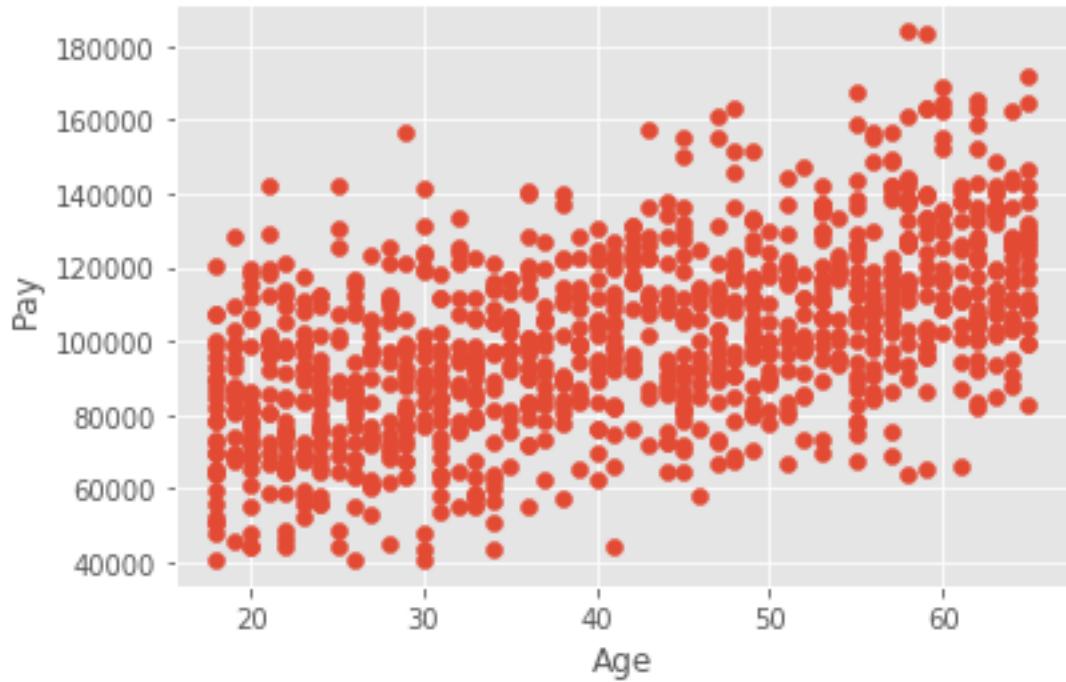
```
Female      96416.831197
Male        104918.678571
Name: pay, dtype: float64
```

```
[12]: pay_by_gender.mean()['Male'] - pay_by_gender.mean()['Female']
```

```
[12]: 8501.847374847363
```

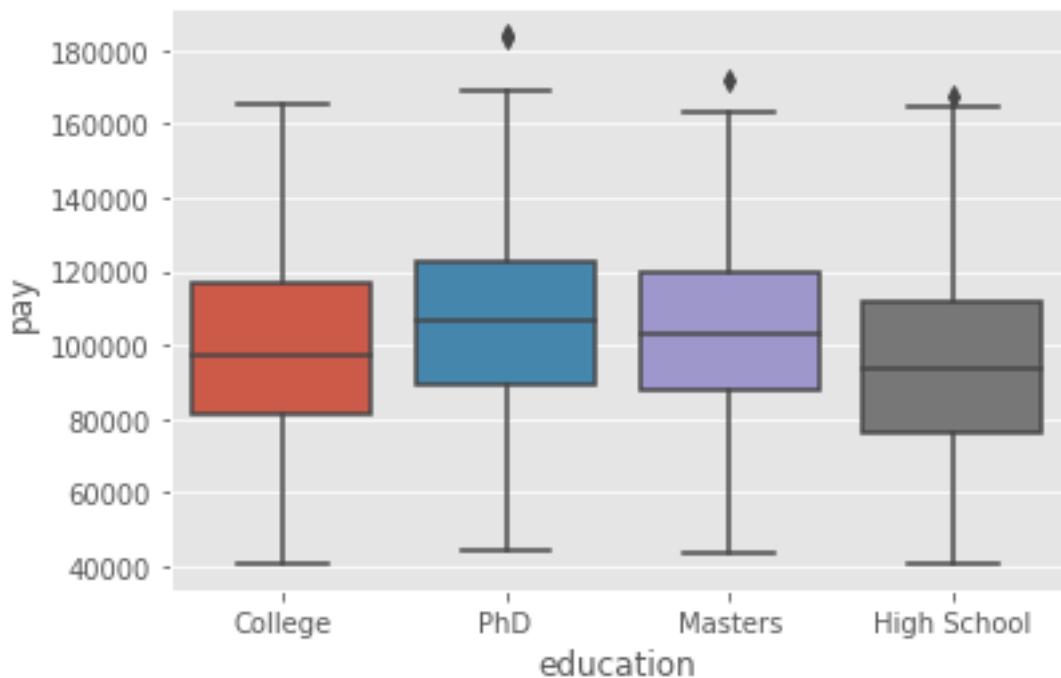
```
[13]: plt.scatter(Data['age'], Data['pay'])
plt.title("Pay vs. Age", fontsize=20, verticalalignment='bottom');
plt.xlabel("Age");
plt.ylabel("Pay");
```

Pay vs. Age



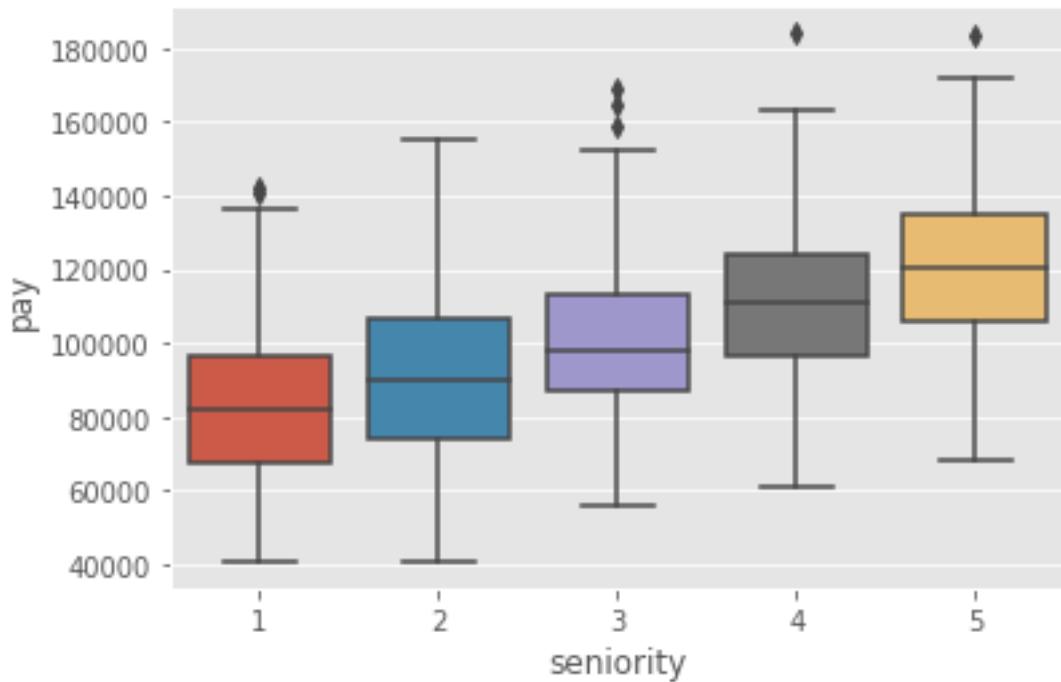
```
[14]: sns.boxplot(x='education', y = 'pay', data = Data)
plt.title("Pay vs. Education", fontsize=20, verticalalignment='bottom');
```

Pay vs. Education



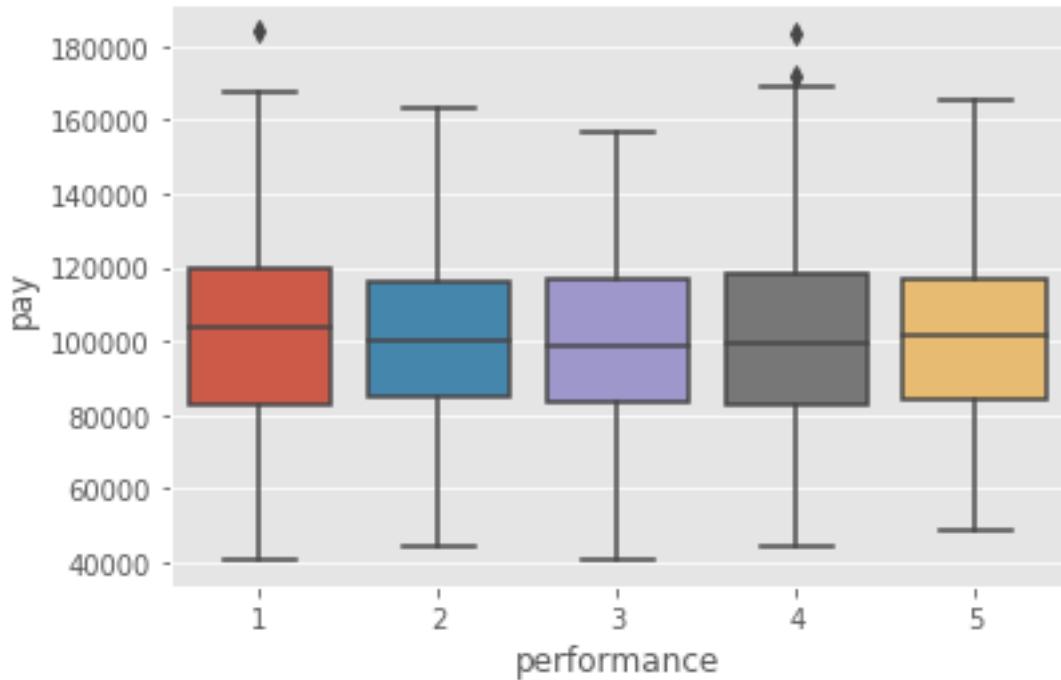
```
[15]: sns.boxplot(x='seniority', y = 'pay', data = Data)
plt.title("Pay vs. Seniority", fontsize=20, verticalalignment='bottom');
```

Pay vs. Seniority

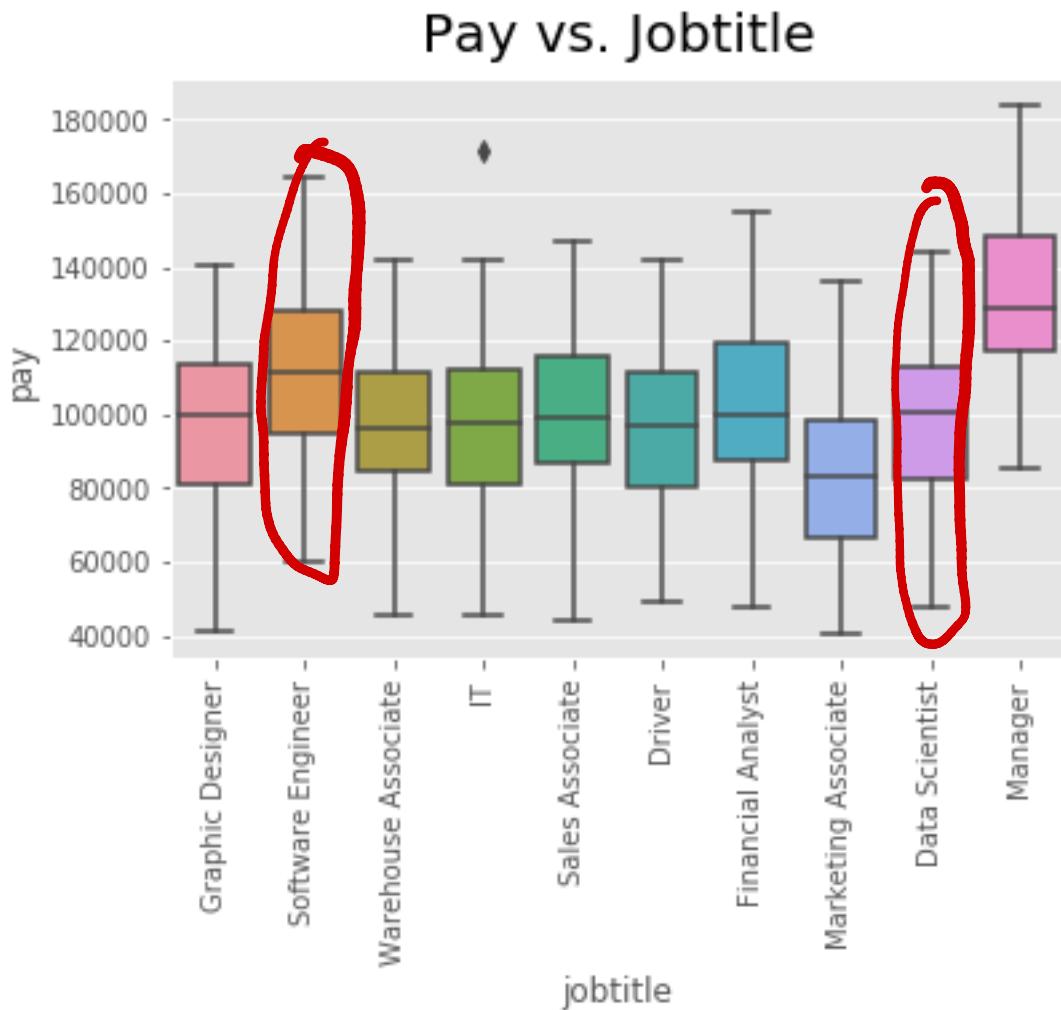


```
[16]: sns.boxplot(x='performance', y = 'pay', data = Data)
plt.title("Pay vs. Performance", fontsize=20, verticalalignment='bottom');
```

Pay vs. Performance



```
[17]: sns.boxplot(x='jobtitle', y = 'pay', data = Data)
plt.title("Pay vs. Jobtitle", fontsize=20, verticalalignment='bottom')
plt.xticks(rotation=90);
```



****Answer:**** There is a significant variation of pay across some of the attributes. Pay increases with seniority. Among job titles, managers are paid relatively high on average. Marketing associates are paid much less. Pay vs. education is similar for all categories except for people with only a high school education. They are paid on average much less. There is a linear relationship between pay and age.

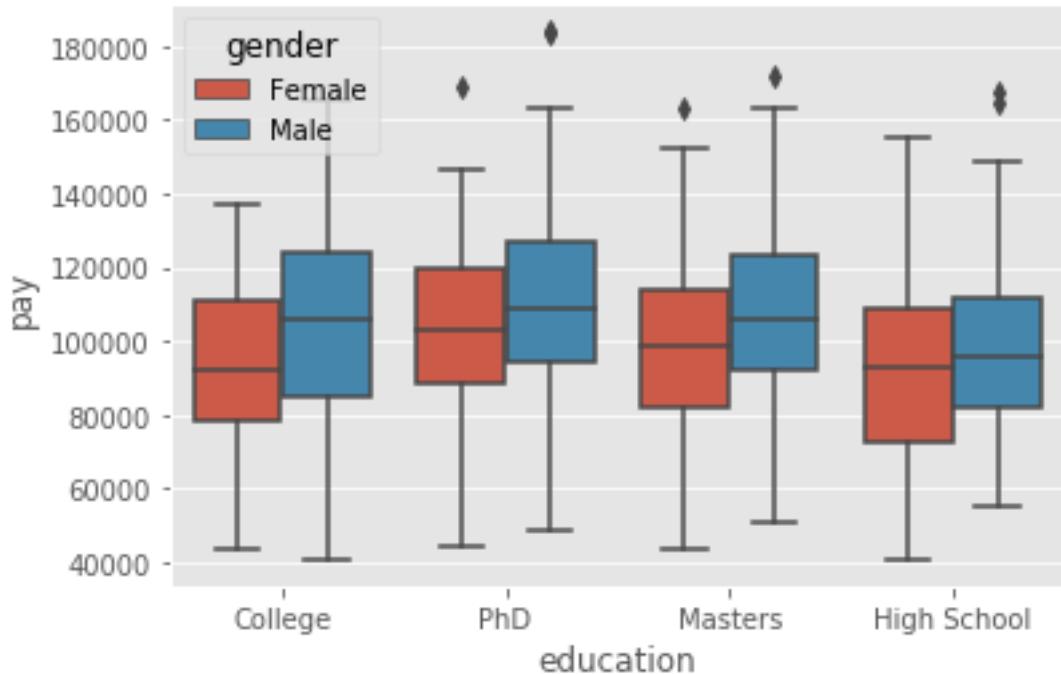
1.2.2 Exercise 2:

Now, let's make the same plots as in Exercise 1.3, but additionally differentiate by gender. What patterns do you observe?

Answer.

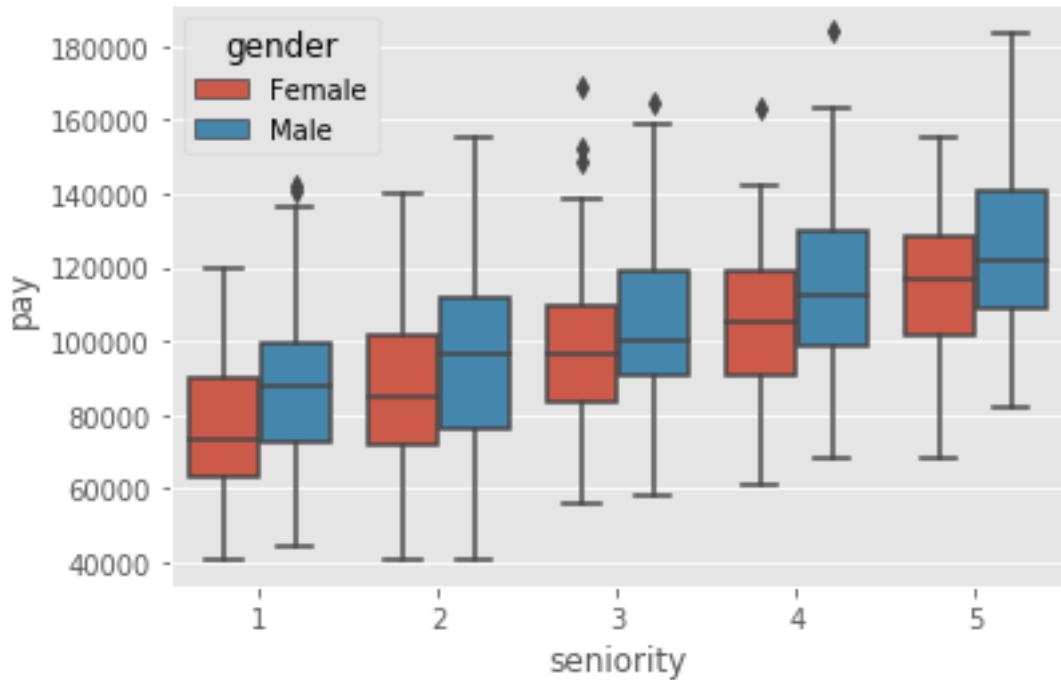
```
[18]: sns.boxplot(x='education', y = 'pay', hue = 'gender', data = Data)
plt.title("Pay vs. Education", fontsize=20, verticalalignment='bottom');
```

Pay vs. Education



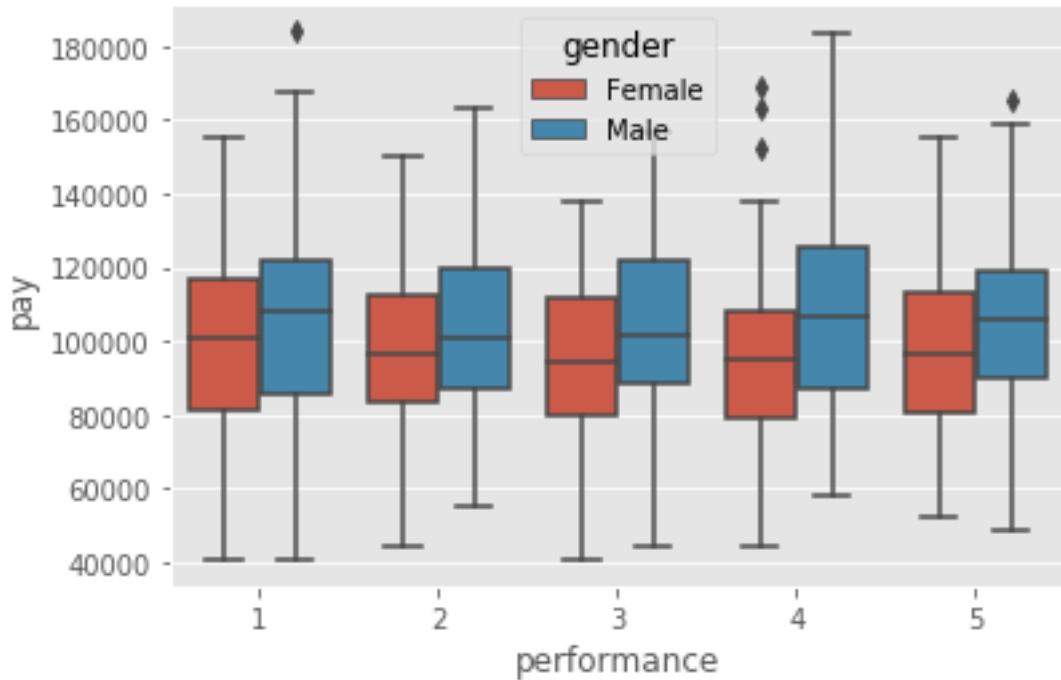
```
[19]: sns.boxplot(x='seniority', y = 'pay', hue = 'gender', data = Data)
plt.title("Pay vs. Seniority", fontsize=20, verticalalignment='bottom');
```

Pay vs. Seniority



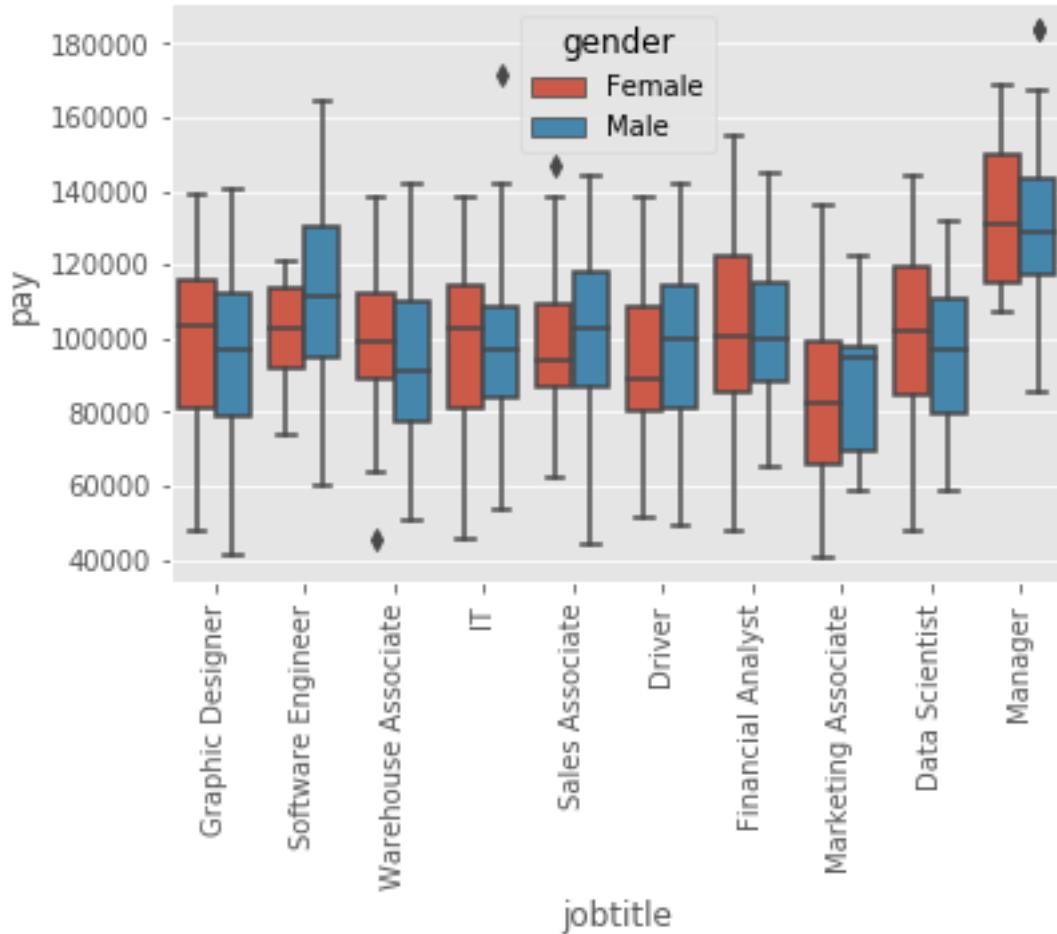
```
[20]: sns.boxplot(x='performance', y = 'pay', hue = 'gender', data = Data)
plt.title("Pay vs. Performance", fontsize=20, verticalalignment='bottom');
```

Pay vs. Performance



```
[21]: sns.boxplot(x='jobtitle', y = 'pay', hue = 'gender', data = Data)
plt.title("Pay vs. Jobtitle", fontsize=20, verticalalignment='bottom')
plt.xticks(rotation=90);
```

Pay vs. Jobtitle



There is quite some variation of pay across various categories (such as seniority, performance, jobtitle etc.). Within each category, there is some difference between pay for men and women, but not as much as across categories.

1.2.3 Exercise 3: (10 mts)

Consider the following statements:

- I. Based on the boxplot and the table, the average salary for men is 8500 USD more than that of women in this organization.
- II. Men are paid significantly more than women due solely to gender differences.

Which of the following conclusions about these statements is correct?

- a. Both Statements I and II are correct.
- b. Statement I is correct, but Statement II is not.

- c. Statement I is not correct, but Statement II is.
- d. Both Statements I and II are incorrect.

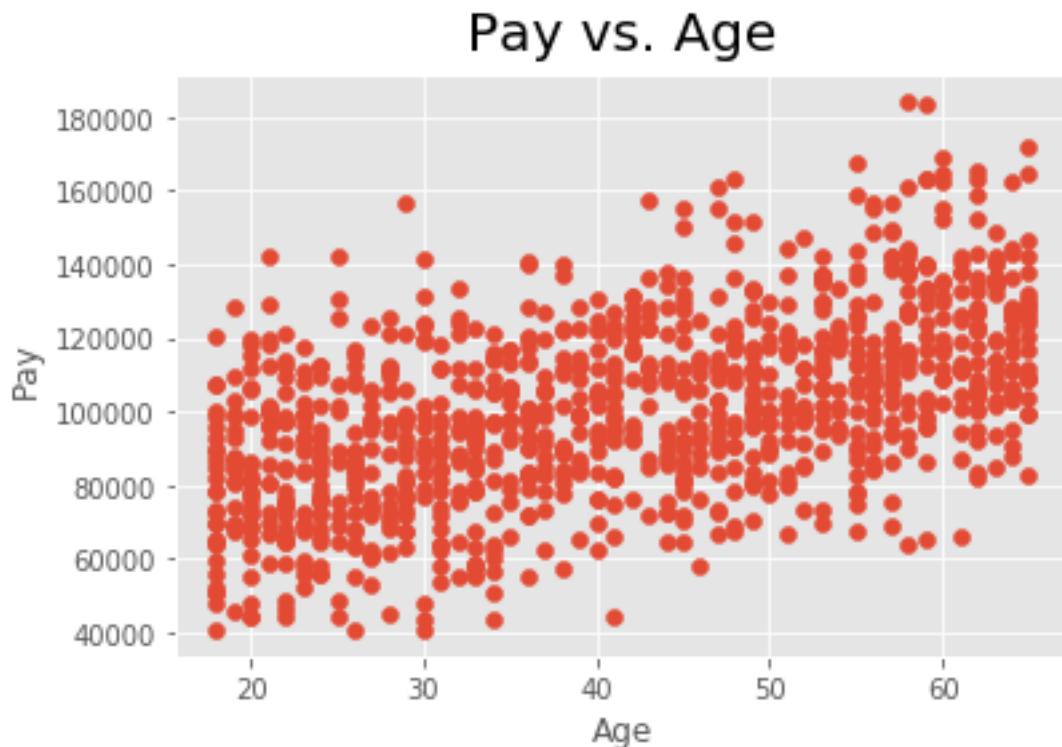
Answer. Statement I is true just as the tables and plot show. However, we cannot conclude that Statement II is correct. This is because there might be other factors that contribute to the higher pay for men. For instance, the men in this data set might all be in senior positions, while most of the women are in junior positions. Pay usually increases with seniority. Thus if we do not account for seniority while calculating wage gap, we could be comparing apples to oranges.

1.3 What are the variables that influence pay?

As we discussed in Exercise 3 and also saw in the data exploration, even though there is a significant pay gap between the genders, there are also some other factors at work which are driving this difference. Thus, ignoring these factors while addressing pay discrimination could lead to wrong or misleading conclusions about how to go about addressing the issue.

How do we take into account the influence of the other variables on pay? What are these variables? A good place to start is to take inspiration from our exploratory data analysis before. The only numerical variable in this dataset is age, so let's start with a scatter plot of pay vs. age:

```
[22]: plt.scatter(Data['age'], Data['pay'])
plt.title("Pay vs. Age", fontsize=20, verticalalignment='bottom')
plt.xlabel("Age")
plt.ylabel("Pay");
```



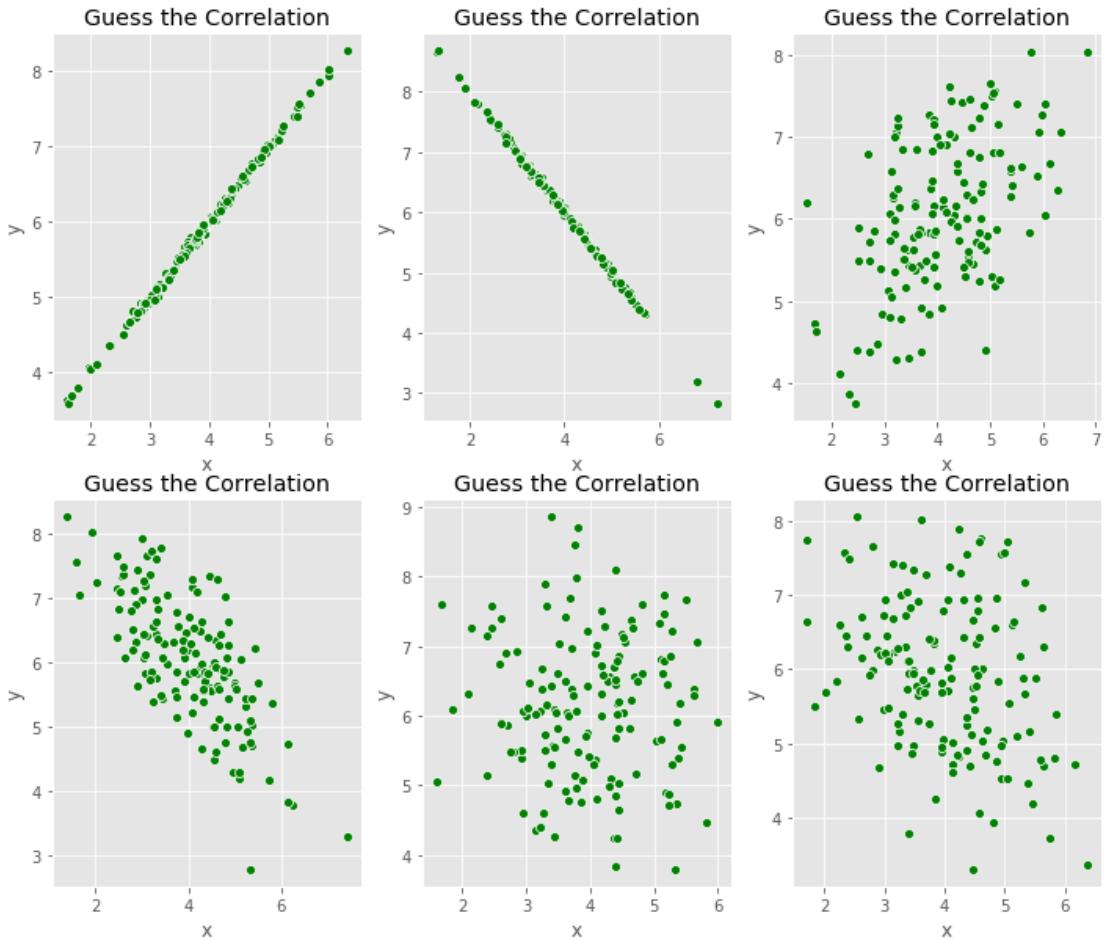
1.3.1 Correlation measures the linear relationship between two variables

We see a linear relationship between pay and age. Pay seems to be positively correlated with age; i.e. the older someone is, the more they tend to get paid. Thus, it could be that there are more men in our dataset that are older and the pay difference between men and women we see could be a consequence of this. The quantity correlation measures how linear the relationship between two variables is. Guess the correlation between pay and age.

1.3.2 Guess the correlation (5 mts)

To get warmed up on the concept of correlation, here are some simulated scenarios. In each of these cases, guess the correlation:

```
[23]: plt.figure(figsize=(12,10))
rho = [0.999, -0.999, 0.5, -0.7, 0.001, -0.3]
cor_list = []
np.random.seed(10)
for i, r in enumerate(rho):
    plt.subplot(2,3,i+1)
    mean, cov = [4, 6], [(1, r), (r, 1)]
    x, y = np.random.multivariate_normal(mean, cov, 150).T
    ax = sns.scatterplot(x=x, y=y, color="g")
    cor_list.append(np.corrcoef(x,y)[0,1])
    plt.xlabel("x")
    plt.ylabel("y")
    plt.title("Guess the Correlation")
```



```
[24]: cor_list
```

```
[24]: [0.9990081442281228,
-0.9992506121564401,
0.5047051300834758,
-0.722898221972322,
-0.040411257687132,
-0.3208318426338665]
```

1.3.3 Question:

If the correlation between two variables is zero, then the two variables are unrelated. Is this true or false?

This is false. There could be non-linear relationships between the variables, and these may not be captured by correlation, which is a linear measure. In such cases, advanced models that capture non-linear relationships are required.

1.3.4 Correlation matrix (10 mts)

To find the variables that have the greatest influence on pay, we can compute a correlation matrix which measures the pairwise correlations between any two variables:

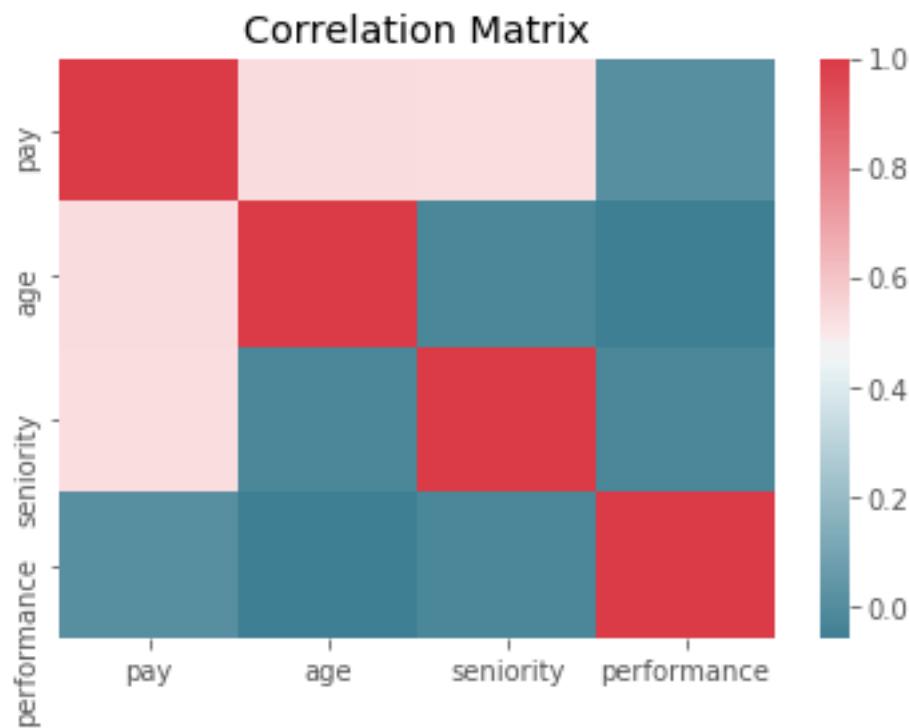
```
[25]: corr_mat = Data[['pay', 'age', 'seniority','performance']].corr()  
corr_mat
```

```
[25]:          pay      age  seniority  performance  
pay    1.000000  0.533715  0.530307  0.014155  
age    0.533715  1.000000 -0.021414 -0.056875  
seniority  0.530307 -0.021414  1.000000 -0.021127  
performance  0.014155 -0.056875 -0.021127  1.000000
```

We can also visualize the correlation matrix above as a heatmap:

```
[26]: sns.heatmap(corr_mat,cmap=sns.diverging_palette(220, 10, as_cmap=True))  
plt.title("Correlation Matrix")
```

```
[26]: Text(0.5,1,'Correlation Matrix')
```

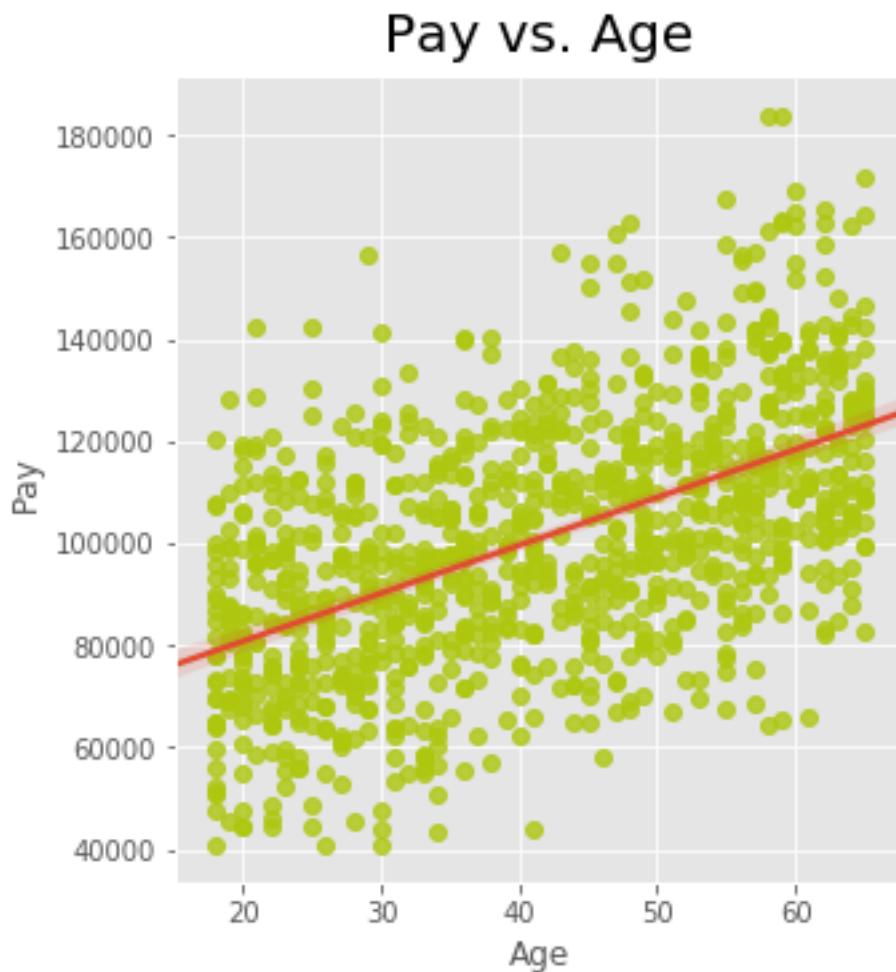


1.4 Using linear models to account for variables correlated with pay (10 mts)

Once we identify some independent variables that are correlated with the dependent variable, a linear model can be used to capture this relationship quantitatively. A linear model does this by finding a line that **best fits** the data points:

```
[27]: sns.lmplot(x = 'age', y = 'pay', data = Data, scatter_kws = {'color': (174/  
255,199/255,14/255)})  
plt.title("Pay vs. Age", fontsize=20, verticalalignment='bottom')  
plt.xlabel("Age")  
plt.ylabel("Pay")
```

```
[27]: Text(-15.075,0.5,'Pay')
```



A line has two parameters – intercept (β_0) and slope (β_1). Thus a linear model for pay vs. age can be represented as

$$\text{pay} = \beta_0 + \beta_1 \text{age} + \text{error.}$$

The interpretation of the coefficient β_1 is the following: an increase of one year in age will on average lead to a change of β_1 USD in pay. The intercept β_0 can be thought of as a sort of “baseline” pay.

1.4.1 Interactive Visualization (10 mts)

The following interactive visualization gives you a few scenarios. Try to find the line of best fit in each scenario:

```
[28]: from bqplot import *
from IPython.display import display
from bqplot.interacts import (
    FastIntervalSelector, IndexSelector, BrushIntervalSelector,
    BrushSelector, MultiSelector, LassoSelector, PanZoom, HandDraw
)
import ipywidgets as widgets
from ipywidgets import *

def run_scenario(scenario='1'):

    #determining which scenario this is
    if (scenario == '1'):
        # defining x,y coordinates, scenario 1
        x_coor = [18, 20, 22, 24, 26, 28, 30, 32, 34, 36, 38, 40]
        y_coor = [50000, 53000, 60000, 59000, 63000, 70000, 67000, 80000, 75000, 88000, 90000, 92000]
    elif (scenario=='2'):
        # defining x,y coordinates, scenario 1
        x_coor = [18, 26, 28, 30, 32, 34, 36, 38, 40, 20, 24]
        y_coor = [50000, 63000, 70000, 67000, 80000, 75000, 88000, 90000, 92000, 85000, 70000]
    elif (scenario=='3'):
        # defining x,y coordinates, scenario 1
        x_coor = [18, 20, 22, 32, 34, 36, 38, 40, 38, 36]
        y_coor = [50000, 53000, 60000, 80000, 75000, 88000, 90000, 92000, 50000, 52000]
    elif (scenario=='4'):
        # defining x,y coordinates, scenario 1
        x_coor = [18, 20, 22, 24, 26, 28, 30, 38, 40, 35, 22]
        y_coor = [50000, 53000, 60000, 59000, 63000, 70000, 67000, 90000, 92000, 55000, 90000]

    # defining linear scale for x,y axes
    x_sc = LinearScale(min=min(x_coor), max=max(x_coor))
    y_sc = LinearScale(min=min(y_coor), max=max(y_coor))

    # creating x,y axes using linear scales above
```

```

    ax_x = Axis(label='Age', scale=x_sc, grid_lines='dashed', num_ticks=8, u
    ↵label_color='blue')
    ax_y = Axis(label='Income', scale=y_sc, orientation='vertical', u
    ↵label_color='blue', label_offset='50px')

#adding scatter marks
names = []

for i in range(0, len(x_coor)):
    name = str(x_coor[i]) + ', ' + str(y_coor[i])
    names.append(name)

def_tt = Tooltip(fields=['x', 'y'], formats=['', '.2f'])

scatter_dots = Scatter(x=x_coor, y=y_coor, scales={'x': x_sc, 'y': y_sc}, u
    ↵tooltip=def_tt, unhovered_style={'opacity': 0.5})
#     scatter2 = Scatter(x=[x_coor[1]], y=[y_coor[1]], scales={'x': x_sc, 'y': u
    ↵y_sc}, enable_move=True)

# defining in regression line
sum_x = 0
sum_y = 0
sum_xy = 0
sum_x_squared = 0
sum_y_squared = 0
n = len(x_coor)

for i in range(0, len(x_coor)):
    sum_x += x_coor[i]
    sum_y += y_coor[i]
    sum_xy += sum_x * sum_y
    sum_x_squared += sum_x * sum_x
    sum_y_squared += sum_y * sum_y

a = ((sum_y * sum_x_squared) - (sum_x * sum_xy)) / ((n * sum_x_squared) - u
    ↵(sum_x * sum_x))
b = ((n * sum_xy) - (sum_x * sum_y)) / ((n * sum_x_squared) - (sum_x * u
    ↵sum_x))

y_hat_values = []
for x in x_coor:
    new_y_val = a + (b * x)
    y_hat_values.append(new_y_val)

```

```

regression_line = Lines(x=x_coor, y=y_hat_values, scales={'x': x_sc, 'y': y_sc}, colors=['red'], visible=False)

#user interaction, drawing a line

user_line_x = [19, 40]
user_line_y = [60000, 60000]

starting_user_line = Lines(x=user_line_x, y=user_line_y, scales={'x': x_sc, 'y': y_sc}, colors=['green'])

scatter = Scatter(x=user_line_x, y=user_line_y, scales={'x': x_sc, 'y': y_sc}, unhovered_style={'opacity': 0.5}, enable_move=True, colors=['black'])

#create figures and render
fig = Figure(marks=[scatter_dots, scatter, starting_user_line, regression_line], axes=[ax_x, ax_y], title="Find the Regression Line of Best Fit")

#add in dropdown menu
dropdown = widgets.Dropdown(
    options=['1', '2', '3', '4'],
    value=scenario,
    description='Scenario:',
    disabled=False,
)

out = Output()
display(out)

def handler(x):
    print('hi')

@out.capture()
def observe_scatter_x_y(point, change):

    x_array = point.x
    y_array = point.y
    find_change_x = change['point']['x']
    find_change_y = change['point']['y']
    hovered_point = point.hovered_point

    if (hovered_point == 0):

```

```

        starting_user_line.set_trait('x', [find_change_x, x_array[1]])
        starting_user_line.set_trait('y', [find_change_y, y_array[1]])
#
#           starting_user_line.y[0] = find_change_y
#           starting_user_line.x[0] = starting_user_line.x.pop(0)
#           starting_user_line.y = starting_user_line.y.pop(0)
#           starting_user_line.x = starting_user_line.x.insert(0,_
#           ↪find_change_y)
#           starting_user_line.y = starting_user_line.y.insert(0,_
#           ↪find_change_y)
        elif (hovered_point == 1):
            starting_user_line.set_trait('x', [x_array[0], find_change_x])
            starting_user_line.set_trait('y', [y_array[0], find_change_y])
#
#           starting_user_line.x = starting_user_line.x.pop(1)
#           starting_user_line.y = starting_user_line.y.pop(1)
#           starting_user_line.x = starting_user_line.x.insert(1,_
#           ↪find_change_y)
#           starting_user_line.y = starting_user_line.y.insert(1,_
#           ↪find_change_y)

#
#           if change['name'] == 'x':
#               line.x = line.x + change['new'] - change['old']
#           else:
#               line.y = line.y + change['new'] - change['old']

scatter.on_drag(observe_scatter_x_y)

@out.capture()
def on_dropdown_change(b):
    fig.close()
    button.close()
    dropdown.close()
    run_scenario(dropdown.value)

dropdown.observe(on_dropdown_change, 'value')

#add clicking action for submit
button = widgets.Button(description="See Answer", button_style='primary', □
↪margin_left='40px', layout=Layout(display='flex', button_color='white', □
↪justify_content='center', width='300px'))

def on_button_clicked(b):
    regression_line.set_trait('visible', True)

button.on_click(on_button_clicked)

```

```

box_layout = Layout(display='flex',
                    justify_content='center',
                    width='50%')

box = Box(children=[button], layout=box_layout)

display.dropdown, fig, box

run_scenario()

```

ModuleNotFoundError

↳ last

```

<ipython-input-28-0e1cbd099b01> in <module>()
      1
----> 2 from bqplot import *
      3 from IPython.display import display
      4 from bqplot.interacts import (
      5     FastIntervalSelector, IndexSelector, BrushIntervalSelector,

```

ModuleNotFoundError: No module named 'bqplot'

1.4.2 Residuals (5 mts)

The residual represents the difference between the actual observation and the one predicted by the linear regression line, as shown below:

Instructions: - Hover over the data points and you will see their (x,y) coordinates - Select (by clicking) the point you believe to have the largest residual - When you have selected, click ‘See Answer’ to see the result

```
[29]: from bqplot import *
from IPython.display import display
from bqplot.interacts import (
    FastIntervalSelector, IndexSelector, BrushIntervalSelector,
    BrushSelector, MultiSelector, LassoSelector, PanZoom, HandDraw
)
import ipywidgets as widgets
from ipywidgets import *
```

```

# defining x,y coordinates, scenario 1
x_coor = [18, 20, 22, 24, 26, 28, 30, 32, 34, 36, 38, 40, 39]
y_coor = [50000, 53000, 60000, 59000, 63000, 70000, 67000, 80000, 75000, 88000, 90000, 92000, 52000]

# defining linear scale for x,y axes
x_sc = LinearScale(min=min(x_coor), max=max(x_coor))
y_sc = LinearScale(min=min(y_coor), max=max(y_coor))

# creating x,y axes using linear scales above
ax_x = Axis(label='Age', scale=x_sc, grid_lines='dashed', num_ticks=8, label_color='blue')
ax_y = Axis(label='Income', scale=y_sc, orientation='vertical', label_color='blue', label_offset='50px')

#adding scatter marks
names = []

for i in range(0, len(x_coor)):
    name = str(x_coor[i]) + ', ' + str(y_coor[i])
    names.append(name)

def_tt = Tooltip(fields=['x', 'y'], formats=['', '.2f'])

scatter = Scatter(x=x_coor, y=y_coor, scales={'x': x_sc, 'y': y_sc}, opacity=[0.5], tooltip=def_tt, selected_style={'fill':'red', 'stroke': 'red'}, selected=[5])

out = Output()
display(out)

@out.capture()
def on_scatter_click(marks, dot):
    marks.set_trait('selected', [dot['data']['index']])

scatter.on_element_click(on_scatter_click)

# defining in regression line
sum_x = 0
sum_y = 0
sum_xy = 0
sum_x_squared = 0
sum_y_squared = 0
n = len(x_coor)

```

```

for i in range(0, len(x_coor)):
    sum_x += x_coor[i]
    sum_y += y_coor[i]
    sum_xy += sum_x * sum_y
    sum_x_squared += sum_x * sum_x
    sum_y_squared += sum_y * sum_y

a = ((sum_y * sum_x_squared) - (sum_x * sum_xy)) / ((n * sum_x_squared) - (sum_x * sum_x))
b = ((n * sum_xy) - (sum_x * sum_y)) / ((n * sum_x_squared) - (sum_x * sum_x))

y_hat_values = []

for x in x_coor:
    new_y_val = a + (b * x)
    y_hat_values.append(new_y_val)

regression_line = Lines(x=x_coor, y=y_hat_values, scales={'x': x_sc, 'y': y_sc}, colors=['red'])

# create figures and render
fig = Figure(marks=[scatter, regression_line], axes=[ax_x, ax_y], title='Find the Largest Residual')

# add clicking action for submit
button = widgets.Button(description="See Answer", button_style='primary', margin_left='40px', layout=Layout(display='flex', button_color='white', justify_content='center', width='300px'))

@out.capture()
def on_button_clicked(b):
    answer = 12
    scatter.set_trait('selected', [answer])
    residual_line=Lines(x=[x_coor[answer], x_coor[answer]], y=[y_coor[answer], y_hat_values[answer]], scales={'x': x_sc, 'y': y_sc}, colors=['red'])
    fig.set_trait('marks', [scatter, regression_line, residual_line])

button.on_click(on_button_clicked)

box_layout = Layout(display='flex', justify_content='center', width='50%')

```

```

box = Box(children=[button], layout=box_layout)

display(fig, box)

```

ModuleNotFoundError

Traceback (most recent call last)

```

<ipython-input-29-7752376e1dfa> in <module>()
----> 1 from bqplot import *
      2 from IPython.display import display
      3 from bqplot.interacts import (
      4     FastIntervalSelector, IndexSelector, BrushIntervalSelector,
      5     BrushSelector, MultiSelector, LassoSelector, PanZoom, HandDraw

```

ModuleNotFoundError: No module named 'bqplot'

1.5 Interpreting the output of a linear model (10 mts)

The linear model of pay vs. age can be fit as follows:

```
[30]: model1 = 'pay~age'
lm1 = sm.ols(formula = model1, data = Data).fit()
print(lm1.summary())
```

OLS Regression Results

Dep. Variable:	pay	R-squared:	0.285			
Model:	OLS	Adj. R-squared:	0.284			
Method:	Least Squares	F-statistic:	397.5			
Date:	Wed, 13 Nov 2019	Prob (F-statistic):	1.04e-74			
Time:	16:22:35	Log-Likelihood:	-11384.			
No. Observations:	1000	AIC:	2.277e+04			
Df Residuals:	998	BIC:	2.278e+04			
Df Model:	1					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	6.206e+04	2062.885	30.085	0.000	5.8e+04	6.61e+04
age	939.2501 ^{USD}	47.109	19.938	0.000	846.806	1031.694

Omnibus:	6.360	Durbin-Watson:	1.905
Prob(Omnibus):	0.042	Jarque-Bera (JB):	6.421
Skew:	0.182	Prob(JB):	0.0403
Kurtosis:	2.853	Cond. No.	134.
=====			

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

1.5.1 What should we focus on in the regression table?

Although the output table above contains a lot of information, we only need to focus on a small number of quantities. These are the coefficients (the estimates of β_0 and β_1), as well as their p - and R -squared values.

1.5.2 Coefficients

The intercept β_0 is about 62060 USD. This can be thought of as the baseline pay. (Frequently, the intercept does not have a meaningful interpretation – that is okay.) The slope (the coefficient β_1 for the age) is 939.25. The interpretation of this coefficient is as explained before: if an employee becomes one year older, his or her pay is expected to increase by 939.25 USD on average.

1.5.3 p -values

The null hypothesis we are testing here is

$$H_0 : \beta_1 = 0$$

Si $p < 0.05$, el coeficiente β_1 es estadísticamente significativo

and the alternative is

$$H_a : \beta_1 \neq 0.$$

The p -value of β_1 (given under the column: " $P > |t|$ ") is 0.000. Thus, it is statistically significant at the 0.05 level, and we reject the null hypothesis. This implies that age does explain some of the differences in pay.

1.6 R -squared (15 mts)

One of the key quantities that should be paid attention to while interpreting a regression table is the quantity R -squared. Note that the table shows R -squared and adjusted R -squared. We will focus on R -squared. This quantity is always going to be between 0 and 1. For the pay vs. age model, this quantity is $0.285 = 28.5\%$. In this organization, there is a large variation in pay – different employees are paid different amounts.

An R -squared of 28.5% in this linear model means that this observed variation in pay is not due to random chance; rather, the age factor systematically explains 28.5% of this variation in pay. The higher the R -squared, the higher the percentage of observed variation that can be explained by the

model. Since model1 only explains about 28.5% of the variation, this motivates us to investigate if factors other than age can be used to explain the pay differences.

1.6.1 Relation between correlation and R -squared

```
[31]: corr_age_pay = np.corrcoef(Data['pay'], Data['age'])[0,1]
corr_age_pay*corr_age_pay
```

```
[31]: 0.2848512885110829
```

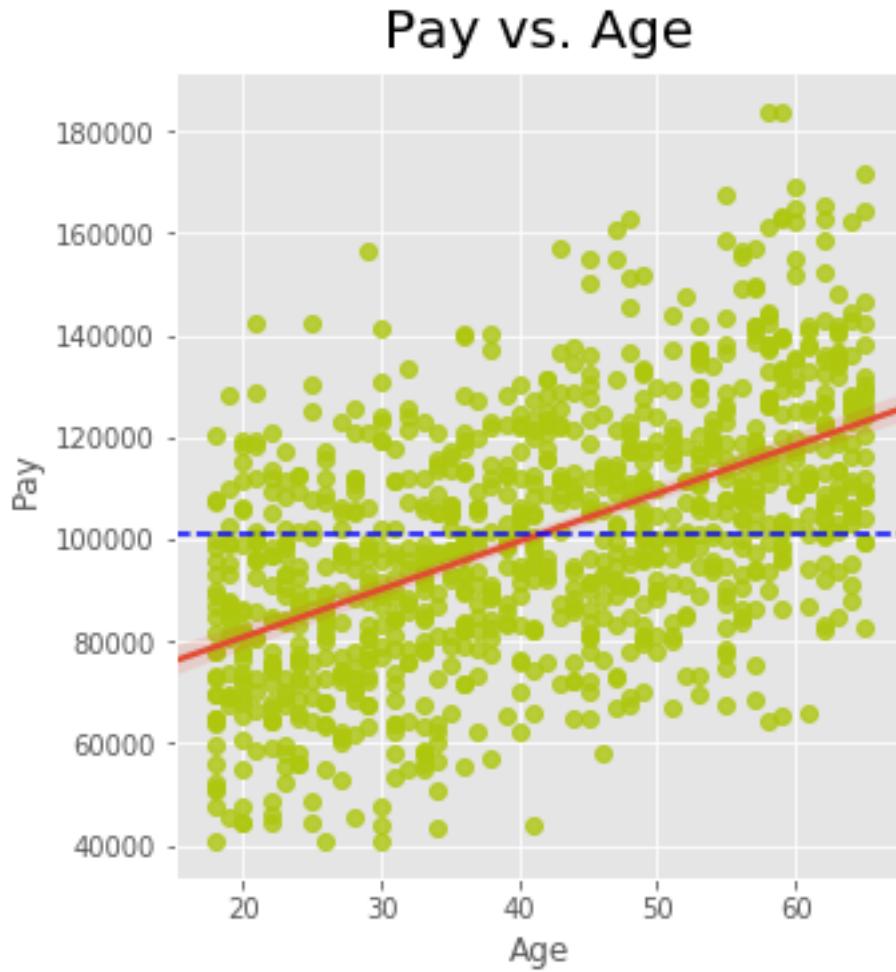
1.6.2 Geometric interpretation of R -Squared

The formula for R -squared is the following:

$$R^2 = 1 - \frac{\text{RSS Residual Sum of Squares}}{\text{TSS Total sum of squares}}$$

Here **RSS** and **TSS** denote the Residual and Total Sum of Squares respectively. To understand the above formula carefully, let us plot pay vs. age again.

```
[32]: sns.lmplot(x = 'age', y = 'pay', data = Data, scatter_kws = {'color': (174/
    ↪255,199/255,14/255)})
plt.title("Pay vs. Age", fontsize=20, verticalalignment='bottom')
plt.axhline(Data['pay'].mean(), ls='--', color = 'b')
plt.xlabel("Age")
plt.ylabel("Pay");
```



Consider the naive model0 given by

$$\text{pay} = \beta_0 + \text{error}$$

The best guess for β_0 is just the average pay. This is given by the blue line. The sum of the squares of the residuals here is called the Total Sum of Squares (TSS). R-squared measures how well the regression line of model1 (pay vs. age) given by the red line explains the observed variation as compared to model0. The sum of residual squares for this model is the RSS.

1.6.3 Calculating R-Squared by writing code

```
[33]: # you can 'peek' inside the object lm1 using the dir command
dir(lm1)
# for instance if you want to know what the fitted values are
## lm.fittedvalues
```

```
[33]: ['HC0_se',
'HC1_se',
'HC2_se',
'HC3_se',
'_HCCM',
'__class__',
'__delattr__',
'__dict__',
'__dir__',
'__doc__',
'__eq__',
'__format__',
'__ge__',
'__getattribute__',
'__gt__',
'__hash__',
'__init__',
'__init_subclass__',
'__le__',
'__lt__',
'__module__',
'__ne__',
'__new__',
'__reduce__',
'__reduce_ex__',
'__repr__',
'__setattr__',
'__sizeof__',
'__str__',
'__subclasshook__',
'__weakref__',
'_cache',
'_data_attr',
'_get_robustcov_results',
'_is_nested',
'_wexog_singular_values',
'aic',
'bic',
'bse',
'centered_tss',
'compare_f_test',
'compare_lm_test',
'compare_lr_test',
'condition_number',
'conf_int',
'conf_int_el',
'cov_HC0',
```

```
'cov_HC1',
'cov_HC2',
'cov_HC3',
'cov_kwds',
'cov_params',
'cov_type',
'df_model',
'df_resid',
'diagn',
'eigenvals',
'el_test',
'ess',
'f_pvalue',
'f_test',
'fittedvalues',
'fvalue',
'get_influence',
'get_prediction',
'get_robustcov_results',
'initialize',
'k_constant',
'llf',
'load',
'model',
'mse_model',
'mse_resid',
'mse_total',
'nobs',
'normalized_cov_params',
'outlier_test',
'params',
'predict',
'pvalues',
'remove_data',
'resid',
'resid_pearson',
'rsquared',
'rsquared_adj',
'save',
'scale',
'ssr',
'summary',
'summary2',
't_test',
't_test_pairwise',
'tvalues',
'uncentered_tss',
```

```
'use_t',
'wald_test',
'wald_test_terms',
'wresid']
```

```
[34]: ## Write code to find residuals from fitted values
res = Data['pay'] - lm1.fittedvalues
res[0:5]
```

```
[34]: 0    -26666.935762
1     37818.313854
2     19568.814110
3     37387.563982
4     22301.063216
dtype: float64
```

```
[35]: ## check if this is the same as that stored in lm1.resid
lm1.resid[0:5]
```

```
[35]: 0    -26666.935762
1     37818.313854
2     19568.814110
3     37387.563982
4     22301.063216
dtype: float64
```

```
[36]: # The mean for pay
mean_pay = Data['pay'].mean()

## The total sum of squares:
TSS = sum((Data['pay'] - mean_pay)*(Data['pay'] - mean_pay))
TSS

## Residual sum of squares
RSS = sum(res*res)
```

```
[37]: R_squared = (TSS - RSS)/TSS
R_squared
```

```
[37]: 0.28485128851108443
```

1.7 Looking at age and gender: Categorical variables (10 mts)

Now that we have seen that age explains some of the relationship with pay, let us consider a model in which we take age and gender into account simultaneously. Age is a numeric variable (e.g., 26.5, 32). In contrast, gender takes only two values – male and female. Such variables are called

categorical variables. The way we interpret the coefficients of factor variables in the linear model is slightly different from those of numeric variables:

```
[38]: model2 = 'pay~age + gender'
lm2      = sm.ols(formula = model2, data = Data).fit()
print(lm2.summary())
```

OLS Regression Results					
Dep. Variable:	pay	R-squared:	0.319		
Model:	OLS	Adj. R-squared:	0.317		
Method:	Least Squares	F-statistic:	233.2		
Date:	Wed, 13 Nov 2019	Prob (F-statistic):	8.10e-84		
Time:	16:22:44	Log-Likelihood:	-11359.		
No. Observations:	1000	AIC:	2.272e+04		
Df Residuals:	997	BIC:	2.274e+04		
Df Model:	2				
Covariance Type:	nonrobust				
<hr/>					
==					
	coef	std err	t	P> t	[0.025
0.975]					
<hr/>					
--					
Intercept	5.674e+04	2151.480	26.373	0.000	5.25e+04
6.1e+04					
gender[T.Male]	9279.3180	1317.787	7.042	0.000	6693.364
1.19e+04					
age	948.5266	46.022	20.610	0.000	858.216
1038.837					
<hr/>					
Omnibus:	9.898	Durbin-Watson:	1.871		
Prob(Omnibus):	0.007	Jarque-Bera (JB):	9.345		
Skew:	0.197	Prob(JB):	0.00935		
Kurtosis:	2.737	Cond. No.	148.		
<hr/>					

Warnings:

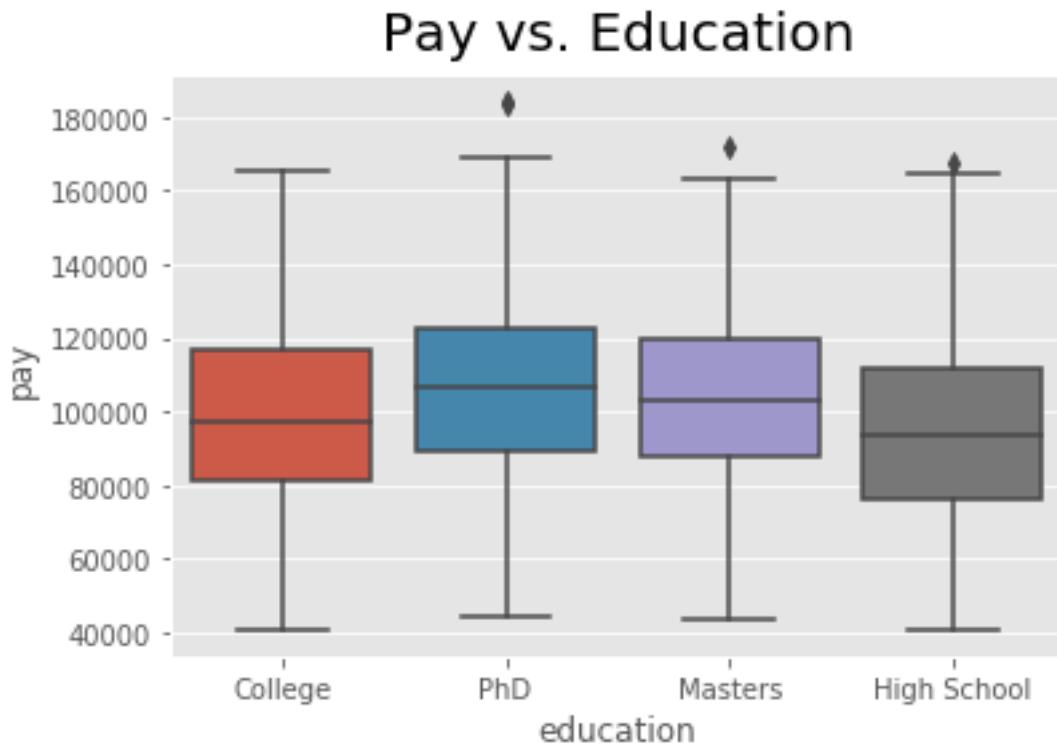
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

The interpretation of the coefficient of age is the same as before: if age increases by one year, the pay is expected to increase by 948.5 USD. Now, focus on the coefficient of gender. It shows male (T.male) only, because the category female is taken as the default category. (Note that the choice of default category doesn't matter – we could easily have chosen to make male the default category and hence the coefficient for gender would be T.female.) The coefficient 9279.3180 is interpreted as follows: for employees of the same age, on average, men make 9279.3180 USD more than women.

But we still haven't satisfactorily answered our main question yet. So far, we have only accounted

for age in addition to gender for explaining pay gaps. There are still a few more factors that could affect pay. We consider education next. The following plot shows that employees with at least a college education are paid more:

```
[39]: sns.boxplot(x='education', y = 'pay', data = Data)
plt.title("Pay vs. Education", fontsize=20, verticalalignment='bottom');
```



1.7.1 Exercise 4: (5 mts)

Give the code for fitting a linear model with pay as the dependent variable and age, gender, and education as the independent variables. Compare the R -squared of this with that of model2. What conclusions can we make?

Answer. The R -squared has increased from 28.5% to 34.5%. This means education helps explain a bit more variation that is not explained by age and gender.

```
[40]: model3 = 'pay~age + gender + education'
lm3    = sm.ols(formula = model3, data = Data).fit()
print(lm3.summary())
```

```
OLS Regression Results
=====
Dep. Variable:          pay    R-squared:       0.345
```

Model:	OLS	Adj. R-squared:	0.342		
Method:	Least Squares	F-statistic:	104.7		
Date:	Wed, 13 Nov 2019	Prob (F-statistic):	8.74e-89		
Time:	16:22:46	Log-Likelihood:	-11340.		
No. Observations:	1000	AIC:	2.269e+04		
Df Residuals:	994	BIC:	2.272e+04		
Df Model:	5				
Covariance Type:	nonrobust				
<hr/>					
<hr/>					
	coef	std err	t	P> t	[0.025
0.975]					
<hr/>					
<hr/>					
Intercept	5.562e+04	2378.880	23.380	0.000	5.1e+04
6.03e+04					
gender[T.Male]	8789.0942	1297.996	6.771	0.000	6241.967
1.13e+04					
education[T.High School]	-3906.3104	1817.025	-2.150	0.032	-7471.956
-340.664					
education[T.Masters]	4171.0770	1836.165	2.272	0.023	567.871
7774.283					
education[T.PhD]	6727.2845	1867.621	3.602	0.000	3062.352
1.04e+04					
age	942.4566	45.209	20.846	0.000	853.740
1031.173					
<hr/>					
Omnibus:	8.469	Durbin-Watson:	1.891		
Prob(Omnibus):	0.014	Jarque-Bera (JB):	8.283		
Skew:	0.195	Prob(JB):	0.0159		
Kurtosis:	2.782	Cond. No.	206.		
<hr/>					

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

1.7.2 Exercise 5:

Which of the following statements is false? After accounting for age and gender:

- (a) Employees with college education are paid 3906 USD more on average than those with just a high school degree.
- (b) Employees with just a high school degree are paid 7500.1 USD less on average compared to those with a masters degree.
- (c) Employees with a PhD degree earn on average 2556.20 USD more than those with a masters

degree.

- (d) Employees with a masters degree earn on average 8077.3 USD more than those with just a high school degree.

Answer. (a) Correct. For the linear model: ‘pay~age + gender + education’, education is a factor variable, and college education is the default due to the alphabetical order. We get the coefficient of high school to be -3906.3104.

- (b) Incorrect! For the linear model: ‘pay~age + gender + education’, the difference of the coefficients of high school and masters degree gives the required answer. This is because (high school - college) - (masters - college) = (high school - masters).
- (c) Incorrect! For the linear model: ‘pay~age + gender + education’, the difference of the coefficients of masters and PhD degree gives the required answer. This is because (PhD - college) - (masters - college) = (PhD - masters).
- (d) Incorrect! For the linear model: ‘pay~age + gender + education’, the difference of the coefficients of high school and masters degree gives the required answer. This is because (high school - college) - (masters - college) = (high school - masters).

1.8 Integrated model accounting for all variables (15 mts)

Let us account for all the other factors that could explain pay gaps at once. What can you conclude?:

```
[41]: model4 = 'pay~jobtitle + age + performance + education + department + seniority  
          ↵+ gender'  
lm4      = sm.ols(formula = model4, data = Data).fit()  
print(lm4.summary())
```

OLS Regression Results

Dep. Variable:	pay	R-squared:	0.841
Model:	OLS	Adj. R-squared:	0.838
Method:	Least Squares	F-statistic:	259.6
Date:	Wed, 13 Nov 2019	Prob (F-statistic):	0.00
Time:	16:22:48	Log-Likelihood:	-10631.
No. Observations:	1000	AIC:	2.130e+04
Df Residuals:	979	BIC:	2.141e+04
Df Model:	20		
Covariance Type:	nonrobust		

	coef	std err	t	P> t
[0.025 0.975]				
-----	-----	-----	-----	-----
Intercept	2.203e+04	1933.534	11.392	0.000
1.82e+04 2.58e+04				
jobtitle[T.Driver]	-3928.9812	1447.166	-2.715	0.007

-6768.886	-1089.076				
jobtitle[T.Financial Analyst]	3417.7090	1388.276	2.462	0.014	
693.370	6142.048				
jobtitle[T.Graphic Designer]	-2457.6992	1420.886	-1.730	0.084	
-5246.031	330.633				
jobtitle[T.IT]	-2149.7022	1427.414	-1.506	0.132	
-4950.846	651.442				
jobtitle[T.Manager]	3.16e+04	1471.445	21.476	0.000	
2.87e+04	3.45e+04				
jobtitle[T.Marketing Associate]	-1.701e+04	1385.795	-12.277	0.000	
-1.97e+04	-1.43e+04				
jobtitle[T.Sales Associate]	263.4456	1435.261	0.184	0.854	
-2553.096	3079.988				
jobtitle[T.Software Engineer]	1.339e+04	1413.182	9.473	0.000	
1.06e+04	1.62e+04				
jobtitle[T.Warehouse Associate]	-564.0171	1452.967	-0.388	0.698	
-3415.305	2287.271				
education[T.High School]	-1435.1693	908.773	-1.579	0.115	
-3218.536	348.198				
education[T.Masters]	4717.9971	914.421	5.160	0.000	
2923.546	6512.448				
education[T.PhD]	6026.2867	929.705	6.482	0.000	
4201.842	7850.731				
department[T.Engineering]	3267.9358	1036.509	3.153	0.002	
1233.901	5301.970				
department[T.Management]	2957.8290	1033.031	2.863	0.004	
930.619	4985.039				
department[T.Operations]	-481.4968	1014.616	-0.475	0.635	
-2472.570	1509.577				
department[T.Sales]	6193.4962	1020.679	6.068	0.000	
4190.526	8196.466				
gender[T.Male]	392.3244	715.798	0.548	0.584	
-1012.351	1797.000				
age	948.9464	22.527	42.126	0.000	
904.740	993.152				
performance	1156.8797	228.188	5.070	0.000	
709.085	1604.674				
seniority	9903.7103	231.490	42.782	0.000	
9449.436	1.04e+04				
<hr/>					
Omnibus:	3.500	Durbin-Watson:	1.982		
Prob(Omnibus):	0.174	Jarque-Bera (JB):	3.555		
Skew:	-0.130	Prob(JB):	0.169		
Kurtosis:	2.866	Cond. No.	453.		
<hr/>					

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly

El genero no es
significativamente
estadisticamente

specified.

1.8.1 Exercise 6:

Taking all of the factors in to account as in model4, which of the following jobs pay the highest?

- (a) Marketing Associate
- (b) Software Engineer
- (c) Manager
- (d) Graphic Designer

Answer.

- (a) Incorrect.
- (b) Incorrect.
- (c) Correct!
- (d) Incorrect.

1.9 Revisiting the question of pay discrimination

Now that we have looked at and accounted for various attributes that are correlated with pay, let us now revisit the question of what is driving pay discrimination. model4, which accounts for all of the variables in the dataset, has an R -squared of 84.1%. This is definitely a vast improvement from model1 (pay vs. age) which has an R -squared of 28.5%. Thus compared to model1, model4 explains $84.1\% - 28.5\% = 55.6\%$ more of the variation seen in the data.

1.9.1 Exercise 7:

Based on the analysis we have done so far (model1 - model4), consider the following statements.

I. After accounting for job title, education, performance and age, the proportion of pay difference attributable solely to gender is small. ✓

- II. There is evidence that pay discrimination between men and women is due solely to gender. •
- III. There is reason to believe that there could be a disproportionate amount of women in lower paying jobs such as marketing, while there could be more men in higher paying jobs, such as manager. ✓

Choose the correct answer:

- (a) I is correct, II and III are incorrect. •
- (b) II is correct, I and III are incorrect. •
- (c) I is incorrect, II and III are correct. •
- (d) I and III are correct, II is incorrect. ✓

Answer.

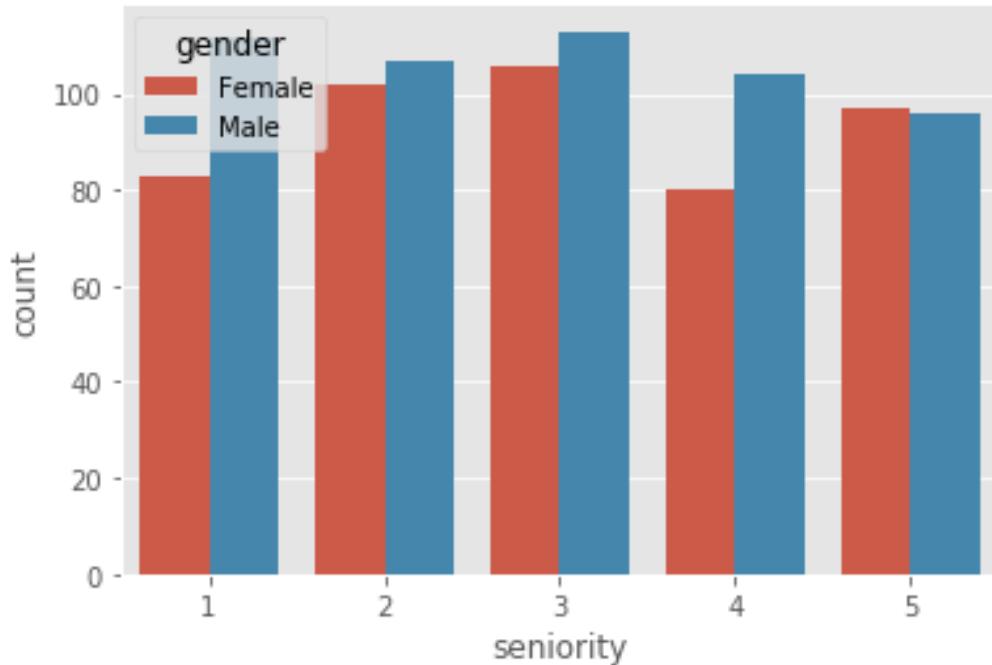
- (a) Incorrect. We know from model4 that the coefficient of gender is not significant (392.3244 USD, and this coefficient is not significant). Thus, I is correct and II is incorrect. But from the first plot we made, we saw that the average salary difference between men and women in this organization is about 8500 USD. The key explanation for this discrepancy is that women are in lower paying jobs than men. Thus III is correct.
- (b) Incorrect. We know from model4 that the coefficient of gender is not significant (392.3244 USD, and this coefficient is not significant). Thus, I is correct and II is incorrect. But from the first plot we made, we saw that the average salary difference between men and women in this organization is about 8500 USD. The key explanation for this discrepancy is that women are in lower paying jobs than men. Thus III is correct.
- (c) Incorrect. We know from model4 that the coefficient of gender is not significant (392.3244 USD, and this coefficient is not significant). Thus, I is correct and II is incorrect. But from the first plot we made, we saw that the average salary difference between men and women in this organization is about 8500 USD. The key explanation for this discrepancy is that women are in lower paying jobs than men. Thus III is correct.
- (d) Correct! We know from model4 that the coefficient of gender is not significant (392.3244 USD, and this coefficient is not significant). Thus, I is correct and II is incorrect. But from the first plot we made, we saw that the average salary difference between men and women in this organization is about 8500 USD. The key explanation for this discrepancy is that women are in lower paying jobs than men. Thus III is correct.

1.10 Investigating the distribution of gender across seniority and job types

Motivated by Exercise 7, let's look into how women are distributed across various factors. The following plot shows that men and women are similarly distributed by seniority:

```
[42]: sns.countplot(x='seniority', hue = 'gender', data = Data)
plt.title("Seniority distribution for men vs. women", fontsize=20, u
          ↴verticalalignment='bottom');
```

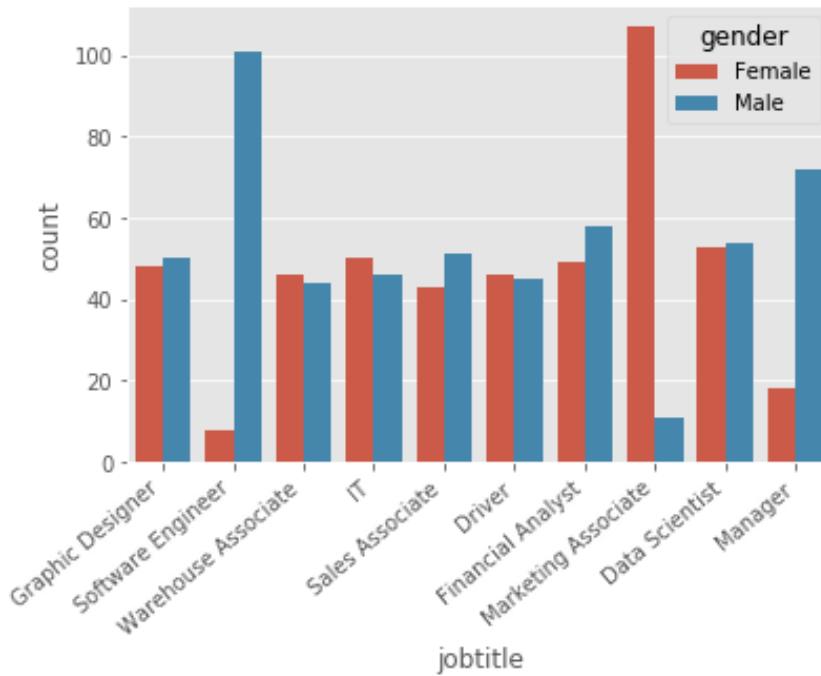
Seniority distribution for men vs. women



However, looking at the distribution of women across various job types tells a different story. From the following barplot, we see that women are disproportionately underrepresented in software engineering and managerial roles. Also, women are disproportionately overrepresented in the marketing associate job title:

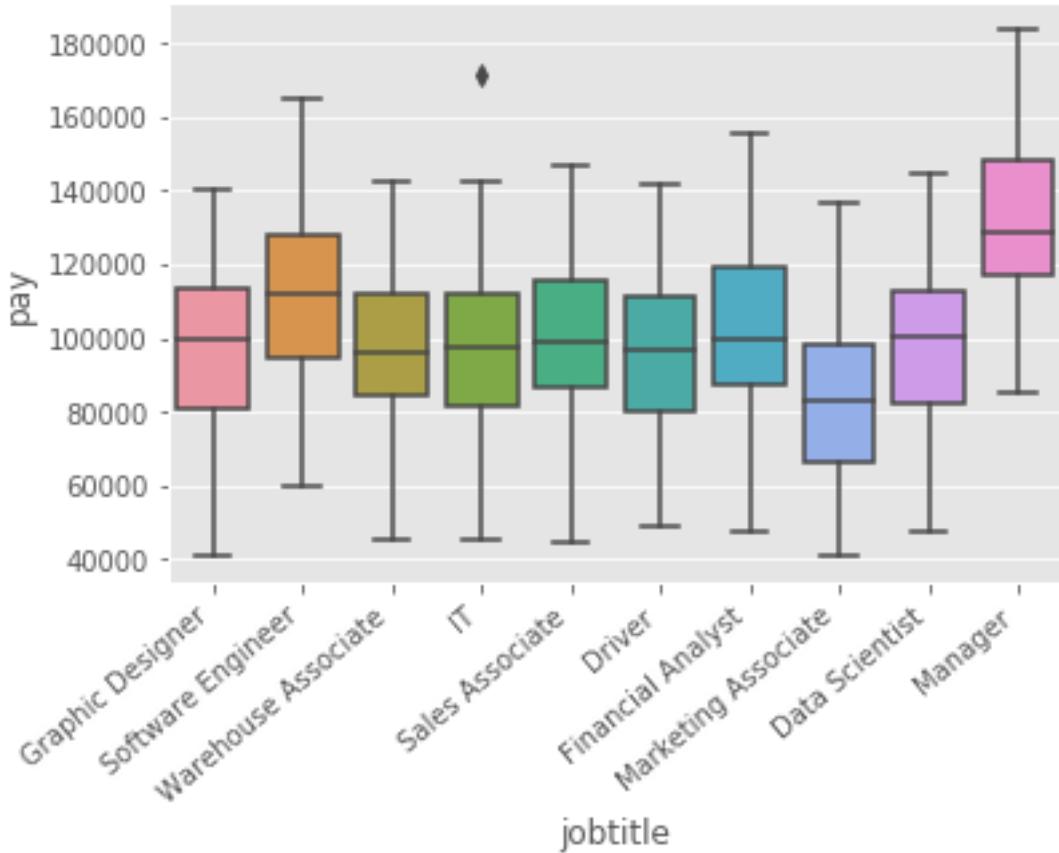
```
[43]: pl = sns.countplot(x="jobtitle", hue= "gender", data= Data)
pl.set_xticklabels(pl.get_xticklabels(), rotation=40, ha="right");
plt.title(" Distribution of men vs. women across various roles", fontsize=20,
          verticalalignment='bottom');
```

Distribution of men vs. women across various roles



```
[44]: pl2 = sns.boxplot(x="jobtitle", y = "pay", data = Data)
pl2.set_xticklabels(pl2.get_xticklabels(), rotation=40, ha="right");
plt.title("Pay distribution across jobtitles", fontsize=20,
          verticalalignment='bottom');
```

Pay distribution across jobtitles



The above chart shows that manager and marketing associate have the highest and lowest average salaries, respectively. We also see from the previous plot that the proportion of women are the lowest and the highest in manager and marketing associate roles, respectively. This visually explains the output of our regression model about why there is no evidence of pay discrepancy after taking into account all of these factors.

1.11 Conclusions (15 mts)

We used the techniques of linear regression to determine whether or not there existed gender-based pay discrimination within your organization. We modeled the effect of various independent variables (in this case seniority, age, performance, and job title) to explain the observed variation in a dependent variable (in this case, pay). We looked at the R -squared statistic of our linear models to help us measure what percentage of observed variation in pay was explained by the independent variables.

We saw that the difference in average salary between men and women is about 8500 USD in this data. However, this difference became 400 USD and is statistically indistinguishable (p -value = 0.584) from zero after controlling for the other factors correlated with pay. However, further

exploration of the data suggested that women are disproportionately overrepresented in the lowest-paying jobs, while men are disproportionately overrepresented in the highest-paying jobs.

Thus, an investigation into the hiring, promotion, and job placement practices of men and women is warranted. In your report to the human resources department, you should ask them to look into the following questions:

1. Are women choosing to or being forced to take lower paying jobs?
2. Are women being discriminated against in hiring processes for higher-paying jobs?

1.12 Takeaways (5 mts)

In this case, you learned how to leverage your skills in exploratory data analysis to build an effective linear model that accounted for several factors related to the outcome of interest (pay). Crucially, we learned that:

1. Looking directly at the relationship between the outcome of interest and the independent variable of interest is not enough – there may be several confounding factors.
2. Conducting EDA before doing any model-building is important to discover and account for these confounding factors that could be driving the differences in the outcome of interest.
3. R -squared is an important quantity that explains how well your model explained the observed variation. It can be used to compare different models.

→ EDA: Exploratory Data Analysis