

# How should we price homes in Seattle?

```
In [1]: # Ignore user warnings
import warnings
warnings.simplefilter("ignore", UserWarning)

# Load relevant packages
import pandas as pd
import numpy as np
from scipy import stats
import matplotlib.pyplot as plt
import seaborn as sns
sns.set_style('darkgrid')
import statsmodels.api as sm
import statsmodels.formula.api as smf
import pylab

%matplotlib inline
plt.style.use('ggplot')
```

## Introduction (5 mts)

**Business Context.** You have been hired as a data scientist by a large real estate company in their Seattle office. Your job is to assist Seattle residents willing to sell their home with determining an optimal price to sell their property at in order to maximize their proceeds while still being able to find willing buyers. To do this, the firm would like you to build a pricing model for Seattle real estate, in order to maximize the probability of helping residents close sales (and thus maximizing commissions for the firm).

**Business Problem.** Your task is to **build a model that uses past sales data in Seattle to recommend an optimal sell price for any particular property.**

**Analytical Context.** The provided dataset was retrieved from Kaggle (<https://www.kaggle.com/harlfoxem/housesalesprediction>) and includes sales prices of houses in the state of Washington (King county, where Seattle is located) between May 2014 and May 2015. As we have learned, the primary tool to predict a response variable is the multiple regression model. However, sometimes the assumptions of a linear model are not met by our data. We will learn a set of strategies to mitigate some common issues that appear during regression analysis.

The case is structured as follows: you will (1) conduct basic EDA of some of the variables to determine that standard linear regression is not sufficient; (2) learn about variable transformations and use these to improve the initial model; and finally (3) learn how to incorporate interaction effects (which are themselves a form of variable transformation involving two or more variables) into our model.

## Data exploration (15 mts)

Let's start by reviewing the columns of the dataset and what they mean:

1. **id**: identification for a house
2. **date**: date house was sold
3. **price**: price house was sold at
4. **bedrooms**: number of bedrooms
5. **bathrooms**: number of bathrooms
6. **sqft\_living**: square footage of the home
7. **sqft\_lot**: square footage of the lot
8. **floors**: total floors (levels) in house
9. **waterfront**: whether or not the house has a view of a waterfront
10. **view**: whether or not the house has been viewed
11. **condition**: how good the condition of the house is
12. **grade**: overall grade given to the housing unit, based on King County grading system
13. **sqft\_above**: square footage of the house apart from basement
14. **sqft\_basement**: square footage of the basement
15. **yr\_builtin**: year house was built

16. **yr\_renovated**: year house was renovated
17. **zipcode**: zipcode of the house
18. **lat**: latitude coordinate of the house
19. **long**: longitude coordinate of the house

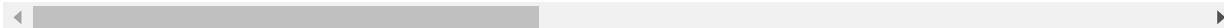
```
In [2]: houses = pd.read_csv('kc_house_data.csv')
```

```
In [3]: houses.head()
```

Out[3]:

	<b>id</b>	<b>date</b>	<b>price</b>	<b>bedrooms</b>	<b>bathrooms</b>	<b>sqft_living</b>	<b>sqft_lot</b>	<b>floors</b>	<b>w</b>
0	7129300520	20141013T000000	221900.0	3	1.00	1180	5650	1.0	
1	6414100192	20141209T000000	538000.0	3	2.25	2570	7242	2.0	
2	5631500400	20150225T000000	180000.0	2	1.00	770	10000	1.0	
3	2487200875	20141209T000000	604000.0	4	3.00	1960	5000	1.0	
4	1954400510	20150218T000000	510000.0	3	2.00	1680	8080	1.0	

5 rows × 21 columns



## Exercise 1: (5 mts)

Analyze the distribution of house prices using `.describe()`, a QQ plot, and a histogram plot. Does it look Gaussian?

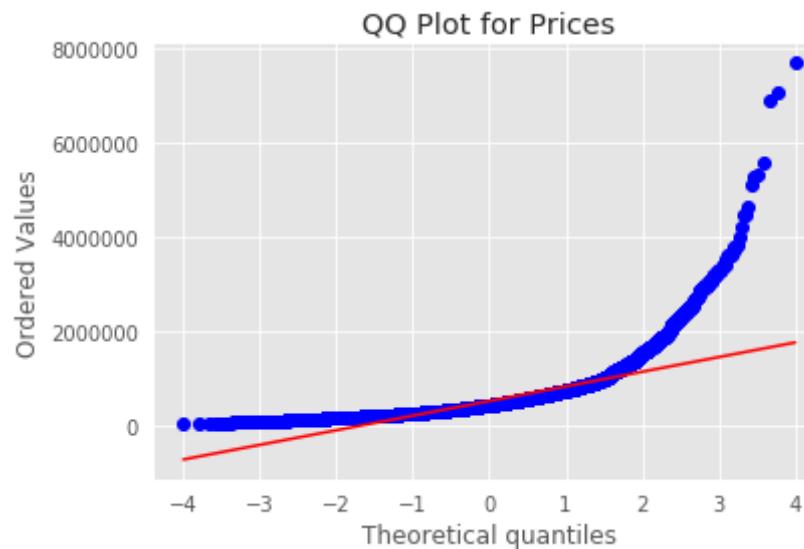
**Answer.** One possible solution is shown below:

```
In [4]: houses['price'].describe()
```

```
Out[4]: count    2.161300e+04
mean      5.400881e+05
std       3.671272e+05
```

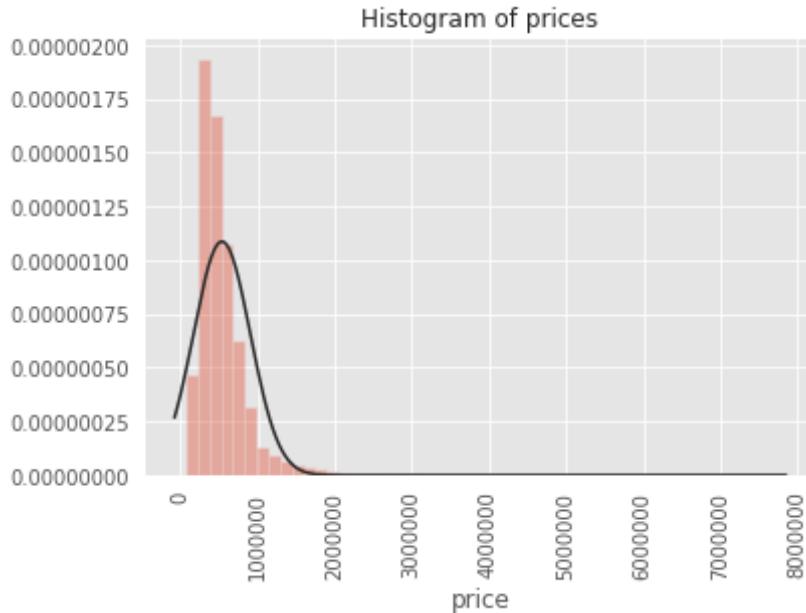
```
min      7.500000e+04
25%     3.219500e+05
50%     4.500000e+05
75%     6.450000e+05
max     7.700000e+06
Name: price, dtype: float64
```

```
In [5]: ## QQ plot of price
stats.probplot(houses['price'], dist = "norm", plot = plt)
plt.title("QQ Plot for Prices")
plt.show()
```



```
In [6]: ## histogram plot of price
sns.distplot(houses['price'], fit=stats.laplace, kde=False)
sns.distplot(houses['price'], fit=stats.norm, kde=False)
sns.set(color_codes=True)
plt.xticks(rotation=90)
plt.title("Histogram of prices")
```

```
Out[6]: Text(0.5, 1.0, 'Histogram of prices')
```



The distribution does not look Gaussian. Looking at both the QQ plot and the histogram, we can see that the distribution of our data is heavily skewed.

## Exercise 2: (5 mts)

Analyze the relationship between house prices and price per square foot of living space. What can you conclude? (Hint: use the `lmplot()` function in the `seaborn` library.)

**Answer.** We can create a regression plot to visualize this relationship:

```
In [7]: ## linear relation between sqft_living and price
sns.lmplot(x='sqft_living',y='price',data=houses,
            line_kws = {'color': "red"}, aspect= 2)
plt.title("Price vs. Sqft_living");
```



Given the way that house price vs. price per square foot seems to "fan out", we see that the relationship does not appear to be linear. In fact, it is not immediately obvious what sort of relationship is exhibited at all here.

## Variable transformation (30 mts)

We have seen in Exercise 1 that the distribution of house prices is not Gaussian, and that this may be contributing to the "fanning out" effect we observed in Exercise 2. We want to find a way to remove the "fanning out" effect, as it implies that a linear fit becomes less and less suitable, with higher and higher variance from the line of best fit for large values of the predictor and response variables. A common method of addressing this issue is to transform the response variable and/or the predictor variable. Such a **variable transformation** involves applying a known function to one or more of these variables to achieve conditions that are suitable for the application of a linear model.

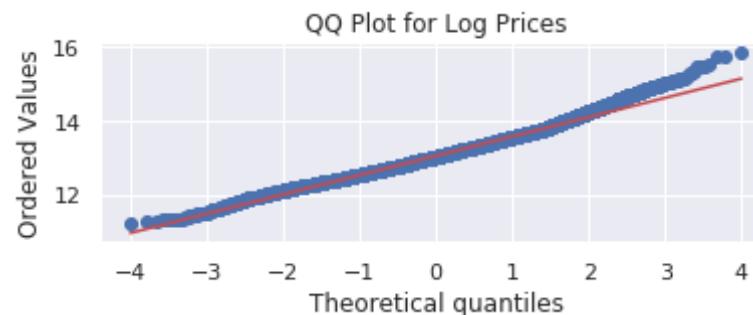
Typical mathematical functions used to transform variables include powers (quadratic, cubic, square root, etc.), logarithms, and trigonometric functions. Let's start with the logarithmic transformation to see if we can achieve some results.

### Exercise 3: (5 mts)

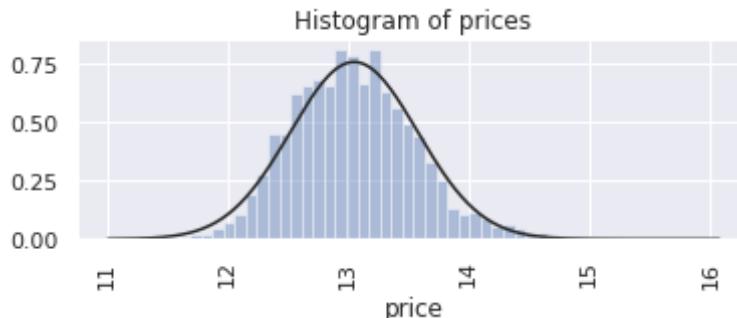
Take the logarithm of house prices and create plots to ascertain if this makes the distribution of the transformed variable roughly Gaussian.

**Answer.** One possible solution is shown below:

```
In [8]: ## QQ plot of price
plt.subplot(2,1,1)
stats.probplot(np.log(houses['price']), dist = "norm", plot = plt)
plt.title("QQ Plot for Log Prices")
plt.show()
plt.subplot(2,1,2)
sns.distplot(np.log(houses['price']), fit=stats.norm, kde=False)
sns.set(color_codes=True)
plt.xticks(rotation=90)
plt.title("Histogram of prices")
```



```
Out[8]: Text(0.5, 1.0, 'Histogram of prices')
```



```
In [9]: np.log(houses['price']).describe()
```

```
Out[9]: count    21613.000000
mean      13.047817
std       0.526685
min     11.225243
25%    12.682152
50%    13.017003
75%    13.377006
max     15.856731
Name: price, dtype: float64
```

We can see from both the QQ plot and the histogram that the distribution is far closer to normal.

## Building a linear model with transformed variables (15 mts)

Of course, we aren't just restricted to applying the logarithmic transformation to house prices; we can do it to any other variable in our dataset. Let's transform both house prices and price per square foot by this method and interpret the resulting linear model:

```
In [10]: mod1 = smf.ols(formula='np.log(price) ~ np.log(sqft_living)', data=houses).fit()
print(mod1.summary())
```

OLS Regression Results

```
=====
=====
Dep. Variable: np.log(price) R-squared:
0.456
Model: OLS Adj. R-squared:
0.455
Method: Least Squares F-statistic: 1.
808e+04
Date: Fri, 29 May 2020 Prob (F-statistic):
0.00
Time: 16:02:55 Log-Likelihood:
-10240.
No. Observations: 21613 AIC: 2.
048e+04
Df Residuals: 21611 BIC: 2.
050e+04
Df Model: 1

Covariance Type: nonrobust

=====
=====

```

	coef	std err	t	P> t
[0.025 0.975]				
-----	-----	-----	-----	-----
Intercept	6.7299	0.047	143.001	0.000
6.638 6.822				
np.log(sqrt_living)	0.8368	0.006	134.459	0.000
0.825 0.849				
=====	=====	=====	=====	=====
Omnibus: 1.978	123.344	Durbin-Watson:		
Prob(Omnibus): 113.759	0.000	Jarque-Bera (JB):		
Skew: 1.98e-25	0.142	Prob(JB):		
Kurtosis: 2.787	Cond. No.			

### Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

We have to be mindful of how we interpret the coefficients. Although we could say that our results tell us that a 1 unit increase in the logarithm of living space will result in a 0.836 increase in the logarithm of the price, this is a very mechanical and not at all intuitive interpretation.

Mathematics can help us come up with a more intuitive interpretation. Note that the fit our above model has come up with is  $\log \text{price} = 0.84 * \log \text{sqft\_living} + 6.73$ . Exponentiating both sides, we get  $\text{price} = e^{6.73} * \text{sqft\_living}^{0.84}$ . This is a nonlinear relationship, so it's not as straightforward as "increasing sqft living by 1 means that price goes up by X".

However, we can try reframing this in percentage terms; i.e. how does a 1 percent increase in `sqft_living` affect price? We can plug  $sqft\_living_0 = 1.01 * sqft\_living$  into this equation to get

$price_0 = e^{6.73} * 1.01^{0.84} * sqft\_living^{0.84} = 1.01^{0.84} * price \approx 1.0084 * price$ ;  
 i.e. a 1 percent increase in living space results in a 0.84 percent increase in price. This  
 percentage vs. percentage change comparison is known as **elasticity**.

Let's now build a linear model where the logarithmic transform is only applied to the house prices:

```
In [11]: mod2 = smf.ols(formula='np.log(price) ~ sqft_living', data=houses).fit()
()
print(mod2.summary())
```

## OLS Regression Results

Dep. Variable: np.log(price) R-squared:

0.483

Model: OLS Adj. R-squared:

0.483

Method: Least Squares F-statistic: 2.

023e+04

Date: Fri, 29 May 2020 Prob (F-statistic):

0.00

Time: 16:02:57 Log-Likelihood:

-9670.2

No. Observations: 21613 AIC: 1.

934e+04

Df Residuals: 21611 BIC: 1.

936e+04

Df Model: 1

Covariance Type: nonrobust

=====

=====

	coef	std err	t	P> t	[0.025
0.975]					
-----	-----	-----	-----	-----	-----
Intercept	12.2185	0.006	1916.883	0.000	12.206
12.231					
sqft_living	0.0004	2.8e-06	142.233	0.000	0.000
0.000					
-----	-----	-----	-----	-----	-----
Omnibus:		3.128	Durbin-Watson:		
1.979					
Prob(Omnibus):		0.209	Jarque-Bera (JB):		
3.149					
Skew:		0.027	Prob(JB):		
0.207					
Kurtosis:		2.974	Cond. No.		
5.63e+03					
-----	-----	-----	-----	-----	-----
=====	=====	=====	=====	=====	=====

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 5.63e+03. This might indicate that there are strong multicollinearity or other numerical problems.

The interpretation of the regression coefficient is once again different. We interpret the coefficient as a **semi-elasticity**, where an absolute increase in `sqft_living` (because it has not had the logarithm function applied to it) corresponds to a percentage increase `price`. Specifically, here we can say that an increase in living space by 1 square foot leads to a 0.04% percent increase in `price`.

### Exercise 4: (10 mts)

Using the `sns.lmplot()` function, determine which of the above two models is "more linear".

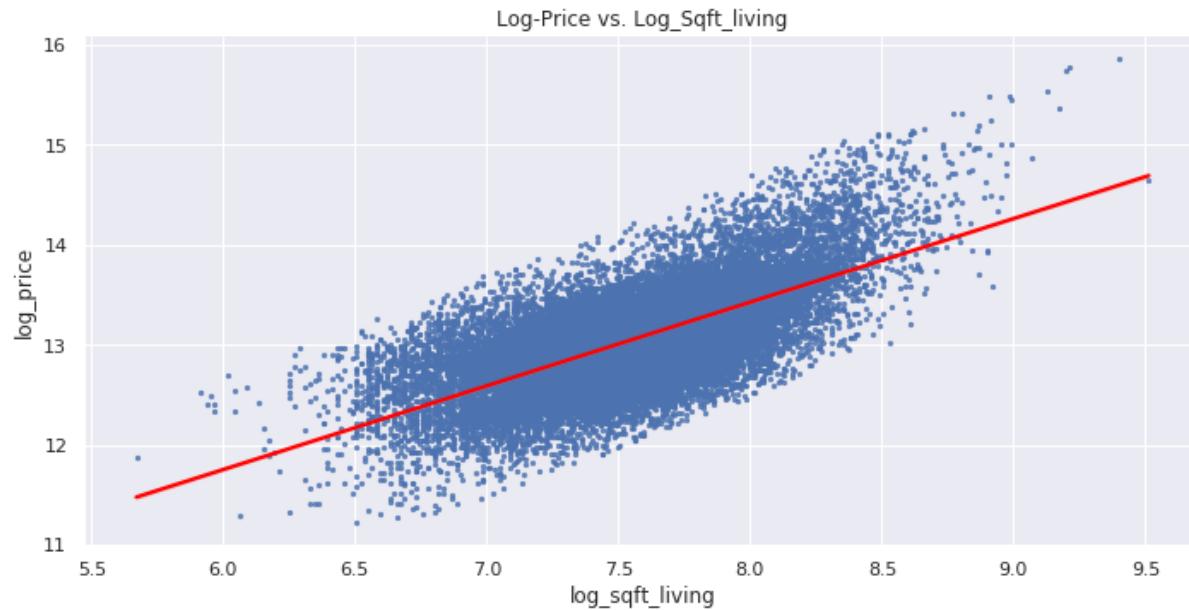
**Answer.** One possible solution is given below:

```
In [12]: houses['log_price'] = np.log(houses['price'])
houses['log_sqft_living'] = np.log(houses['sqft_living'])
```

```
In [13]: ## linear relation between sqft_living and log-price
sns.lmplot(x='sqft_living',y='log_price',data=houses,
            line_kws = {'color': "red"}, aspect= 2, scatter_kws={"s": 5
})
plt.title("Log-Price vs. Sqft_living");
```



```
In [14]: ## linear relation between log-sqft_living and log-price
sns.lmplot(x='log_sqft_living',y='log_price',data=houses,
            line_kws = {'color': "red"}, aspect= 2, scatter_kws={"s": 5})
plt.title("Log-Price vs. Log_Sqft_living");
```



We can see from these plots that the data points of the log-log model cluster more uniformly around the line of best fit across different levels of the predictor variable as compared to the other model, suggesting that the log-log model is more linear.

### Box-Cox transformation (5 mts)

Logarithmic transformations are just one of the possible transformations that we discussed. Earlier, we mentioned powers (e.g. squares, cubes, square roots, etc.) as well as trigonometric functions. In some cases, choosing a transformation can be straightforward (e.g. the logarithm because it is easily interpretable); other times, it is much more difficult. A formal way to decide on which transformation to use is to estimate the coefficient  $\lambda$  of the Box-Cox transformation:

$$BC(\lambda) = \frac{Y^\lambda - 1}{\lambda}$$

If the estimate of  $\lambda$  is close to 2, we can use the quadratic transformation; if it is close to 0.5, the square root transformation; if it is close to zero or less than zero (negative), the logarithmic transformation; etc. In our case, we have:

```
In [15]: price,fitted_lambda = stats.boxcox(houses['price'])  
round(fitted_lambda,2)
```

```
Out[15]: -0.23
```

This is less than zero, so it would seem that using the logarithmic transformation is sensible.

## Multiple linear regression with transformed variables (25 mts)

Of course, as we have seen from the previous case, it doesn't make sense to restrict ourselves to modeling house prices based on only one predictor variable. Let's add in several more variables, some transformed and some not:

### Exercise 5: (5 mts)

Fit a linear model of `log price` vs. `log sqft_living`, `log sqft_lot`, `bedrooms`, `floors`, `bathrooms`, `waterfront`, `condition`, `view`, `grade`, `yr_built`, `lat`, and `long`. Provide interpretations for the coefficients of `log sqft_living` and `waterfront`.

**Answer.** One possible solution is given below:

```
In [16]: mod3 = smf.ols (formula = 'np.log(price) ~ np.log(sqft_living)+ np.log  
(sqft_lot) +bedrooms + floors + bathrooms +waterfront + condition + wat  
erfront + view + grade + yr_built + lat + long ', data = houses).fit()  
print(mod3.summary())
```

OLS Regression Results

```

=====
=====
Dep. Variable: np.log(price) R-squared:
0.762
Model: OLS Adj. R-squared:
0.762
Method: Least Squares F-statistic:
5772.
Date: Fri, 29 May 2020 Prob (F-statistic):
0.00
Time: 16:03:06 Log-Likelihood:
-1284.6
No. Observations: 21613 AIC:
2595.
Df Residuals: 21600 BIC:
2699.
Df Model: 12

Covariance Type: nonrobust

=====
=====
```

		coef	std err	t	P> t
[0.025	0.975]				
-----					
Intercept		-39.9251	2.036	-19.609	0.000
3.916	-35.934				-4
np.log(sqft_living)		0.4096	0.009	46.679	0.000
0.392	0.427				
np.log(sqft_lot)		-0.0076	0.002	-3.070	0.002
0.012	-0.003				-
bedrooms		-0.0256	0.002	-10.265	0.000
0.031	-0.021				-
floors		0.0490	0.004	11.302	0.000
0.040	0.057				
bathrooms		0.0705	0.004	17.458	0.000
0.063	0.078				

waterfront		0.3911	0.022	17.680	0.000
0.348	0.434				
condition		0.0552	0.003	18.832	0.000
0.049	0.061				
view		0.0715	0.003	27.024	0.000
0.066	0.077				
grade		0.1858	0.003	74.252	0.000
0.181	0.191				
yr_built		-0.0038	8.66e-05	-44.115	0.000
0.004	-0.004				-
lat		1.3443	0.013	100.583	0.000
1.318	1.371				
long		0.0673	0.015	4.446	0.000
0.038	0.097				

---

Omnibus:	568.294	Durbin-Watson:	
1.981			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	1
119.493			
Skew:	0.182	Prob(JB):	8.
04e-244			
Kurtosis:	4.054	Cond. No.	
2.30e+06			

---

=====

#### Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.3e+06. This might indicate that there are strong multicollinearity or other numerical problems.

All the variables are statistically significant (all  $p$  - values less than 0.01). Overall this linear model explains over 76 percent of the total variability of the response variable.

An increase of one percent in living space leads to an increase of 0.4096 percent in price. A property with a water view has an increase in price of 39.11 percent.

What other factors may impact the price that we have left out? Some that may play a role include proximity to services (hospitals, schools, commercial areas, movie theaters, metro stops...), crime rates, etc. Our dataset does not have a comprehensive list of possible factors; however, we do have some variables that would be interesting to investigate further.

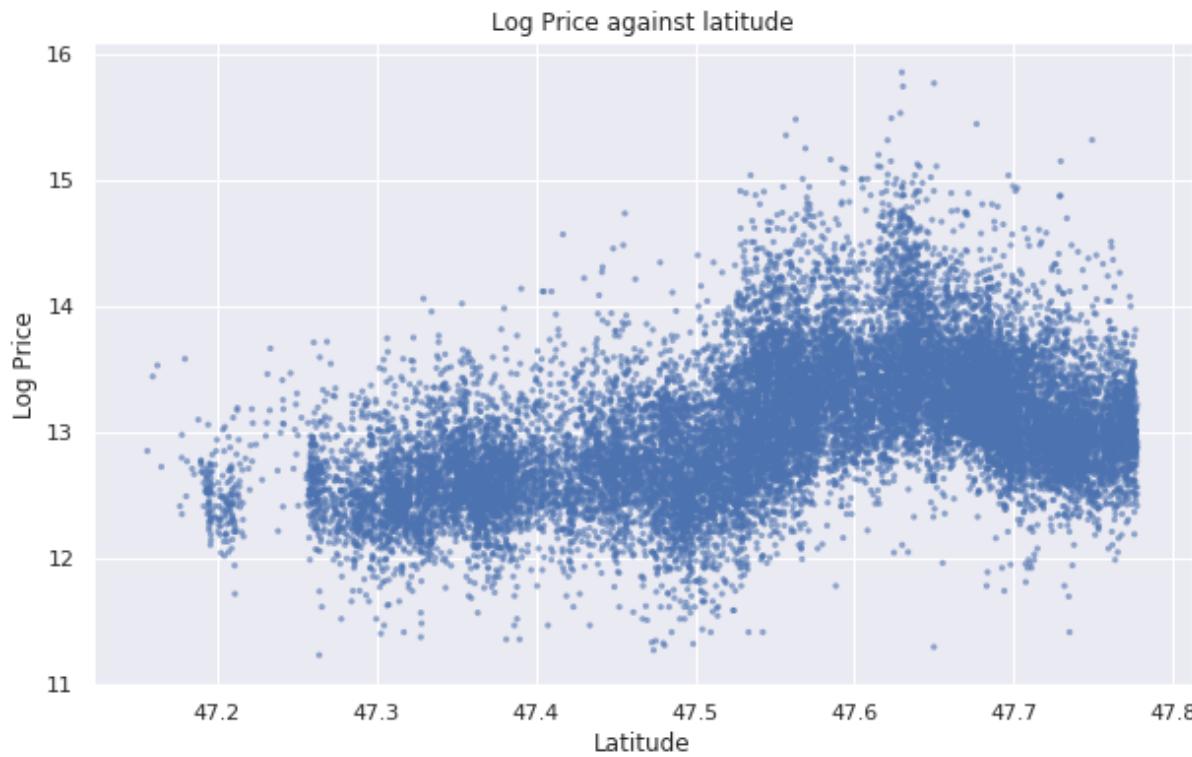
In general, house prices change depending on the location. Two houses with comparable features can be priced very differently depending on the neighborhood and geographic position. In this dataset, we have zipcode and geographic coordinates. Let us start by taking a look at the relationship between latitude and prices.

### Exercise 6: (5 mts)

Examine the relationship between latitude and the logarithm of house prices. What do you observe?

**Answer.** One possible solution is shown below:

```
In [17]: fig = plt.figure(figsize = (10,6))
ax = fig.add_subplot(111)
ax.scatter(houses['lat'],np.log(houses['price']),c ='b',alpha=0.6,edgecolors='none',s=10)
plt.xlabel("Latitude")
plt.ylabel("Log Price")
ax.set_title("Log Price against latitude")
plt.show();
```



We can see that there is a nonlinear relationship between latitude and price. Based on the concave curvature in the right half of the plot above, it seems that adding a quadratic term would be able to help us explain this.

### Exercise 7: (10 mts)

Add the square of the latitude as an additional predictor to the model in Exercise 5. Is the term significant? Use the AIC score (described below) to evaluate whether or not the fit has improved.

**Background on AIC Score:** One of the drawbacks of  $R^2$  is that it can never decrease when the set of predictors is increased. In other words, there is no penalty for continuing to add variables that have little explanatory power. Consequently, selecting predictor variables trying to maximize

$R^2$  can lead to choosing unnecessarily complex and redundant models. How do we choose a model that offers a good quality of fit while minimizing the number of features?

There are several model selection criteria that quantify the quality of a model by managing the tradeoff between goodness-of-fit and simplicity. The most common one is the AIC (Akaike Information Criterion). The AIC penalizes the addition of more terms to a model, so in order for an updated model to have a better AIC, its  $R^2$  needs to improve at least as much as the additional imposed penalty. The smaller a model's AIC, the higher its quality.

For now, do not worry about the technical details behind AIC (although you are free to look them up yourself). In future cases on regularization, you will learn more about the rationale behind why these sorts of estimators matter and how to construct and use them in model-building.

**Answer.** One possible solution is shown below:

In [18]:

```
## lat square effect
mod4 = smf.ols ( formula = 'np.log(price) ~ np.log(sqft_living)+ np.log
(sqft_lot) +bedrooms + floors + bathrooms +waterfront + condition + wat
erfront + view + grade + yr_built + lat + I(lat**2) + long ', data = ho
uses ).fit()
print(mod4.summary())
```

### OLS Regression Results

```
=====
=====
Dep. Variable: np.log(price) R-squared: 0.778
Model: OLS Adj. R-squared: 0.778
Method: Least Squares F-statistic: 5816.
Date: Fri, 29 May 2020 Prob (F-statistic): 0.00
Time: 16:03:11 Log-Likelihood: -554.41
No. Observations: 21613 AIC: 1127
```

113/.						
Df Residuals:		21599	BIC:			
1249.						
Df Model:		13				
Covariance Type:		nonrobust				
<hr/>						
		coef	std err	t	P> t	
[0.025	0.975]					
<hr/>						
Intercept		-7773.2083	199.024	-39.057	0.000	-816
3.311	-7383.106					
np.log(sqft_living)	0.392	0.4084	0.008	48.137	0.000	
0.425						
np.log(sqft_lot)	0.007	0.0122	0.002	4.968	0.000	
0.017						
bedrooms		-0.0253	0.002	-10.498	0.000	-
0.030	-0.021					
floors		0.0488	0.004	11.643	0.000	
0.041	0.057					
bathrooms		0.0660	0.004	16.913	0.000	
0.058	0.074					
waterfront		0.3755	0.021	17.555	0.000	
0.334	0.417					
condition		0.0598	0.003	21.083	0.000	
0.054	0.065					
view		0.0687	0.003	26.853	0.000	
0.064	0.074					
grade		0.1763	0.002	72.508	0.000	
0.172	0.181					
yr_built		-0.0032	8.53e-05	-37.186	0.000	-
0.003	-0.003					
lat		326.2171	8.361	39.019	0.000	30
9.830	342.604					
I(lat ** 2)		-3.4174	0.088	-38.858	0.000	-
3.590	-3.245					
long		-0.0232	0.015	-1.568	0.117	-

```
0.052      0.006
=====
=====
```

Omnibus:	649.112	Durbin-Watson:
1.984		
Prob(Omnibus):	0.000	Jarque-Bera (JB):
420.202		1
Skew:	0.175	Prob(JB):
05e-309		4.
Kurtosis:	4.206	Cond. No.
3.54e+08		

```
=====
=====
```

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 3.54e+08. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [19]: print(mod3.aic)
```

```
2595.2534342787694
```

```
In [20]: print(mod4.aic)
```

```
1136.8143169405448
```

We can see that the coefficient is highly significant. The R-squared has increased by about 1%. Comparing model 3 and model 4 we can see that the AIC has improved from 2597.253 to 1138.814.

Let's now add the zipcode to our model:

```
In [21]: ## location effect with dummy variables
```

```
mod5 = smf.ols ( formula = 'np.log(price) ~ np.log(sqft_living)+ np.log(sqft_lot) +bedrooms + floors + bathrooms +waterfront + condition + waterfront + view + grade + yr_built + lat + I(lat**2) + long + C(zipcode)', data = houses ).fit()
print(mod5.summary())
```

### OLS Regression Results

```
=====
=====
Dep. Variable: np.log(price) R-squared: 0.879
Model: OLS Adj. R-squared: 0.879
Method: Least Squares F-statistic: 1910.
Date: Fri, 29 May 2020 Prob (F-statistic): 0.00
Time: 16:03:15 Log-Likelihood: -6026.8
No. Observations: 21613 AIC: -1.189e+04
Df Residuals: 21530 BIC: -1.123e+04
Df Model: 82
Covariance Type: nonrobust
=====
=====
```

	coef	std err	t	P> t
[0.025 0.975]				
-----	-----	-----	-----	-----
Intercept	-5882.6548	526.169	-11.180	0.000
3.984 -4851.325				-691
C(zipcode)[T.98002]	0.0200	0.017	1.208	0.227
0.012 0.052				-
C(zipcode)[T.98003]	-0.0119	0.015	-0.804	0.421
0.041 0.017				-

C(zipcode)[T.98004]	0.9101	0.028	32.355	0.000
0.855 0.965				
C(zipcode)[T.98005]	0.5146	0.030	17.169	0.000
0.456 0.573				
C(zipcode)[T.98006]	0.4547	0.026	17.704	0.000
0.404 0.505				
C(zipcode)[T.98007]	0.4419	0.031	14.293	0.000
0.381 0.503				
C(zipcode)[T.98008]	0.4523	0.030	15.326	0.000
0.394 0.510				
C(zipcode)[T.98010]	0.3120	0.025	12.308	0.000
0.262 0.362				
C(zipcode)[T.98011]	0.2602	0.037	7.108	0.000
0.188 0.332				
C(zipcode)[T.98014]	0.1927	0.041	4.747	0.000
0.113 0.272				
C(zipcode)[T.98019]	0.2136	0.040	5.393	0.000
0.136 0.291				
C(zipcode)[T.98022]	0.3399	0.024	13.880	0.000
0.292 0.388				
C(zipcode)[T.98023]	-0.0631	0.014	-4.648	0.000
0.090 -0.036				-
C(zipcode)[T.98024]	0.3140	0.037	8.503	0.000
0.242 0.386				
C(zipcode)[T.98027]	0.3753	0.026	14.191	0.000
0.323 0.427				
C(zipcode)[T.98028]	0.2015	0.036	5.667	0.000
0.132 0.271				
C(zipcode)[T.98029]	0.4724	0.030	15.928	0.000
0.414 0.530				
C(zipcode)[T.98030]	0.0020	0.017	0.118	0.906
0.032 0.036				-
C(zipcode)[T.98031]	-0.0240	0.019	-1.292	0.196
0.060 0.012				-
C(zipcode)[T.98032]	-0.1287	0.020	-6.313	0.000
0.169 -0.089				-
C(zipcode)[T.98033]	0.5712	0.031	18.518	0.000
0.511 0.632				
C(zipcode)[T.98034]	0.3229	0.033	9.881	0.000

0.259	0.387				
C(zipcode)[T.98038]		0.1920	0.019	9.995	0.000
0.154	0.230				
C(zipcode)[T.98039]		1.0861	0.037	29.279	0.000
1.013	1.159				
C(zipcode)[T.98040]		0.6648	0.026	25.771	0.000
0.614	0.715				
C(zipcode)[T.98042]		0.0526	0.016	3.199	0.001
0.020	0.085				
C(zipcode)[T.98045]		0.3142	0.036	8.834	0.000
0.245	0.384				
C(zipcode)[T.98052]		0.4500	0.031	14.297	0.000
0.388	0.512				
C(zipcode)[T.98053]		0.4481	0.034	13.292	0.000
0.382	0.514				
C(zipcode)[T.98055]		-0.0070	0.021	-0.328	0.743
0.049	0.035				-
C(zipcode)[T.98056]		0.1444	0.023	6.239	0.000
0.099	0.190				
C(zipcode)[T.98058]		0.0387	0.020	1.912	0.056
0.001	0.078				-
C(zipcode)[T.98059]		0.2042	0.023	8.974	0.000
0.160	0.249				
C(zipcode)[T.98065]		0.3722	0.033	11.197	0.000
0.307	0.437				
C(zipcode)[T.98070]		0.0356	0.025	1.443	0.149
0.013	0.084				-
C(zipcode)[T.98072]		0.3015	0.036	8.287	0.000
0.230	0.373				
C(zipcode)[T.98074]		0.3943	0.031	12.847	0.000
0.334	0.454				
C(zipcode)[T.98075]		0.4312	0.030	14.342	0.000
0.372	0.490				
C(zipcode)[T.98077]		0.2904	0.038	7.676	0.000
0.216	0.365				
C(zipcode)[T.98092]		0.0821	0.015	5.565	0.000
0.053	0.111				
C(zipcode)[T.98102]		0.6994	0.032	21.700	0.000
0.636	0.763				

C(zipcode)[T.98103]	0.5426	0.030	18.208	0.000
0.484 0.601				
C(zipcode)[T.98105]	0.6933	0.031	22.540	0.000
0.633 0.754				
C(zipcode)[T.98106]	0.0775	0.024	3.215	0.001
0.030 0.125				
C(zipcode)[T.98107]	0.5516	0.031	17.928	0.000
0.491 0.612				
C(zipcode)[T.98108]	0.1021	0.026	3.900	0.000
0.051 0.153				
C(zipcode)[T.98109]	0.7096	0.032	22.107	0.000
0.647 0.772				
C(zipcode)[T.98112]	0.7930	0.029	27.537	0.000
0.737 0.849				
C(zipcode)[T.98115]	0.5516	0.030	18.263	0.000
0.492 0.611				
C(zipcode)[T.98116]	0.4565	0.026	17.471	0.000
0.405 0.508				
C(zipcode)[T.98117]	0.5155	0.031	16.889	0.000
0.456 0.575				
C(zipcode)[T.98118]	0.2292	0.024	9.680	0.000
0.183 0.276				
C(zipcode)[T.98119]	0.6845	0.030	22.566	0.000
0.625 0.744				
C(zipcode)[T.98122]	0.5481	0.028	19.758	0.000
0.494 0.602				
C(zipcode)[T.98125]	0.3041	0.032	9.399	0.000
0.241 0.367				
C(zipcode)[T.98126]	0.2688	0.025	10.968	0.000
0.221 0.317				
C(zipcode)[T.98133]	0.1944	0.033	5.827	0.000
0.129 0.260				
C(zipcode)[T.98136]	0.3992	0.025	15.937	0.000
0.350 0.448				
C(zipcode)[T.98144]	0.4202	0.026	15.981	0.000
0.369 0.472				
C(zipcode)[T.98146]	0.0226	0.023	0.985	0.325
0.022 0.067				-
C(zipcode)[T.98148]	-0.0319	0.029	-1.090	0.276
				-

0.089	0.025					
C(zipcode)[T.98155]	0.1852	0.035	5.322	0.000		
0.117	0.253					
C(zipcode)[T.98166]	0.0839	0.021	4.013	0.000		
0.043	0.125					
C(zipcode)[T.98168]	-0.1574	0.022	-7.066	0.000	-	
0.201	-0.114					
C(zipcode)[T.98177]	0.3235	0.035	9.275	0.000		
0.255	0.392					
C(zipcode)[T.98178]	-0.0625	0.023	-2.716	0.007	-	
0.108	-0.017					
C(zipcode)[T.98188]	-0.0918	0.023	-3.999	0.000	-	
0.137	-0.047					
C(zipcode)[T.98198]	-0.0715	0.017	-4.166	0.000	-	
0.105	-0.038					
C(zipcode)[T.98199]	0.5506	0.029	18.691	0.000		
0.493	0.608					
np.log(sqft_living)	0.4229	0.006	66.616	0.000		
0.410	0.435					
np.log(sqft_lot)	0.0703	0.002	33.767	0.000		
0.066	0.074					
bedrooms	-0.0148	0.002	-8.202	0.000	-	
0.018	-0.011					
floors	0.0194	0.003	5.966	0.000		
0.013	0.026					
bathrooms	0.0383	0.003	13.145	0.000		
0.033	0.044					
waterfront	0.4698	0.016	29.193	0.000		
0.438	0.501					
condition	0.0418	0.002	19.413	0.000		
0.038	0.046					
view	0.0640	0.002	32.841	0.000		
0.060	0.068					
grade	0.1117	0.002	58.285	0.000		
0.108	0.115					
yr_built	-0.0006	7.15e-05	-7.745	0.000	-	
0.001	-0.000					
lat	245.3030	22.157	11.071	0.000	20	
1.873	288.733					

```

I(lat ** 2)           -2.5750      0.233     -11.045      0.000      -
3.032      -2.118
long                  -0.4050      0.052     -7.749      0.000      -
0.507      -0.303
=====
=====

Omnibus:                1422.131   Durbin-Watson:
2.001
Prob(Omnibus):          0.000    Jarque-Bera (JB):      5
908.533
Skew:                   -0.193    Prob(JB):
0.00
Kurtosis:                5.532    Cond. No.
1.27e+09
=====
=====

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is
correctly specified.
[2] The condition number is large, 1.27e+09. This might indicate that t
here are
strong multicollinearity or other numerical problems.

```

From the output, we can see how many of the different ZIP codes exert a significant effect. The R-squared of the model has increased dramatically to 87.92%.

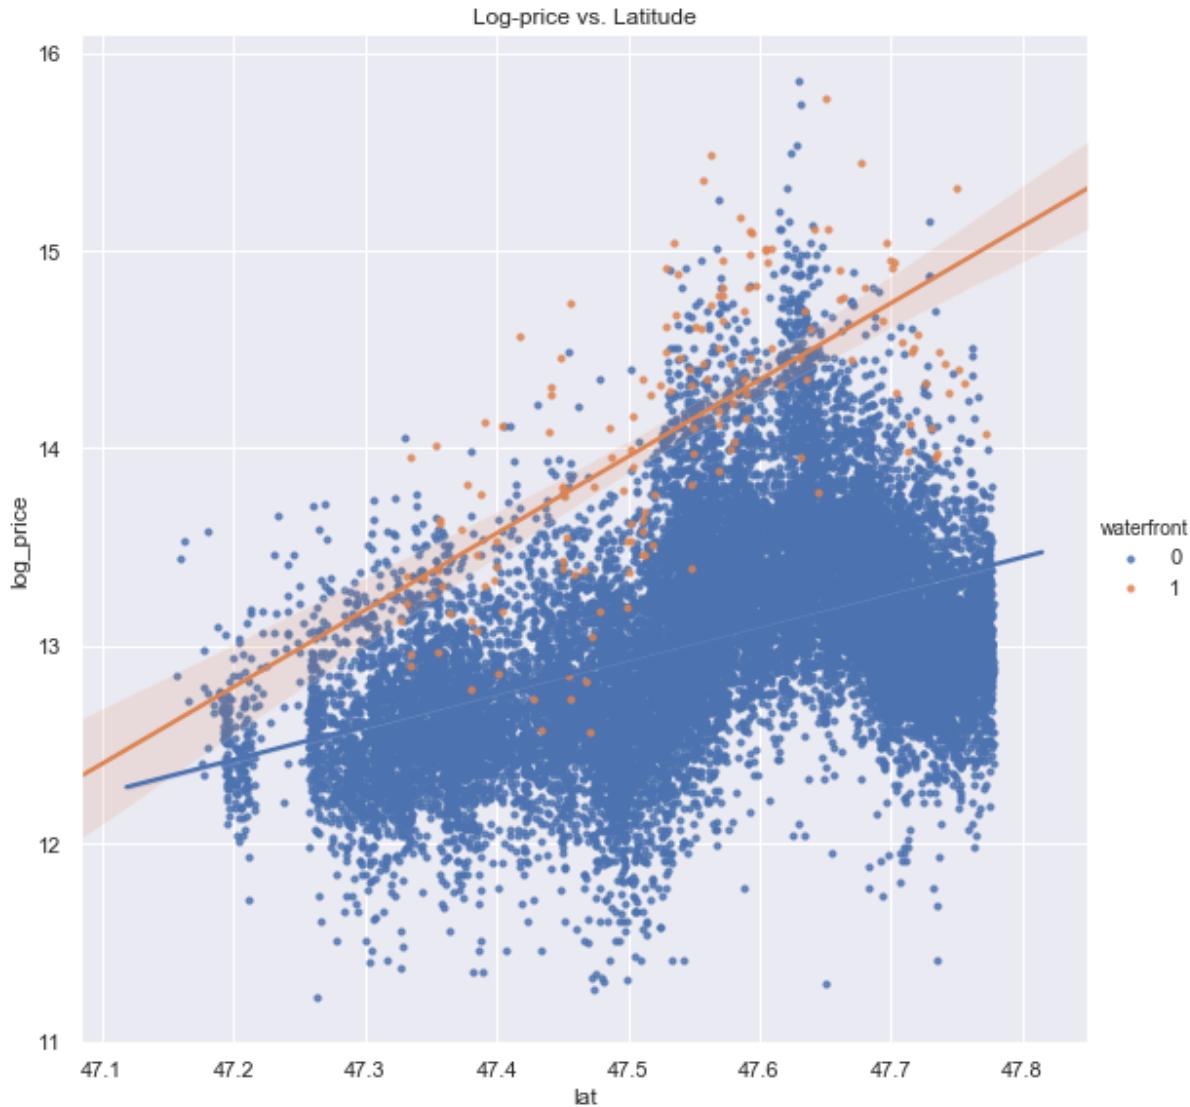
## Modeling interaction effects (50 mts)

As we have seen during the EDA cases, interaction effects can complicate the perceived effect of the predictor variables on the outcome of interest. Let's dig into potential interactions by looking at three of the predictors in tandem: `waterfront`, geographic position (`lat` and `long`), and `sqft_living`. Specifically, is the effect of geographic position and `sqft_living` different among the houses that have a waterfront view vs. those that do not?

To study this interaction effect, let's fit two separate regression lines for each subgroup (here, the two subgroups are having a waterfront or not):

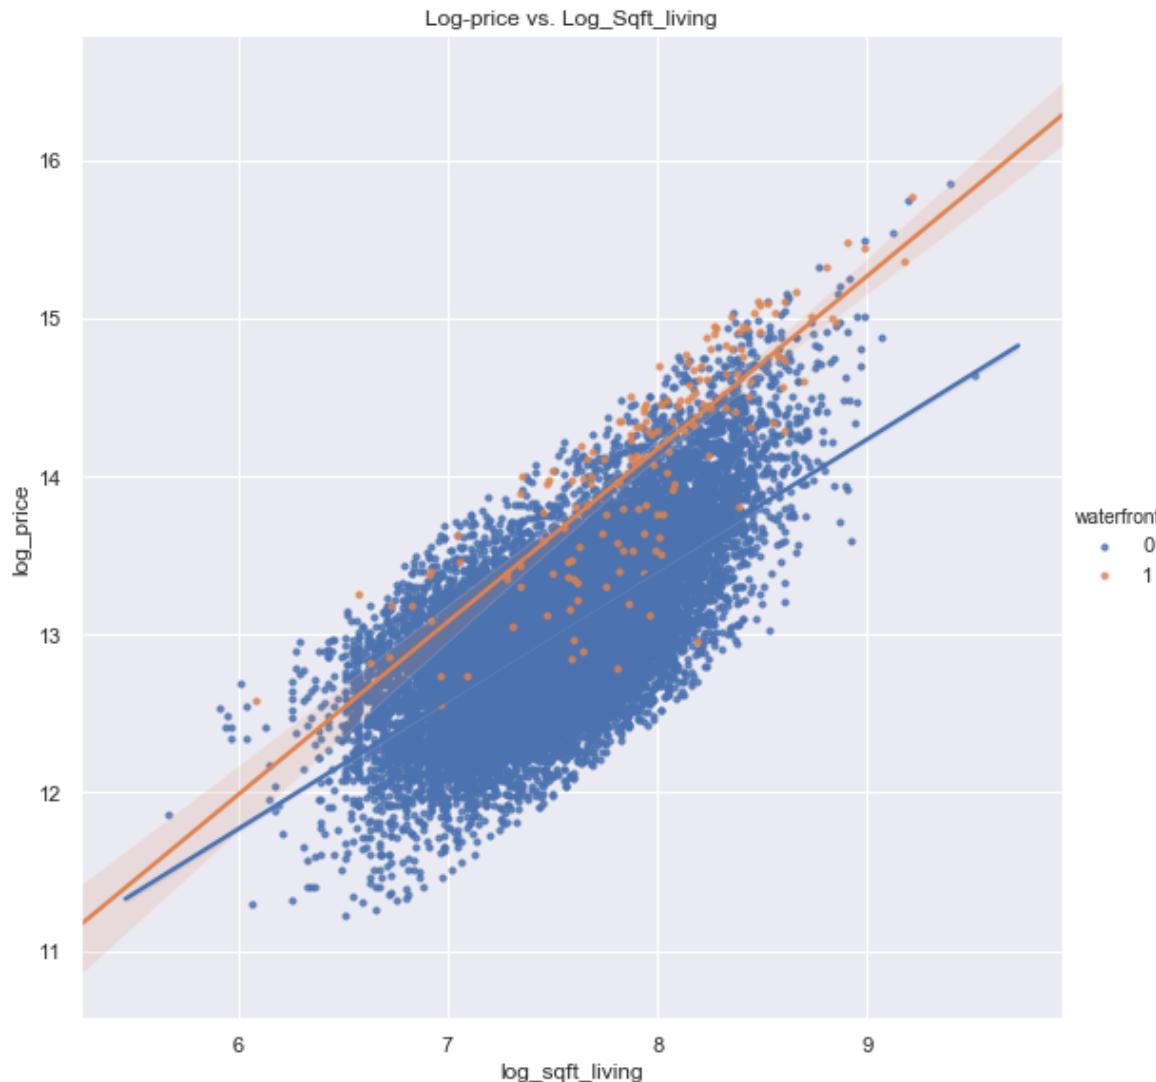
```
In [61]: sns.lmplot(x = 'lat', y= 'log_price', data = houses,
                  hue="waterfront", height=8, scatter_kws={"s": 10})
plt.title("Log-price vs. Latitude")
```

Out[61]: Text(0.5, 1, 'Log-price vs. Latitude')



```
In [62]: sns.lmplot(x='log_sqft_living', y='log_price', data=houses,
                  hue="waterfront", height=8, scatter_kws={"s": 10})
plt.title("Log-price vs. Log_Sqft_living")
```

```
Out[62]: Text(0.5, 1, 'Log-price vs. Log_Sqft_living')
```



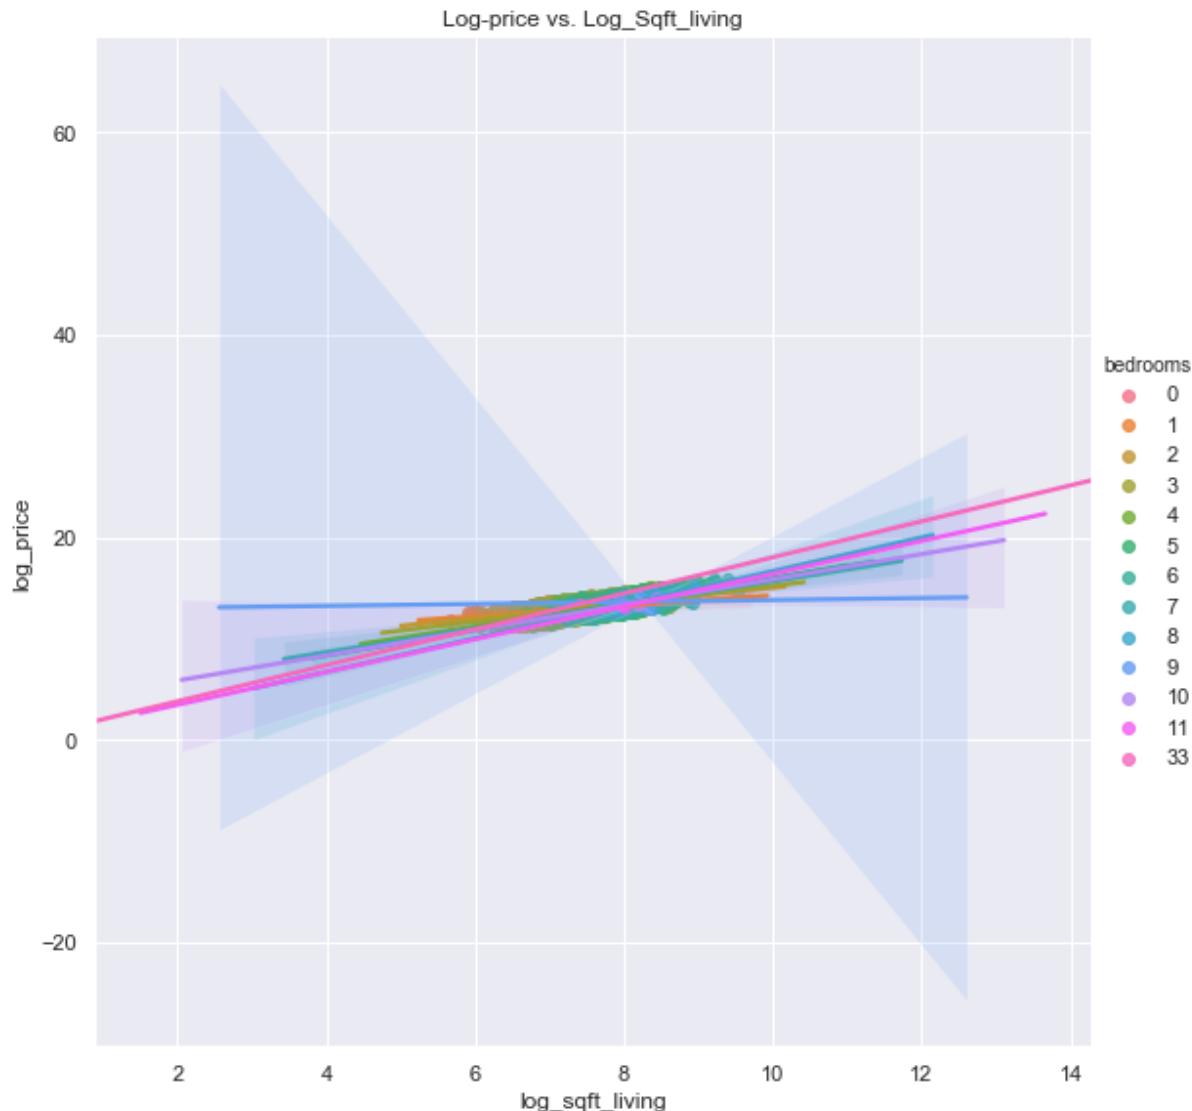
We see that the effects of `latitude` and the logarithm of `sqft_living` are more pronounced for the subgroup of houses with a waterfront view.

## Exercise 8: (5 mts)

Using the `lmplot()` function in `seaborn`, create a regression plot of the logarithm of `price` vs. the logarithm of `sqft_living`, interacting on the number of bedrooms. What patterns do you observe?

**Answer.** One possible solution is shown below:

```
In [63]: sns.lmplot(x = 'log_sqft_living', y= 'log_price', data = houses, hue="bedrooms", height=8)
plt.title("Log-price vs. Log_Sqft_living");
```



There is not a clear monotonic pattern in regression line slope as the number of bedrooms increases, so we will say that the number of bedrooms does *not* interact with the relationship between price and square footage.

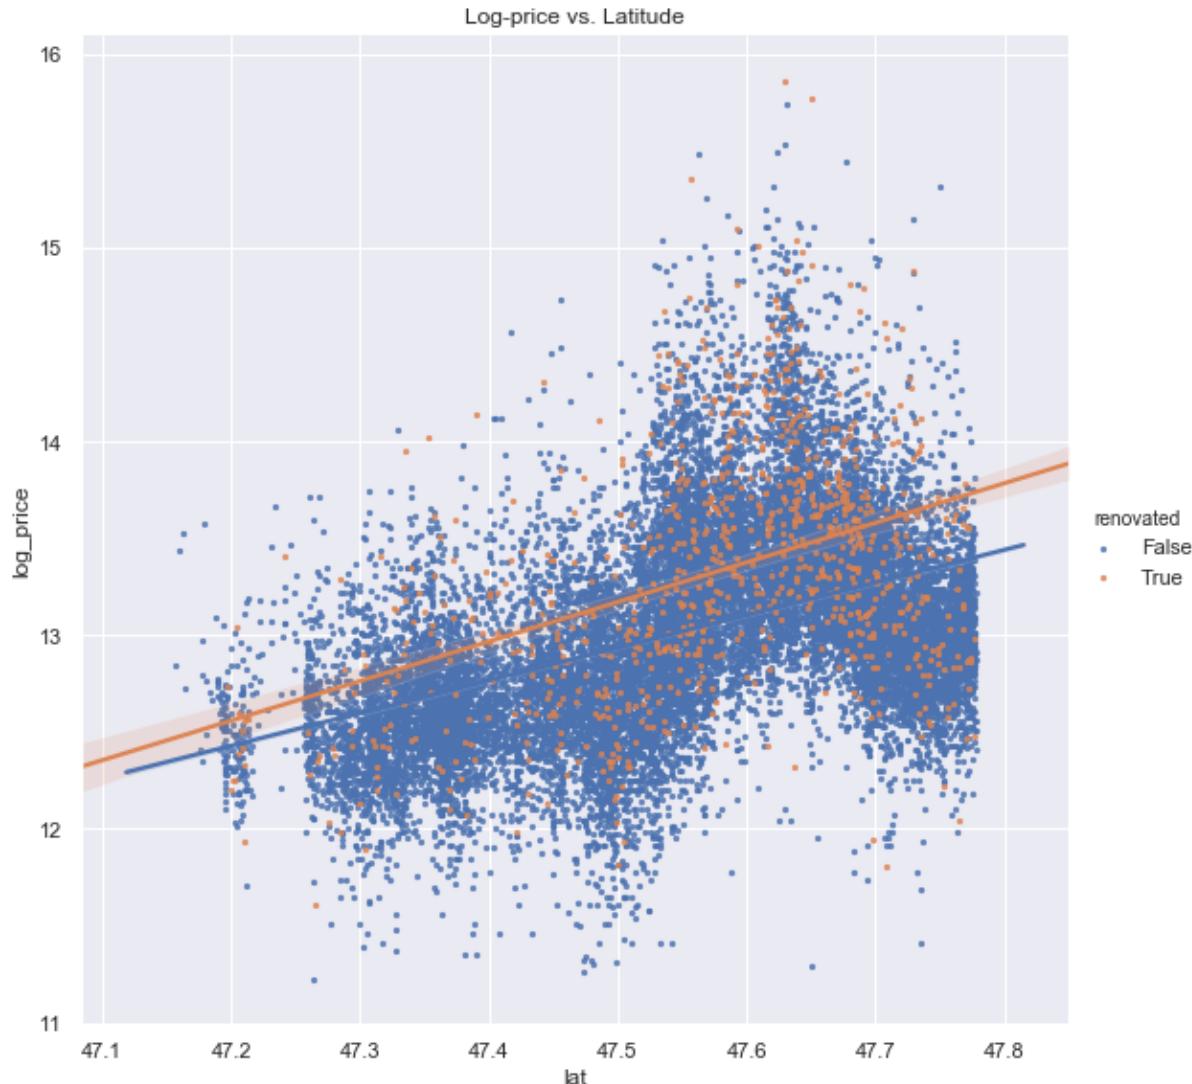
### Exercise 9: (10 mts)

Create a regression plot to see if a house's renovation status interacts with the relationship between the logarithm of `price` and the logarithm of `sqft_living`. What can you conclude?

**Answer.** One possible solution is shown below:

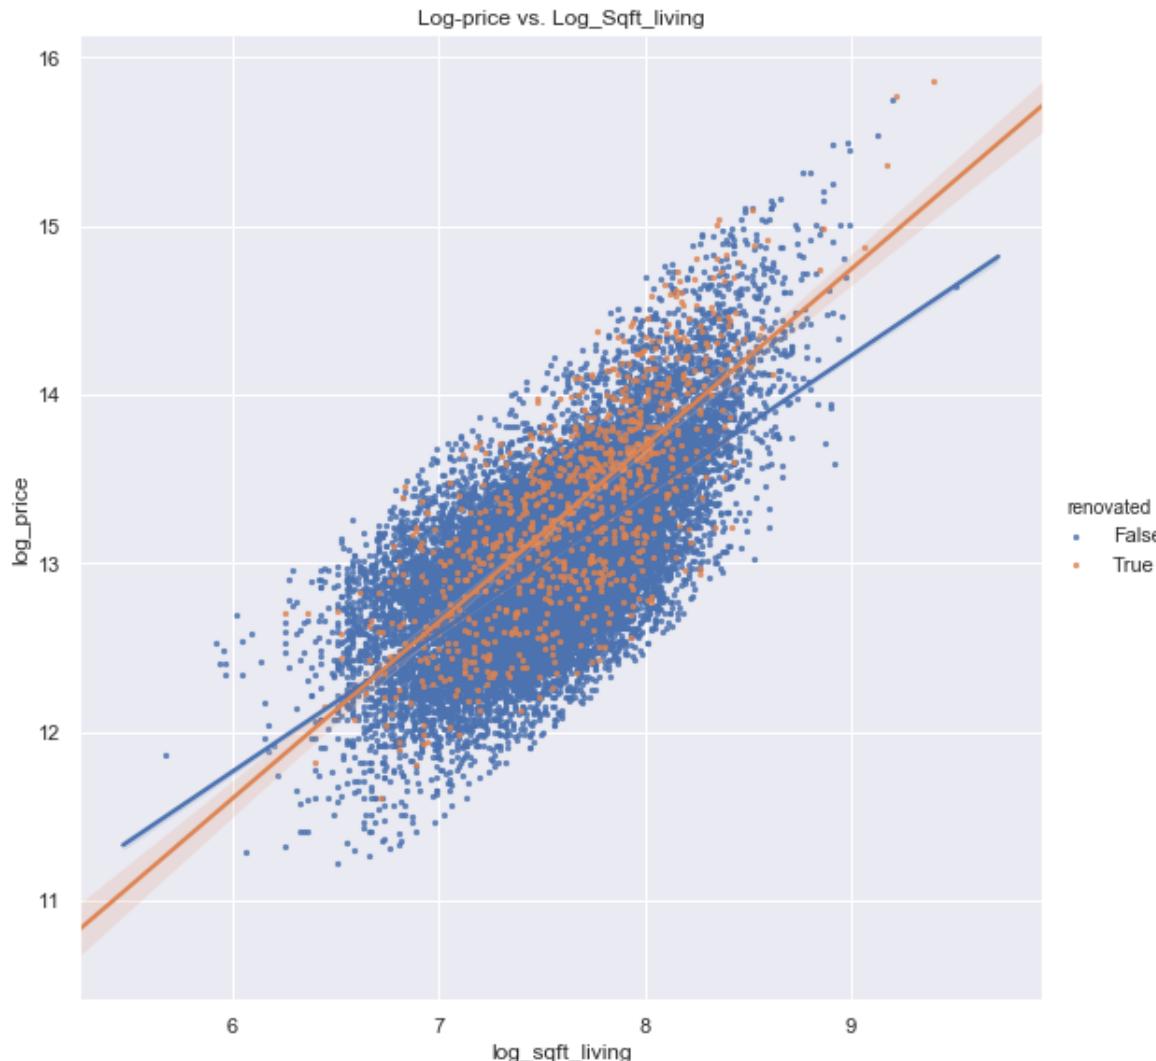
```
In [64]: houses['renovated'] = houses['yr_renovated'] >0
```

```
In [65]: sns.lmplot(x = 'lat', y= 'log_price', data = houses, hue="renovated", height=8, scatter_kws={"s": 5});  
plt.title("Log-price vs. Latitude");
```



```
In [66]: sns.lmplot(x = 'log_sqft_living', y= 'log_price', data = houses,
                  hue="renovated", height=8, scatter_kws={"s": 5})
plt.title("Log-price vs. Log_Sqft_living")
```

```
Out[66]: Text(0.5, 1, 'Log-price vs. Log_Sqft_living')
```



We can see that the relationship between the predictor and outcome variables is more pronounced for houses that have been renovated.

We verify this with the following statistical models:

In [67]: *fito os meus*

```
formula = ('np.log(price) ~ lat*C(waterfront)')
mod5_1 = smf.ols(formula=formula, data=houses).fit()
print(mod5_1.summary())
```

### OLS Regression Results

Dep. Variable:	np.log(price)	R-squared:		
0.236				
Model:	OLS	Adj. R-squared:		
0.236				
Method:	Least Squares	F-statistic:		
2228.				
Date:	Thu, 14 Nov 2019	Prob (F-statistic):		
0.00				
Time:	01:50:02	Log-Likelihood:		
-13898.				
No. Observations:	21613	AIC:		
780e+04		2.		
Df Residuals:	21609	BIC:		
784e+04		2.		
Df Model:	3			
Covariance Type:	nonrobust			
	coef	std err	t	P> t
[0.025 0.975]				
-----	-----	-----	-----	-----
Intercept	-68.0880	1.078	-63.180	0.000
-70.200 -65.976				
C(waterfront)[T.1]	-102.3040	14.909	-6.862	0.000
131.526 -73.082				-
lat	1.7058	0.023	75.280	0.000
1.661 1.750				
lat:C(waterfront)[T.1]	2.1753	0.314	6.936	0.000
1.561 2.790				

```
=====
=====
Omnibus:                      1283.881   Durbin-Watson:
    1.957
Prob(Omnibus):                 0.000   Jarque-Bera (JB):      1
    927.815
Skew:                           0.510   Prob(JB):
    0.00
Kurtosis:                      4.049   Cond. No.
    2.27e+05
=====
=====
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.  
[2] The condition number is large, 2.27e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [68]: formula = ('np.log(price) ~ np.log(sqft_living)*C(waterfront)')
mod5_2 = smf.ols(formula=formula, data=houses).fit()
print(mod5_1.summary())
```

### OLS Regression Results

```
=====
=====
Dep. Variable:          np.log(price)   R-squared:
    0.236
Model:                  OLS            Adj. R-squared:
    0.236
Method:                 Least Squares   F-statistic:
    2228.
Date:                   Thu, 14 Nov 2019   Prob (F-statistic):
    0.00
Time:                   01:50:02        Log-Likelihood:
    -13898.
No. Observations:      21613         AIC:                  2.
```

780e+04  
Df Residuals: 21609 BIC: 2.  
784e+04  
Df Model: 3  
  
Covariance Type: nonrobust

		coef	std err	t	P> t
	[0.025 0.975]				
Intercept		-68.0880	1.078	-63.180	0.000
-70.200	-65.976				
C(waterfront)[T.1]		-102.3040	14.909	-6.862	0.000
131.526	-73.082				
lat		1.7058	0.023	75.280	0.000
1.661	1.750				
lat:C(waterfront)[T.1]		2.1753	0.314	6.936	0.000
1.561	2.790				

Omnibus:	1283.881	Durbin-Watson:
1.957		
Prob(Omnibus):	0.000	Jarque-Bera (JB):
927.815		1
Skew:	0.510	Prob(JB):
0.00		
Kurtosis:	4.049	Cond. No.
2.27e+05		

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.27e+05. This might indicate that t

here are strong multicollinearity or other numerical problems.

For the latitude model (mod5\_1) the difference between the two intercepts is -102.303 and the difference between the two slopes is 2.175. They are both statistically significant, as we can conclude from the very small p-values. For the log of square feet model (mod5\_2) the difference between the two intercepts is -1.415 and the difference between the two slopes is 0.272. Also in this case they are both statistically significant, as we can conclude from the very small p-values.

We can conclude that there is a significant interaction effect between the effect of geographic position and house size when compared for properties with and without a waterfront view. In particular the price of houses with a waterfront increase faster with size and latitude.

## Incorporating interaction effects into a linear model (15 mts)

Of course, the above methodology is very inefficient for two reasons:

1. It can only incorporate one interaction effect at a time.
2. It requires fitting multiple linear regression models, depending on the value(s) of the interacting term.

Let's start with our base model which includes all of the other variables we have discussed before, along with separate fixed effects for `waterfront` and `renovated` (but no interaction):

```
In [69]: formula = ('np.log(price) ~ np.log(sqft_living)+ np.log(sqft_lot) + bed  
rooms + floors + bathrooms '  
    '+ waterfront + condition + C(waterfront) + view + grade + y  
r_built + lat + I(lat**2) '  
    '+ long + C(zipcode)+ C(renovated))  
mod6 = smf.ols(formula=formula, data=houses).fit()  
print(mod6.summary())
```

OLS Regression Results

```

=====
Dep. Variable: np.log(price) R-squared:
0.880
Model: OLS Adj. R-squared:
0.879
Method: Least Squares F-statistic:
1894.
Date: Thu, 14 Nov 2019 Prob (F-statistic):
0.00
Time: 01:50:04 Log-Likelihood:
6064.5
No. Observations: 21613 AIC: -1.
196e+04
Df Residuals: 21529 BIC: -1.
129e+04
Df Model: 83

Covariance Type: nonrobust
=====
```

	coef	std err	t	P> t
[0.025 0.975]				
-----	-----	-----	-----	-----
Intercept	-6000.5172	525.440	-11.420	0.000
0.419 -4970.615				-703
C(waterfront)[T.1]	0.2312	0.008	28.748	0.000
0.215 0.247				
C(zipcode)[T.98002]	0.0206	0.017	1.244	0.213
0.012 0.053				-
C(zipcode)[T.98003]	-0.0109	0.015	-0.738	0.460
0.040 0.018				-
C(zipcode)[T.98004]	0.9089	0.028	32.366	0.000
0.854 0.964				
C(zipcode)[T.98005]	0.5169	0.030	17.277	0.000
0.458 0.576				
C(zipcode)[T.98006]	0.4552	0.026	17.754	0.000

0.405	0.505				
C(zipcode)[T.98007]		0.4437	0.031	14.375	0.000
0.383	0.504				
C(zipcode)[T.98008]		0.4550	0.029	15.443	0.000
0.397	0.513				
C(zipcode)[T.98010]		0.3080	0.025	12.169	0.000
0.258	0.358				
C(zipcode)[T.98011]		0.2647	0.037	7.243	0.000
0.193	0.336				
C(zipcode)[T.98014]		0.1971	0.041	4.863	0.000
0.118	0.277				
C(zipcode)[T.98019]		0.2179	0.040	5.512	0.000
0.140	0.295				
C(zipcode)[T.98022]		0.3415	0.024	13.971	0.000
0.294	0.389				
C(zipcode)[T.98023]		-0.0630	0.014	-4.649	0.000
0.090	-0.036				-
C(zipcode)[T.98024]		0.3142	0.037	8.522	0.000
0.242	0.386				
C(zipcode)[T.98027]		0.3769	0.026	14.277	0.000
0.325	0.429				
C(zipcode)[T.98028]		0.2064	0.036	5.812	0.000
0.137	0.276				
C(zipcode)[T.98029]		0.4756	0.030	16.065	0.000
0.418	0.534				
C(zipcode)[T.98030]		0.0029	0.017	0.169	0.866
0.031	0.036				-
C(zipcode)[T.98031]		-0.0236	0.019	-1.276	0.202
0.060	0.013				-
C(zipcode)[T.98032]		-0.1286	0.020	-6.320	0.000
0.169	-0.089				-
C(zipcode)[T.98033]		0.5725	0.031	18.593	0.000
0.512	0.633				
C(zipcode)[T.98034]		0.3271	0.033	10.024	0.000
0.263	0.391				
C(zipcode)[T.98038]		0.1926	0.019	10.044	0.000
0.155	0.230				
C(zipcode)[T.98039]		1.0816	0.037	29.207	0.000
1.009	1.154				

C(zipcode)[T.98040]	0.6621	0.026	25.708	0.000
0.612 0.713				
C(zipcode)[T.98042]	0.0514	0.016	3.132	0.002
0.019 0.084				
C(zipcode)[T.98045]	0.3168	0.036	8.921	0.000
0.247 0.386				
C(zipcode)[T.98052]	0.4534	0.031	14.432	0.000
0.392 0.515				
C(zipcode)[T.98053]	0.4513	0.034	13.410	0.000
0.385 0.517				
C(zipcode)[T.98055]	-0.0055	0.021	-0.259	0.796 -
0.047 0.036				
C(zipcode)[T.98056]	0.1439	0.023	6.231	0.000
0.099 0.189				
C(zipcode)[T.98058]	0.0378	0.020	1.871	0.061 -
0.002 0.077				
C(zipcode)[T.98059]	0.2037	0.023	8.968	0.000
0.159 0.248				
C(zipcode)[T.98065]	0.3751	0.033	11.303	0.000
0.310 0.440				
C(zipcode)[T.98070]	0.0316	0.025	1.281	0.200 -
0.017 0.080				
C(zipcode)[T.98072]	0.3056	0.036	8.412	0.000
0.234 0.377				
C(zipcode)[T.98074]	0.3975	0.031	12.976	0.000
0.337 0.458				
C(zipcode)[T.98075]	0.4342	0.030	14.467	0.000
0.375 0.493				
C(zipcode)[T.98077]	0.2957	0.038	7.828	0.000
0.222 0.370				
C(zipcode)[T.98092]	0.0827	0.015	5.615	0.000
0.054 0.112				
C(zipcode)[T.98102]	0.7057	0.032	21.929	0.000
0.643 0.769				
C(zipcode)[T.98103]	0.5476	0.030	18.404	0.000
0.489 0.606				
C(zipcode)[T.98105]	0.6996	0.031	22.780	0.000
0.639 0.760				
C(zipcode)[T.98106]	0.0773	0.024	3.215	0.001

0.030	0.125				
C(zipcode)[T.98107]	0.5557	0.031	18.090	0.000	
0.495	0.616				
C(zipcode)[T.98108]	0.1054	0.026	4.032	0.000	
0.054	0.157				
C(zipcode)[T.98109]	0.7152	0.032	22.316	0.000	
0.652	0.778				
C(zipcode)[T.98112]	0.7971	0.029	27.722	0.000	
0.741	0.853				
C(zipcode)[T.98115]	0.5557	0.030	18.429	0.000	
0.497	0.615				
C(zipcode)[T.98116]	0.4564	0.026	17.498	0.000	
0.405	0.508				
C(zipcode)[T.98117]	0.5206	0.030	17.083	0.000	
0.461	0.580				
C(zipcode)[T.98118]	0.2321	0.024	9.816	0.000	
0.186	0.278				
C(zipcode)[T.98119]	0.6880	0.030	22.718	0.000	
0.629	0.747				
C(zipcode)[T.98122]	0.5506	0.028	19.883	0.000	
0.496	0.605				
C(zipcode)[T.98125]	0.3074	0.032	9.518	0.000	
0.244	0.371				
C(zipcode)[T.98126]	0.2699	0.024	11.032	0.000	
0.222	0.318				
C(zipcode)[T.98133]	0.1978	0.033	5.936	0.000	
0.132	0.263				
C(zipcode)[T.98136]	0.3992	0.025	15.964	0.000	
0.350	0.448				
C(zipcode)[T.98144]	0.4225	0.026	16.096	0.000	
0.371	0.474				
C(zipcode)[T.98146]	0.0206	0.023	0.903	0.367	-
0.024	0.065				
C(zipcode)[T.98148]	-0.0299	0.029	-1.026	0.305	-
0.087	0.027				
C(zipcode)[T.98155]	0.1894	0.035	5.452	0.000	
0.121	0.257				
C(zipcode)[T.98166]	0.0814	0.021	3.896	0.000	
0.040	0.122				

C(zipcode)[T.98168]	-0.1559	0.022	-7.012	0.000	-
0.200 -0.112					
C(zipcode)[T.98177]	0.3255	0.035	9.351	0.000	
0.257 0.394					
C(zipcode)[T.98178]	-0.0605	0.023	-2.631	0.009	-
0.105 -0.015					
C(zipcode)[T.98188]	-0.0910	0.023	-3.972	0.000	-
0.136 -0.046					
C(zipcode)[T.98198]	-0.0722	0.017	-4.211	0.000	-
0.106 -0.039					
C(zipcode)[T.98199]	0.5527	0.029	18.794	0.000	
0.495 0.610					
C(renovated)[T.True]	0.0579	0.007	8.671	0.000	
0.045 0.071					
np.log(sqft_living)	0.4225	0.006	66.665	0.000	
0.410 0.435					
np.log(sqft_lot)	0.0709	0.002	34.107	0.000	
0.067 0.075					
bedrooms	-0.0141	0.002	-7.827	0.000	-
0.018 -0.011					
floors	0.0175	0.003	5.388	0.000	
0.011 0.024					
bathrooms	0.0349	0.003	11.905	0.000	
0.029 0.041					
waterfront	0.2312	0.008	28.748	0.000	
0.215 0.247					
condition	0.0451	0.002	20.677	0.000	
0.041 0.049					
view	0.0637	0.002	32.732	0.000	
0.060 0.067					
grade	0.1115	0.002	58.283	0.000	
0.108 0.115					
yr_built	-0.0004	7.51e-05	-4.691	0.000	-
0.000 -0.000					
lat	250.2116	22.126	11.308	0.000	20
6.842 293.581					
I(lat ** 2)	-2.6267	0.233	-11.283	0.000	-
3.083 -2.170					
long	-0.4128	0.052	-7.911	0.000	-

```

0.515      -0.311
=====
=====
Omnibus:                  1424.478   Durbin-Watson:
2.001
Prob(Omnibus):            0.000    Jarque-Bera (JB):      5
984.714
Skew:                      -0.188   Prob(JB):
0.00
Kurtosis:                 5.550    Cond. No.
4.47e+19
=====
=====
```

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 9.74e-29. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

In [70]: *## effect of a waterfront view different for houses that were recently renovated*

```

formula = ('np.log(price) ~ np.log(sqft_living)*C(renovated) + np.log(s
qft_lot) + bedrooms + floors + bathrooms '
          ' + condition + view + grade + yr_built + lat*C(waterfront)
          + I(lat**2) '
          '+ long + C(zipcode)')
mod7 = smf.ols (formula=formula, data=houses).fit()
print(mod7.summary())
```

*Taped*

### OLS Regression Results

```

=====
=====
Dep. Variable:      np.log(price)   R-squared:
0.880
Model:                  OLS   Adj. R-squared:
0.879
```

Method: Least Squares F-statistic:  
 1855.  
 Date: Thu, 14 Nov 2019 Prob (F-statistic):  
 0.00  
 Time: 01:50:05 Log-Likelihood:  
 6088.9  
 No. Observations: 21613 AIC: -1.  
 201e+04 BIC: -1.  
 Df Residuals: 21527  
 132e+04  
 Df Model: 85  
 Covariance Type: nonrobust  
 ======  
 ======  
 t P>|t| [0.025 0.975] coef std err  
 Intercept -5894.6846 525.102 -11.2  
 26 0.000 -6923.924 -4865.445  
 C(renovated)[T.True] -0.1618 0.108 -1.4  
 92 0.136 -0.374 0.051  
 C(waterfront)[T.1] -40.1983 6.092 -6.5  
 99 0.000 -52.139 -28.258  
 C(zipcode)[T.98002] 0.0202 0.017 1.2  
 22 0.222 -0.012 0.053  
 C(zipcode)[T.98003] -0.0110 0.015 -0.7  
 44 0.457 -0.040 0.018  
 C(zipcode)[T.98004] 0.9138 0.028 32.5  
 55 0.000 0.859 0.969  
 C(zipcode)[T.98005] 0.5232 0.030 17.4  
 97 0.000 0.465 0.582  
 C(zipcode)[T.98006] 0.4614 0.026 18.0  
 03 0.000 0.411 0.512  
 C(zipcode)[T.98007] 0.4496 0.031 14.5  
 75 0.000 0.389 0.510  
 C(zipcode)[T.98008] 0.4595 0.029 15.6

07	0.000	0.402	0.517			
C(zipcode)[T.98010]				0.3098	0.025	12.2
52	0.000	0.260	0.359			
C(zipcode)[T.98011]				0.2707	0.037	7.4
14	0.000	0.199	0.342			
C(zipcode)[T.98014]				0.2041	0.041	5.0
40	0.000	0.125	0.283			
C(zipcode)[T.98019]				0.2244	0.040	5.6
80	0.000	0.147	0.302			
C(zipcode)[T.98022]				0.3386	0.024	13.8
65	0.000	0.291	0.387			
C(zipcode)[T.98023]				-0.0618	0.014	-4.5
67	0.000	-0.088	-0.035			
C(zipcode)[T.98024]				0.3204	0.037	8.6
97	0.000	0.248	0.393			
C(zipcode)[T.98027]				0.3825	0.026	14.4
96	0.000	0.331	0.434			
C(zipcode)[T.98028]				0.2113	0.035	5.9
57	0.000	0.142	0.281			
C(zipcode)[T.98029]				0.4816	0.030	16.2
76	0.000	0.424	0.540			
C(zipcode)[T.98030]				0.0047	0.017	0.2
73	0.785	-0.029	0.038			
C(zipcode)[T.98031]				-0.0210	0.019	-1.1
37	0.256	-0.057	0.015			
C(zipcode)[T.98032]				-0.1271	0.020	-6.2
52	0.000	-0.167	-0.087			
C(zipcode)[T.98033]				0.5777	0.031	18.7
78	0.000	0.517	0.638			
C(zipcode)[T.98034]				0.3312	0.033	10.1
59	0.000	0.267	0.395			
C(zipcode)[T.98038]				0.1948	0.019	10.1
68	0.000	0.157	0.232			
C(zipcode)[T.98039]				1.0842	0.037	29.2
89	0.000	1.012	1.157			
C(zipcode)[T.98040]				0.6664	0.026	25.8
87	0.000	0.616	0.717			
C(zipcode)[T.98042]				0.0535	0.016	3.2
62	0.001	0.021	0.086			

C(zipcode)[T.98045]			0.3219	0.035	9.0
72 0.000	0.252	0.391			
C(zipcode)[T.98052]			0.4594	0.031	14.6
31 0.000	0.398	0.521			
C(zipcode)[T.98053]			0.4579	0.034	13.6
15 0.000	0.392	0.524			
C(zipcode)[T.98055]			-0.0016	0.021	-0.0
77 0.939	-0.043	0.040			
C(zipcode)[T.98056]			0.1488	0.023	6.4
42 0.000	0.103	0.194			
C(zipcode)[T.98058]			0.0417	0.020	2.0
66 0.039	0.002	0.081			
C(zipcode)[T.98059]			0.2086	0.023	9.1
88 0.000	0.164	0.253			
C(zipcode)[T.98065]			0.3813	0.033	11.4
97 0.000	0.316	0.446			
C(zipcode)[T.98070]			0.0588	0.025	2.3
58 0.018	0.010	0.108			
C(zipcode)[T.98072]			0.3115	0.036	8.5
81 0.000	0.240	0.383			
C(zipcode)[T.98074]			0.4032	0.031	13.1
69 0.000	0.343	0.463			
C(zipcode)[T.98075]			0.4402	0.030	14.6
76 0.000	0.381	0.499			
C(zipcode)[T.98077]			0.3025	0.038	8.0
14 0.000	0.229	0.377			
C(zipcode)[T.98092]			0.0827	0.015	5.6
23 0.000	0.054	0.112			
C(zipcode)[T.98102]			0.7108	0.032	22.1
04 0.000	0.648	0.774			
C(zipcode)[T.98103]			0.5530	0.030	18.6
00 0.000	0.495	0.611			
C(zipcode)[T.98105]			0.7038	0.031	22.9
34 0.000	0.644	0.764			
C(zipcode)[T.98106]			0.0819	0.024	3.4
06 0.001	0.035	0.129			
C(zipcode)[T.98107]			0.5610	0.031	18.2
77 0.000	0.501	0.621			
C(zipcode)[T.98108]			0.1102	0.026	4.2

20	0.000	0.059	0.161			
C(zipcode)[T.98109]				0.7209	0.032	22.5
09	0.000	0.658	0.784			
C(zipcode)[T.98112]				0.8029	0.029	27.9
44	0.000	0.747	0.859			
C(zipcode)[T.98115]				0.5610	0.030	18.6
19	0.000	0.502	0.620			
C(zipcode)[T.98116]				0.4617	0.026	17.7
11	0.000	0.411	0.513			
C(zipcode)[T.98117]				0.5258	0.030	17.2
67	0.000	0.466	0.586			
C(zipcode)[T.98118]				0.2367	0.024	10.0
20	0.000	0.190	0.283			
C(zipcode)[T.98119]				0.6936	0.030	22.9
19	0.000	0.634	0.753			
C(zipcode)[T.98122]				0.5560	0.028	20.0
91	0.000	0.502	0.610			
C(zipcode)[T.98125]				0.3099	0.032	9.6
05	0.000	0.247	0.373			
C(zipcode)[T.98126]				0.2744	0.024	11.2
25	0.000	0.227	0.322			
C(zipcode)[T.98133]				0.2028	0.033	6.0
93	0.000	0.138	0.268			
C(zipcode)[T.98136]				0.4039	0.025	16.1
63	0.000	0.355	0.453			
C(zipcode)[T.98144]				0.4270	0.026	16.2
78	0.000	0.376	0.478			
C(zipcode)[T.98146]				0.0258	0.023	1.1
31	0.258	-0.019	0.071			
C(zipcode)[T.98148]				-0.0268	0.029	-0.9
21	0.357	-0.084	0.030			
C(zipcode)[T.98155]				0.1921	0.035	5.5
38	0.000	0.124	0.260			
C(zipcode)[T.98166]				0.0895	0.021	4.2
85	0.000	0.049	0.131			
C(zipcode)[T.98168]				-0.1520	0.022	-6.8
37	0.000	-0.196	-0.108			
C(zipcode)[T.98177]				0.3296	0.035	9.4
77	0.000	0.261	0.398			

	C(zipcode)[T.98178]			-0.0550	0.023	-2.3
97	0.017	-0.100	-0.010			
	C(zipcode)[T.98188]			-0.0875	0.023	-3.8
21	0.000	-0.132	-0.043			
	C(zipcode)[T.98198]			-0.0646	0.017	-3.7
63	0.000	-0.098	-0.031			
	C(zipcode)[T.98199]			0.5579	0.029	18.9
85	0.000	0.500	0.616			
	np.log(sqft_living)			0.4213	0.006	66.2
85	0.000	0.409	0.434			
	np.log(sqft_living):C(renovated)[T.True]			0.0289	0.014	2.0
34	0.042	0.001	0.057			
	np.log(sqft_lot)			0.0708	0.002	34.0
49	0.000	0.067	0.075			
	bedrooms			-0.0142	0.002	-7.8
55	0.000	-0.018	-0.011			
	floors			0.0174	0.003	5.3
45	0.000	0.011	0.024			
	bathrooms			0.0345	0.003	11.7
63	0.000	0.029	0.040			
	condition			0.0452	0.002	20.7
31	0.000	0.041	0.049			
	view			0.0634	0.002	32.6
25	0.000	0.060	0.067			
	grade			0.1114	0.002	58.2
85	0.000	0.108	0.115			
	yr_built			-0.0003	7.51e-05	-4.6
43	0.000	-0.000	-0.000			
	lat			245.7639	22.112	11.1
14	0.000	202.423	289.105			
	lat:C(waterfront)[T.1]			0.8552	0.128	6.6
74	0.000	0.604	1.106			
	I(lat ** 2)			-2.5800	0.233	-11.0
89	0.000	-3.036	-2.124			
	long			-0.4142	0.052	-7.9
44	0.000	-0.516	-0.312			
<hr/>						
<hr/>						
Omnibus:						
1409.500 Durbin-Watson:						

```
2.001

Prob(Omnibus):          0.000   Jarque-Bera (JB):      5
902.036
Skew:                  -0.183   Prob(JB):
0.00
Kurtosis:                5.534   Cond. No.
1.27e+09
=====
=====
```

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.27e+09. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [71]: print(mod7.aic)
```

```
-12005.868261991935
```

We can see that both the effect and renovations have a positive impact on price. The effect of a waterfront view is 46.25 percent on prices of comparable homes, while the effect of renovations is 5.794 percent. So far we have looked at global effects of predictors, irrespective of the levels of the other variables. However we might ask, is the effect of a waterfront view different for houses that were recently renovated? To answer this question we need to add an interaction term.

```
In [72]: formula = ('np.log(price) ~ np.log(sqft_living)*waterfront + np.log(sqf
t_living)*renovated + np.log(sqft_lot)'
                 '+ bedrooms + floors + bathrooms '
                 '+ waterfront + condition + view + grade + yr_built + lat +
                 I(lat**2) + long + C(zipcode)')
mod8= smf.ols(formula=formula, data=houses).fit()
print(mod8.summary())
```

### OLS Regression Results

=====									
=====									
Dep. Variable:	np.log(price)		R-squared:						
0.880									
Model:	OLS		Adj. R-squared:						
0.879									
Method:	Least Squares		F-statistic:						
1851.									
Date:	Thu, 14 Nov 2019		Prob (F-statistic):						
0.00									
Time:	01:50:06		Log-Likelihood:						
6068.8									
No. Observations:	21613		AIC:	-1.					
197e+04									
Df Residuals:	21527		BIC:	-1.					
128e+04									
Df Model:	85								
Covariance Type:	nonrobust								
=====									
=====									
P> t	[0.025	0.975]	coef	std err	t				
-----									
Intercept			-5996.9104	525.362	-11.415				
0.000	-7026.658	-4967.163							
renovated[T.True]			-0.1488	0.109	-1.368				
0.171	-0.362	0.064							
C(zipcode)[T.98002]			0.0200	0.017	1.214				
0.225	-0.012	0.052							
C(zipcode)[T.98003]			-0.0110	0.015	-0.748				
0.454	-0.040	0.018							
C(zipcode)[T.98004]			0.9079	0.028	32.329				
0.000	0.853	0.963							
C(zipcode)[T.98005]			0.5170	0.030	17.283				

0.000	0.458	0.576			
C(zipcode)[T.98006]			0.4555	0.026	17.769
0.000	0.405	0.506	0.4433	0.031	14.364
C(zipcode)[T.98007]			0.4537	0.029	15.400
0.000	0.383	0.504			
C(zipcode)[T.98008]			0.3081	0.025	12.176
0.000	0.396	0.511			
C(zipcode)[T.98010]			0.2647	0.037	7.244
0.000	0.259	0.358			
C(zipcode)[T.98011]			0.1967	0.041	4.853
0.000	0.193	0.336			
C(zipcode)[T.98014]			0.2174	0.040	5.501
0.000	0.117	0.276			
C(zipcode)[T.98019]			0.3414	0.024	13.965
0.000	0.140	0.295			
C(zipcode)[T.98022]			-0.0629	0.014	-4.645
0.000	0.293	0.389			
C(zipcode)[T.98023]			0.3138	0.037	8.513
0.000	-0.089	-0.036			
C(zipcode)[T.98024]			0.3767	0.026	14.270
0.000	0.242	0.386			
C(zipcode)[T.98027]			0.2060	0.036	5.803
0.000	0.325	0.428			
C(zipcode)[T.98028]			0.4753	0.030	16.055
0.000	0.136	0.276			
C(zipcode)[T.98029]			0.0026	0.017	0.151
0.000	0.417	0.533			
C(zipcode)[T.98030]			-0.0240	0.019	-1.298
0.880	-0.031	0.036			
C(zipcode)[T.98031]			-0.1290	0.020	-6.341
0.194	-0.060	0.012			
C(zipcode)[T.98032]			0.5721	0.031	18.583
0.000	-0.169	-0.089			
C(zipcode)[T.98033]			0.3264	0.033	10.003
0.000	0.512	0.632			
C(zipcode)[T.98034]			0.1922	0.019	10.024
0.000	0.262	0.390			
C(zipcode)[T.98038]			0.230		
0.000	0.155				

C(zipcode)[T.98039]			1.0794	0.037	29.137
0.000	1.007	1.152			
C(zipcode)[T.98040]			0.6606	0.026	25.646
0.000	0.610	0.711			
C(zipcode)[T.98042]			0.0513	0.016	3.123
0.002	0.019	0.083			
C(zipcode)[T.98045]			0.3160	0.036	8.901
0.000	0.246	0.386			
C(zipcode)[T.98052]			0.4533	0.031	14.428
0.000	0.392	0.515			
C(zipcode)[T.98053]			0.4511	0.034	13.405
0.000	0.385	0.517			
C(zipcode)[T.98055]			-0.0058	0.021	-0.271
0.786	-0.047	0.036			
C(zipcode)[T.98056]			0.1434	0.023	6.209
0.000	0.098	0.189			
C(zipcode)[T.98058]			0.0376	0.020	1.863
0.062	-0.002	0.077			
C(zipcode)[T.98059]			0.2036	0.023	8.964
0.000	0.159	0.248			
C(zipcode)[T.98065]			0.3748	0.033	11.296
0.000	0.310	0.440			
C(zipcode)[T.98070]			0.0400	0.025	1.606
0.108	-0.009	0.089			
C(zipcode)[T.98072]			0.3052	0.036	8.403
0.000	0.234	0.376			
C(zipcode)[T.98074]			0.3974	0.031	12.973
0.000	0.337	0.457			
C(zipcode)[T.98075]			0.4345	0.030	14.477
0.000	0.376	0.493			
C(zipcode)[T.98077]			0.2959	0.038	7.835
0.000	0.222	0.370			
C(zipcode)[T.98092]			0.0826	0.015	5.610
0.000	0.054	0.111			
C(zipcode)[T.98102]			0.7050	0.032	21.911
0.000	0.642	0.768			
C(zipcode)[T.98103]			0.5472	0.030	18.393
0.000	0.489	0.605			
C(zipcode)[T.98105]			0.6987	0.031	22.753

0.000	0.639	0.759			
C(zipcode)[T.98106]			0.0770	0.024	3.200
0.001	0.030	0.124			
C(zipcode)[T.98107]			0.5552	0.031	18.077
0.000	0.495	0.615			
C(zipcode)[T.98108]			0.1050	0.026	4.019
0.000	0.054	0.156			
C(zipcode)[T.98109]			0.7151	0.032	22.317
0.000	0.652	0.778			
C(zipcode)[T.98112]			0.7971	0.029	27.728
0.000	0.741	0.853			
C(zipcode)[T.98115]			0.5553	0.030	18.418
0.000	0.496	0.614			
C(zipcode)[T.98116]			0.4565	0.026	17.503
0.000	0.405	0.508			
C(zipcode)[T.98117]			0.5201	0.030	17.069
0.000	0.460	0.580			
C(zipcode)[T.98118]			0.2313	0.024	9.787
0.000	0.185	0.278			
C(zipcode)[T.98119]			0.6878	0.030	22.717
0.000	0.628	0.747			
C(zipcode)[T.98122]			0.5502	0.028	19.873
0.000	0.496	0.605			
C(zipcode)[T.98125]			0.3067	0.032	9.499
0.000	0.243	0.370			
C(zipcode)[T.98126]			0.2695	0.024	11.016
0.000	0.222	0.317			
C(zipcode)[T.98133]			0.1971	0.033	5.919
0.000	0.132	0.262			
C(zipcode)[T.98136]			0.3993	0.025	15.971
0.000	0.350	0.448			
C(zipcode)[T.98144]			0.4215	0.026	16.062
0.000	0.370	0.473			
C(zipcode)[T.98146]			0.0208	0.023	0.909
0.363	-0.024	0.066			
C(zipcode)[T.98148]			-0.0301	0.029	-1.033
0.302	-0.087	0.027			
C(zipcode)[T.98155]			0.1885	0.035	5.429
0.000	0.120	0.257			

C(zipcode)[T.98166]			0.0814	0.021	3.896
0.000	0.040	0.122			
C(zipcode)[T.98168]			-0.1563	0.022	-7.029
0.000	-0.200	-0.113			
C(zipcode)[T.98177]			0.3248	0.035	9.332
0.000	0.257	0.393			
C(zipcode)[T.98178]			-0.0602	0.023	-2.618
0.009	-0.105	-0.015			
C(zipcode)[T.98188]			-0.0912	0.023	-3.979
0.000	-0.136	-0.046			
C(zipcode)[T.98198]			-0.0718	0.017	-4.189
0.000	-0.105	-0.038			
C(zipcode)[T.98199]			0.5526	0.029	18.793
0.000	0.495	0.610			
np.log(sqft_living)			0.4210	0.006	66.165
0.000	0.409	0.434			
np.log(sqft_living):renovated[T.True]			0.0271	0.014	1.904
0.057	-0.001	0.055			
waterfront			-0.0103	0.227	-0.045
0.964	-0.455	0.434			
np.log(sqft_living):waterfront			0.0593	0.028	2.084
0.037	0.004	0.115			
np.log(sqft_lot)			0.0708	0.002	34.042
0.000	0.067	0.075			
bedrooms			-0.0141	0.002	-7.781
0.000	-0.018	-0.011			
floors			0.0175	0.003	5.388
0.000	0.011	0.024			
bathrooms			0.0345	0.003	11.743
0.000	0.029	0.040			
condition			0.0452	0.002	20.717
0.000	0.041	0.049			
view			0.0636	0.002	32.694
0.000	0.060	0.067			
grade			0.1114	0.002	58.199
0.000	0.108	0.115			
yr_built			-0.0003	7.51e-05	-4.608
0.000	-0.000	-0.000			
lat			250.0677	22.123	11.304

```

          0.000    206.705    293.430
I(lat ** 2)      0.000     -3.081     -2.169      -2.6252     0.233    -11.278
long            0.000     -0.513     -0.309      -0.4112     0.052    -7.879
=====
=====
Omnibus:           1424.704 Durbin-Watson:
2.001
Prob(Omnibus):    0.000 Jarque-Bera (JB):      5
990.854
Skew:              -0.187 Prob(JB):
0.00
Kurtosis:         5.552 Cond. No.
1.27e+09
=====
=====
Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is
correctly specified.
[2] The condition number is large, 1.27e+09. This might indicate that t
here are
strong multicollinearity or other numerical problems.

```

In [73]: `print(mod8.aic)`

```
-11965.54786190483
```

### Exercise 10: (15 mts)

Our reference model (mod7) contained the following predictors: `bedrooms` , `floors` ,  
`bathrooms` , `condition` , `view` , `grade` , `yr_built` , `long` , `log(sqft_living) * C(renovated)` , `lat * C(waterfront)` , `I(lat**2)` , `log(sqft_lot)` and  
`C(zipcode)` . The AIC for this model was  $-12005.86$  and the R2 is  $0.8798$ .

Expand this model above by doing the following:

1. Add a term which accounts for the square of the year the house was built
2. Add an interaction term for the presence of a basement in affecting the relationship between the longitude coordinate and the price of a house

Compare the model fit and AIC with the previous model.

**Answer.** One possible solution is shown below:

```
In [74]: formula = ('np.log(price) ~ np.log(sqft_living)*C(renovated) + np.log(s  
qft_lot) + bedrooms + floors + bathrooms '  
    '+ condition + view + grade + yr_built + lat*C(waterfront) +  
    I(lat**2) + long + C(zipcode)'  
    '+ I(yr_built**2)')  
mod9 = smf.ols(formula=formula, data=houses).fit()  
print(mod9.summary())
```

### OLS Regression Results

```
=====
```

Dep. Variable:	np.log(price)	R-squared:	0.884
Model:	OLS	Adj. R-squared:	0.883
Method:	Least Squares	F-statistic:	1899.
Date:	Thu, 14 Nov 2019	Prob (F-statistic):	0.00
Time:	01:50:08	Log-Likelihood:	6428.5
No. Observations:	21613	AIC:	-1.268e+04
Df Residuals:	21526	BIC:	-1.199e+04
Df Model:	86		
Covariance Type:	nonrobust		

t	P> t	[0.025	0.975]	coef	std err
Intercept				-6012.6468	516.949
31	0.000	-7025.905	-4999.389		-11.6
C(renovated)[T.True]				-0.3822	0.107
69	0.000	-0.592	-0.172		-3.5
C(waterfront)[T.1]				-40.7918	5.997
02	0.000	-52.547	-29.037		-6.8
C(zipcode)[T.98002]				0.0233	0.016
36	0.151	-0.009	0.055		1.4
C(zipcode)[T.98003]				0.0072	0.015
98	0.619	-0.021	0.036		0.4
C(zipcode)[T.98004]				0.9283	0.028
85	0.000	0.874	0.982		33.5
C(zipcode)[T.98005]				0.5507	0.029
96	0.000	0.493	0.608		18.6
C(zipcode)[T.98006]				0.4802	0.025
26	0.000	0.431	0.530		19.0
C(zipcode)[T.98007]				0.4816	0.030
44	0.000	0.422	0.541		15.8
C(zipcode)[T.98008]				0.4917	0.029
50	0.000	0.435	0.549		16.9
C(zipcode)[T.98010]				0.3011	0.025
96	0.000	0.252	0.350		12.0
C(zipcode)[T.98011]				0.2840	0.036
98	0.000	0.213	0.354		7.8
C(zipcode)[T.98014]				0.2023	0.040
73	0.000	0.124	0.280		5.0
C(zipcode)[T.98019]				0.2224	0.039
19	0.000	0.146	0.299		5.7
C(zipcode)[T.98022]				0.3422	0.024
30	0.000	0.295	0.389		14.2
C(zipcode)[T.98023]				-0.0471	0.013
31	0.000	-0.073	-0.021		-3.5
C(zipcode)[T.98024]				0.3110	0.036
					8.5

74	0.000	0.240	0.382			
C(zipcode)[T.98027]				0.3871	0.026	14.9
03	0.000	0.336	0.438			
C(zipcode)[T.98028]				0.2208	0.035	6.3
22	0.000	0.152	0.289			
C(zipcode)[T.98029]				0.4922	0.029	16.8
97	0.000	0.435	0.549			
C(zipcode)[T.98030]				0.0048	0.017	0.2
88	0.774	-0.028	0.038			
C(zipcode)[T.98031]				-0.0093	0.018	-0.5
13	0.608	-0.045	0.026			
C(zipcode)[T.98032]				-0.1095	0.020	-5.4
66	0.000	-0.149	-0.070			
C(zipcode)[T.98033]				0.5869	0.030	19.3
74	0.000	0.527	0.646			
C(zipcode)[T.98034]				0.3515	0.032	10.9
50	0.000	0.289	0.414			
C(zipcode)[T.98038]				0.1913	0.019	10.1
41	0.000	0.154	0.228			
C(zipcode)[T.98039]				1.0983	0.036	30.1
36	0.000	1.027	1.170			
C(zipcode)[T.98040]				0.6911	0.025	27.2
51	0.000	0.641	0.741			
C(zipcode)[T.98042]				0.0579	0.016	3.5
81	0.000	0.026	0.090			
C(zipcode)[T.98045]				0.3316	0.035	9.4
91	0.000	0.263	0.400			
C(zipcode)[T.98052]				0.4753	0.031	15.3
73	0.000	0.415	0.536			
C(zipcode)[T.98053]				0.4433	0.033	13.3
88	0.000	0.378	0.508			
C(zipcode)[T.98055]				0.0032	0.021	0.1
55	0.877	-0.038	0.044			
C(zipcode)[T.98056]				0.1495	0.023	6.5
77	0.000	0.105	0.194			
C(zipcode)[T.98058]				0.0579	0.020	2.9
10	0.004	0.019	0.097			
C(zipcode)[T.98059]				0.2036	0.022	9.1
07	0.000	0.160	0.247			

C(zipcode)[T.98065]			0.3734	0.033	11.4
35 0.000	0.309	0.437			
C(zipcode)[T.98070]			0.0264	0.025	1.0
74 0.283	-0.022	0.075			
C(zipcode)[T.98072]			0.3251	0.036	9.0
97 0.000	0.255	0.395			
C(zipcode)[T.98074]			0.4183	0.030	13.8
75 0.000	0.359	0.477			
C(zipcode)[T.98075]			0.4458	0.030	15.0
98 0.000	0.388	0.504			
C(zipcode)[T.98077]			0.3151	0.037	8.4
79 0.000	0.242	0.388			
C(zipcode)[T.98092]			0.0852	0.014	5.8
79 0.000	0.057	0.114			
C(zipcode)[T.98102]			0.6984	0.032	22.0
59 0.000	0.636	0.760			
C(zipcode)[T.98103]			0.5360	0.029	18.3
08 0.000	0.479	0.593			
C(zipcode)[T.98105]			0.7033	0.030	23.2
81 0.000	0.644	0.763			
C(zipcode)[T.98106]			0.0714	0.024	3.0
16 0.003	0.025	0.118			
C(zipcode)[T.98107]			0.5413	0.030	17.9
09 0.000	0.482	0.601			
C(zipcode)[T.98108]			0.0959	0.026	3.7
29 0.000	0.046	0.146			
C(zipcode)[T.98109]			0.7021	0.032	22.2
65 0.000	0.640	0.764			
C(zipcode)[T.98112]			0.7957	0.028	28.1
28 0.000	0.740	0.851			
C(zipcode)[T.98115]			0.5674	0.030	19.1
29 0.000	0.509	0.626			
C(zipcode)[T.98116]			0.4499	0.026	17.5
29 0.000	0.400	0.500			
C(zipcode)[T.98117]			0.5188	0.030	17.3
05 0.000	0.460	0.578			
C(zipcode)[T.98118]			0.2249	0.023	9.6
67 0.000	0.179	0.270			
C(zipcode)[T.98119]			0.6695	0.030	22.4

64	0.000	0.611	0.728			
C(zipcode)[T.98122]				0.5280	0.027	19.3
68	0.000	0.475	0.581			
C(zipcode)[T.98125]				0.3282	0.032	10.3
30	0.000	0.266	0.390			
C(zipcode)[T.98126]				0.2645	0.024	10.9
87	0.000	0.217	0.312			
C(zipcode)[T.98133]				0.2129	0.033	6.4
97	0.000	0.149	0.277			
C(zipcode)[T.98136]				0.3930	0.025	15.9
73	0.000	0.345	0.441			
C(zipcode)[T.98144]				0.4047	0.026	15.6
64	0.000	0.354	0.455			
C(zipcode)[T.98146]				0.0290	0.022	1.2
88	0.198	-0.015	0.073			
C(zipcode)[T.98148]				-0.0134	0.029	-0.4
67	0.640	-0.070	0.043			
C(zipcode)[T.98155]				0.2101	0.034	6.1
50	0.000	0.143	0.277			
C(zipcode)[T.98166]				0.0979	0.021	4.7
56	0.000	0.058	0.138			
C(zipcode)[T.98168]				-0.1458	0.022	-6.6
63	0.000	-0.189	-0.103			
C(zipcode)[T.98177]				0.3481	0.034	10.1
62	0.000	0.281	0.415			
C(zipcode)[T.98178]				-0.0416	0.023	-1.8
39	0.066	-0.086	0.003			
C(zipcode)[T.98188]				-0.0750	0.023	-3.3
25	0.001	-0.119	-0.031			
C(zipcode)[T.98198]				-0.0510	0.017	-3.0
20	0.003	-0.084	-0.018			
C(zipcode)[T.98199]				0.5676	0.029	19.6
19	0.000	0.511	0.624			
np.log(sqft_living)				0.4126	0.006	65.8
52	0.000	0.400	0.425			
np.log(sqft_living):C(renovated)[T.True]				0.0592	0.014	4.2
29	0.000	0.032	0.087			
np.log(sqft_lot)				0.0852	0.002	40.2
04	0.000	0.081	0.089			

bedrooms				-0.0105	0.002	-5.9
16	0.000	-0.014	-0.007			
floors				-0.0125	0.003	-3.6
74	0.000	-0.019	-0.006			
bathrooms				0.0251	0.003	8.6
23	0.000	0.019	0.031			
condition				0.0540	0.002	24.8
61	0.000	0.050	0.058			
view				0.0667	0.002	34.7
70	0.000	0.063	0.070			
grade				0.1081	0.002	57.3
06	0.000	0.104	0.112			
yr_built				-0.1896	0.007	-26.2
59	0.000	-0.204	-0.175			
lat				258.2540	21.773	11.8
61	0.000	215.577	300.931			
lat:C(waterfront)[T.1]				0.8676	0.126	6.8
78	0.000	0.620	1.115			
I(lat ** 2)				-2.7113	0.229	-11.8
35	0.000	-3.160	-2.262			
long				-0.4630	0.051	-9.0
15	0.000	-0.564	-0.362			
I(yr_built ** 2)				4.833e-05	1.84e-06	26.2
12	0.000	4.47e-05	5.19e-05			
<hr/>						
<hr/>						
Omnibus:		1552.582	Durbin-Watson:			
2.005						
Prob(Omnibus):		0.000	Jarque-Bera (JB):			6
876.465						
Skew:		-0.214	Prob(JB):			
0.00						
Kurtosis:		5.730	Cond. No.			
1.64e+12						
<hr/>						
<hr/>						
=====						
=====						
Warnings:						
[1] Standard Errors assume that the covariance matrix of the errors is						

```
correctly specified.  
[2] The condition number is large, 1.64e+12. This might indicate that t  
here are  
strong multicollinearity or other numerical problems.
```

```
In [75]: print(mod9.aic)
```

```
-12682.956456731124
```

```
In [76]: # built dummy variable to separate houses with a basement and houses wi  
th no basement  
houses['has_basement'] = (houses['sqft_basement'] > 0) * 1.0
```

```
In [77]: # estimate a model with an interaction between the longitude coordinate  
# and the presence of a basement.  
formula = ('np.log(price) ~ np.log(sqft_living)*C(renovated) + np.log(s  
qft_lot) + bedrooms + floors + bathrooms '  
         '+ condition + view + grade + yr_built + lat * C(waterfront)  
         + I(lat**2) + long + C(zipcode) '  
         '+ has_basement * long')  
mod10 = smf.ols(formula=formula, data=houses).fit()  
print(mod10.summary())
```

### OLS Regression Results

```
=====
```

Dep. Variable:	np.log(price)	R-squared:
0.881		
Model:	OLS	Adj. R-squared:
0.881		
Method:	Least Squares	F-statistic:
1833.		
Date:	Thu, 14 Nov 2019	Prob (F-statistic):
0.00		
Time:	01:50:10	Log-Likelihood:
6199.6		
No. Observations:	21613	AIC:
		-1.

222e+04  
 Df Residuals: 21525 BIC: -1.  
 152e+04  
 Df Model: 87  
 Covariance Type: nonrobust

			coef	std err
t	P> t	[0.025 0.975]		
Intercept			-6308.3439	523.203
57	0.000	-7333.860 -5282.827		-12.0
C(renovated)[T.True]			-0.2039	0.108
88	0.059	-0.416 0.008		-1.8
C(waterfront)[T.1]			-40.0337	6.064
02	0.000	-51.920 -28.147		-6.6
C(zipcode)[T.98002]			0.0163	0.016
93	0.321	-0.016 0.048		0.9
C(zipcode)[T.98003]			-0.0093	0.015
34	0.526	-0.038 0.019		-0.6
C(zipcode)[T.98004]			0.9153	0.028
43	0.000	0.860 0.970		32.7
C(zipcode)[T.98005]			0.5272	0.030
99	0.000	0.469 0.586		17.6
C(zipcode)[T.98006]			0.4651	0.026
12	0.000	0.415 0.515		18.2
C(zipcode)[T.98007]			0.4502	0.031
55	0.000	0.390 0.510		14.6
C(zipcode)[T.98008]			0.4626	0.029
66	0.000	0.405 0.520		15.7
C(zipcode)[T.98010]			0.3097	0.025
12	0.000	0.260 0.359		12.3
C(zipcode)[T.98011]			0.2827	0.036
72	0.000	0.211 0.354		7.7
C(zipcode)[T.98014]			0.2061	0.040
15	0.000	0.127 0.285		5.1

C(zipcode)[T.98019]				0.2316	0.039	5.8
87 0.000	0.154	0.309				
C(zipcode)[T.98022]				0.3434	0.024	14.1
29 0.000	0.296	0.391				
C(zipcode)[T.98023]				-0.0590	0.013	-4.3
72 0.000	-0.085	-0.033				
C(zipcode)[T.98024]				0.3188	0.037	8.6
96 0.000	0.247	0.391				
C(zipcode)[T.98027]				0.3938	0.026	14.9
55 0.000	0.342	0.445				
C(zipcode)[T.98028]				0.2237	0.035	6.3
33 0.000	0.155	0.293				
C(zipcode)[T.98029]				0.4780	0.029	16.2
30 0.000	0.420	0.536				
C(zipcode)[T.98030]				0.0011	0.017	0.0
65 0.948	-0.032	0.034				
C(zipcode)[T.98031]				-0.0231	0.018	-1.2
55 0.210	-0.059	0.013				
C(zipcode)[T.98032]				-0.1237	0.020	-6.1
12 0.000	-0.163	-0.084				
C(zipcode)[T.98033]				0.5824	0.031	19.0
07 0.000	0.522	0.642				
C(zipcode)[T.98034]				0.3417	0.032	10.5
22 0.000	0.278	0.405				
C(zipcode)[T.98038]				0.1871	0.019	9.8
09 0.000	0.150	0.224				
C(zipcode)[T.98039]				1.0789	0.037	29.2
78 0.000	1.007	1.151				
C(zipcode)[T.98040]				0.6677	0.026	26.0
53 0.000	0.617	0.718				
C(zipcode)[T.98042]				0.0484	0.016	2.9
61 0.003	0.016	0.080				
C(zipcode)[T.98045]				0.3175	0.035	8.9
91 0.000	0.248	0.387				
C(zipcode)[T.98052]				0.4658	0.031	14.8
90 0.000	0.405	0.527				
C(zipcode)[T.98053]				0.4493	0.033	13.4
18 0.000	0.384	0.515				
C(zipcode)[T.98055]				-0.0038	0.021	-0.1

78	0.859	-0.045	0.038			
C(zipcode)[T.98056]				0.1430	0.023	6.2
22	0.000	0.098	0.188			
C(zipcode)[T.98058]				0.0396	0.020	1.9
71	0.049	0.000	0.079			
C(zipcode)[T.98059]				0.1981	0.023	8.7
60	0.000	0.154	0.242			
C(zipcode)[T.98065]				0.3699	0.033	11.2
05	0.000	0.305	0.435			
C(zipcode)[T.98070]				0.0581	0.025	2.3
42	0.019	0.009	0.107			
C(zipcode)[T.98072]				0.3258	0.036	9.0
09	0.000	0.255	0.397			
C(zipcode)[T.98074]				0.4026	0.030	13.2
06	0.000	0.343	0.462			
C(zipcode)[T.98075]				0.4321	0.030	14.4
70	0.000	0.374	0.491			
C(zipcode)[T.98077]				0.3103	0.038	8.2
56	0.000	0.237	0.384			
C(zipcode)[T.98092]				0.0815	0.015	5.5
68	0.000	0.053	0.110			
C(zipcode)[T.98102]				0.7378	0.032	23.0
24	0.000	0.675	0.801			
C(zipcode)[T.98103]				0.5717	0.030	19.3
07	0.000	0.514	0.630			
C(zipcode)[T.98105]				0.7255	0.031	23.7
33	0.000	0.666	0.785			
C(zipcode)[T.98106]				0.0914	0.024	3.8
20	0.000	0.045	0.138			
C(zipcode)[T.98107]				0.5814	0.031	19.0
20	0.000	0.522	0.641			
C(zipcode)[T.98108]				0.1200	0.026	4.6
16	0.000	0.069	0.171			
C(zipcode)[T.98109]				0.7400	0.032	23.2
05	0.000	0.678	0.803			
C(zipcode)[T.98112]				0.8242	0.029	28.7
94	0.000	0.768	0.880			
C(zipcode)[T.98115]				0.5812	0.030	19.3
67	0.000	0.522	0.640			

C(zipcode)[T.98116]				0.4745	0.026	18.2
69 0.000	0.424	0.525		0.5435	0.030	17.9
C(zipcode)[T.98117]				0.2461	0.024	10.4
23 0.000	0.484	0.603		0.7157	0.030	23.7
C(zipcode)[T.98118]				0.5749	0.028	20.8
64 0.000	0.200	0.292		0.3247	0.032	10.1
C(zipcode)[T.98119]				0.2831	0.024	11.6
34 0.000	0.657	0.775		0.2180	0.033	6.5
C(zipcode)[T.98122]				0.4158	0.025	16.6
55 0.000	0.521	0.629		0.4443	0.026	17.0
C(zipcode)[T.98125]				0.0253	0.023	1.1
07 0.000	0.262	0.388		-0.0326	0.029	-1.1
C(zipcode)[T.98126]				0.2059	0.035	5.9
31 0.000	0.235	0.331		0.0887	0.021	4.2
C(zipcode)[T.98133]				-0.1487	0.022	-6.7
78 0.000	0.153	0.283		0.3422	0.035	9.8
C(zipcode)[T.98136]				-0.0515	0.023	-2.2
99 0.000	0.367	0.465		-0.0885	0.023	-3.8
C(zipcode)[T.98144]				-0.0661	0.017	-3.8
07 0.000	0.393	0.495		0.5753	0.029	19.6
C(zipcode)[T.98146]						
12 0.266	-0.019	0.070				
C(zipcode)[T.98148]						
24 0.261	-0.089	0.024				
C(zipcode)[T.98155]						
60 0.000	0.138	0.274				
C(zipcode)[T.98166]						
66 0.000	0.048	0.129				
C(zipcode)[T.98168]						
27 0.000	-0.192	-0.105				
C(zipcode)[T.98177]						
84 0.000	0.274	0.410				
C(zipcode)[T.98178]						
51 0.024	-0.096	-0.007				
C(zipcode)[T.98188]						
84 0.000	-0.133	-0.044				
C(zipcode)[T.98198]						
73 0.000	-0.100	-0.033				
C(zipcode)[T.98199]						

39	0.000	0.518	0.633			
	np.log(sqft_living)			0.4518	0.007	67.2
61	0.000	0.439	0.465			
	np.log(sqft_living):C(renovated)[T.True]			0.0345	0.014	2.4
40	0.015	0.007	0.062			
	np.log(sqft_lot)			0.0669	0.002	31.8
95	0.000	0.063	0.071			
	bedrooms			-0.0156	0.002	-8.6
48	0.000	-0.019	-0.012			
	floors			-0.0062	0.004	-1.7
09	0.087	-0.013	0.001			
	bathrooms			0.0403	0.003	13.6
99	0.000	0.035	0.046			
	condition			0.0462	0.002	21.2
86	0.000	0.042	0.050			
	view			0.0657	0.002	33.8
54	0.000	0.062	0.069			
	grade			0.1069	0.002	55.4
79	0.000	0.103	0.111			
	yr_built			-0.0003	7.49e-05	-3.4
39	0.001	-0.000	-0.000			
	lat			263.2078	22.032	11.9
46	0.000	220.023	306.393			
	lat:C(waterfront)[T.1]			0.8518	0.128	6.6
78	0.000	0.602	1.102			
	I(lat ** 2)			-2.7637	0.232	-11.9
22	0.000	-3.218	-2.309			
	long			-0.4075	0.052	-7.8
38	0.000	-0.509	-0.306			
	has_basement			-7.2390	2.610	-2.7
74	0.006	-12.354	-2.124			
	has_basement:long			-0.0588	0.021	-2.7
55	0.006	-0.101	-0.017			
<hr/>						
<hr/>						
Omnibus:						
2.000						
Prob(Omnibus):						
307.671						
Durbin-Watson:						
1458.357						
Jarque-Bera (JB):						
6						

```
Skew:          -0.186  Prob(JB):  
    0.00  
Kurtosis:      5.620   Cond. No.  
 1.27e+09  
=====
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.  
[2] The condition number is large, 1.27e+09. This might indicate that there are strong multicollinearity or other numerical problems.

The effects are significant in both cases. Both models improve the fit of our reference models, but the addition of a square term in the building year has a stronger impact compared to the interaction between basement and longitude.

```
In [78]: # the r-squared results  
r7 = mod7.rsquared  
r9 = mod9.rsquared  
r10 = mod10.rsquared  
  
# the aic results  
aic7 = mod7.aic  
aic9 = mod9.aic  
aic10 = mod10.aic  
  
print("----- R Squared results -----")  
print("Model 7 -", r7)  
print("Model 9 -", r9)  
print("Model 10 -", r10)  
print("\n----- AIC results -----")  
print("AIC 7 -", aic7)  
print("AIC 9 -", aic9)  
print("AIC 10 -", aic10)
```

----- R Squared results -----

Model 7 - 0.8798452507027756  
Model 9 - 0.8835618615444407  
Model 10 - 0.8810696008728642

----- AIC results -----

AIC 7 - -12005.868261991935  
AIC 9 - -12682.956456731124  
AIC 10 - -12223.229661644618

## Conclusions (5 mts)

In this case, we applied various types of transformations to the predictor and response variables to improve the quality of our linear modeling. In particular, we found that fitting the logarithm of house prices allowed us to get better results. Using our understanding of transformations, we were able to effectively model nonlinear relationships, such as the quadratic relationship between latitude and the log of price. Finally, we tied in our understanding of interaction effects from previous EDA cases in order to directly model and quantify the interaction of renovation and waterfront status on square footage.

## Takeaways (5 mts)

Variable transformations are a powerful technique to improve the quality of our linear models. In particular:

1. Transforming the dependent variable can improve linearity and resolve the problem of uneven variance around the line of best fit.
2. Transforming the independent variables can be useful to improve the quality of the fit, capture nonlinear relationships between the independent and response variables, and test a wider range of hypotheses.
3. Interaction terms are a specific type of variable transformation, involving the product of two other independent variables. They can capture dependencies in the relationship between a predictor variable and the response variable on the value of a third variable.

