

case_9.1

May 29, 2020

1 How important is the income source of an online loan applicant?

```
[1]: import pandas as pd
import numpy as np
from scipy.stats import chi2_contingency
import statsmodels.api as sm
from sklearn.metrics import roc_curve, auc
from sklearn.model_selection import StratifiedKFold
import matplotlib.pyplot as plt
from scipy import interp
from statsmodels.formula.api import ols
import seaborn as sns
from matplotlib.widgets import Slider, Button, RadioButtons
from scipy.optimize import fsolve

#%get_ipython().magic(u'matplotlib inline')
```

1.1 Introduction (5 mts)

Business Context. Online peer-to-peer (P2P) lending has made the practice of borrowing and lending easy. In this form of lending, there is no in-person interview and a borrower can simply fill out an online form and get a loan approved. The information provided solely by a borrower is prone to exaggeration and distortion, especially when it comes to income. Every P2P lending company relies on a well-designed procedure that rejects borrowers with a high chance of not paying off their debts.

Rejecting anyone without a verified source of income is a relevant policy that lending platforms can roll out to help reduce the rate of bad loans. It is natural to suspect that if a person's income source cannot be verified, then they might default on the loan. However, from the perspective of a borrower, the verification process can be cumbersome and time-consuming and they may just switch to another platform due to this inconvenience.

Business Problem. As a data scientist at an emerging P2P lending company, you must answer the following question: **“Should the company verify the income source of an online loan applicant before approving their loan?”**

Analytical Context. The data is downloaded from [LendingClub Statistics](#) and it contains all loans issued from 2007 - 2012 along with their current loan status (fully paid or charged off). There

are ~50 variables that describe the borrowers and the loans; for the sake of reducing complexity, the company has already performed a pre-screening of these variables based on existing analyses from LendingClub to select nine relevant variables, such as annual income, LendingClub credit grade, home ownership, etc. We will use a new technique, **logistic regression**, to answer our question at hand.

The case is structured as follows: you will 1) explore the existing data to get a rough idea of how each variable in the dataset interacts with the current loan status; 2) look at potential confounding effects; 3) learn the basics of logistic regression models; and finally 4) fit a series of logistic regression models to determine whether or not verifying income source is significant.

1.2 Data exploration (45 mts)

Before we begin our exploratory analysis, let's take a look at the data at our disposal:

```
[3]: df = pd.read_csv("Lending_club_cleaned_2.csv")
df.loan_status = df.loan_status.astype(pd.api.types.
    →CategoricalDtype(categories=["Charged Off", "Fully Paid"]))
df.verification_status = df.verification_status.astype(pd.api.types.
    →CategoricalDtype(categories=['Not Verified', 'Source Verified', 'Verified']))
df.emp_length = df.emp_length.astype(pd.api.types.
    →CategoricalDtype(categories=["< 1 year", "1 year", "2 years", "3 years", "4
    →years", \
        "5 years", "6
    →years", "7 years", "8 years", "9 years", \
        "10+ years"]))
df.home_ownership = df.home_ownership.astype(pd.api.types.
    →CategoricalDtype(categories=["RENT", "MORTGAGE", "OWN", "OTHER"]))
df.term = df.term.astype(pd.api.types.CategoricalDtype(categories=["
    →36
    →months", " 60 months"]))
df.grade = df.grade.astype(pd.api.types.
    →CategoricalDtype(categories=["A", "B", "C", "D", "E", "F", "G"]))
df.int_rate = df.int_rate.str.rstrip("%").astype("float")
```

```
[4]: df.shape
```

```
[4]: (38705, 10)
```

```
[5]: df.head(10)
```

	loan_status	annual_inc	verification_status	emp_length	home_ownership	
0	Fully Paid	24000.0	Verified	10+ years	RENT	
1	Charged Off	30000.0	Source Verified	< 1 year	RENT	
2	Fully Paid	12252.0	Not Verified	10+ years	RENT	
3	Fully Paid	49200.0	Source Verified	10+ years	RENT	
4	Fully Paid	80000.0	Source Verified	1 year	RENT	
5	Fully Paid	36000.0	Source Verified	3 years	RENT	

6	Fully Paid	47004.0	Not Verified	8 years	RENT
7	Fully Paid	48000.0	Source Verified	9 years	RENT
8	Charged Off	40000.0	Source Verified	4 years	OWN
9	Charged Off	15000.0	Verified	< 1 year	RENT

	int_rate	loan_amnt	purpose	term	grade
0	10.65	5000	credit_card	36 months	B
1	15.27	2500	car	60 months	C
2	15.96	2400	small_business	36 months	C
3	13.49	10000	other	36 months	C
4	12.69	3000	other	60 months	B
5	7.90	5000	wedding	36 months	A
6	15.96	7000	debt_consolidation	60 months	C
7	18.64	3000	car	36 months	E
8	21.28	5600	small_business	60 months	F
9	12.69	5375	other	60 months	B

We have 38706 records of past transactions in this dataset. Each record corresponds to an approved loan. The first column indicates whether the borrower paid off the loan (fully paid) or not (charged off). Descriptions of the other nine columns are as follows:

annual_inc	The self-reported annual income provided by the borrower during registration.
verification_status	Indicates if income was verified by LC, not verified, or if the income source was verified
emp_length	Employment length in years. Possible values are between 0 and 10 where 0 means less than one year and 10 means ten or more years.
home_ownership	The home ownership status provided by the borrower during registration or obtained from the credit report. Our values are: RENT, OWN, MORTGAGE, OTHER
int_rate	Interest Rate on the loan
loan_amnt	The listed amount of the loan applied for by the borrower. If at some point in time, the credit department reduces the loan amount, then it will be reflected in this value.
purpose	A category provided by the borrower for the loan request.
term	The number of payments on the loan. Values are in months and can be either 36 or 60.
grade	LC assigned loan grade

1.2.1 Relationships between `loan_status` and other variables (25 mts)


We first perform EDA to examine the pairwise relationship between `loan_status` and each of other variables. We have both discrete and continuous variables in the dataset. Let's consider the continuous variables first.

1.2.2 Exercise 1: (5 mts)

Use the `boxplot()` function from the `seaborn` package to visualize the distributions of `annual_inc`, `int_rate` and `loan_amnt` in users that fully paid their debts vs. those that didn't. Do you see a large difference in the distributions?


Answer. One possible solution is shown below:

```
[6]: sns.boxplot(y="annual_inc", x="loan_status", data = df).set_yscale('log');
```



We can see a slight increase in the mean of annual income when comparing fully paid users to charged off users.

```
[7]: sns.boxplot(y="int_rate", x ='loan_status', data = df);
```




→ interves

The difference between charged-off users and full-paid users in the distribution of `int_rate` is more drastic. The boxplot for fully-paid users is lower than the one for charged-off users. This result is as expected as higher interest rates would increase the actual amount of money a user have to pay.

```
[8]: sns.boxplot(y="loan_amnt", x ='loan_status', data = df)
```

```
[8]: <matplotlib.axes._subplots.AxesSubplot at 0x7fb99903df60>
```



We don't see a clear change in the distribution of `loan_amnt`. Users who fully paid their debts seem to have slightly lower amount of loans on average.

We learned in previous hypothesis testing cases that we can use contingency tables to examine the relationship between two discrete variables. A chi-square test can be performed based on the contingency table to verify if the observed relationship is statistically significant. Let's examine all discrete variables in our dataset.

1.2.3 Exercise 2: (10 mts)

Write a function to generate the contingency table between `loan_status` and another variable called `var`. Make sure to include a column for the proportion of fully-paid users in each level of `var`. Use this function to generate the contingency tables to examine the relationships between `loan_status` and all discrete variables in the dataset.

```
[10]: def get_ct(df, var):
    ct_res = pd.crosstab(df[var], df['loan_status'], margins=True)
    ct_res["Fully Paid(%)"] = round(ct_res["Fully Paid"]/ct_res["All"]*100, 2)
    return ct_res.drop(columns="All")
```

```
[11]: get_ct(df, 'verification_status')
```

	Charged Off	Fully Paid	Fully Paid(%)
Not Verified	2050	14449	87.58

Source Verified	1413	8406	85.61
Verified	1977	10410	84.04
All	5440	33265	85.94

[12]: `get_ct(df, 'emp_length')`

loan_status	Charged Off	Fully Paid	Fully Paid(%)
emp_length			
< 1 year	641	3947	86.03
1 year	460	2787	85.83
2 years	571	3823	87.01
3 years	558	3540	86.38
4 years	466	2978	86.47
5 years	459	2826	86.03
6 years	309	1922	86.15
7 years	263	1512	85.18
8 years	206	1279	86.13
9 years	159	1100	87.37
10+ years	1348	7551	84.85
All	5440	33265	85.94

[13]: `get_ct(df, 'home_ownership')`

loan_status	Charged Off	Fully Paid	Fully Paid(%)
home_ownership			
RENT	2752	15755	85.13
MORTGAGE	2268	14987	86.86
OWN	402	2443	85.87
OTHER	18	80	81.63
All	5440	33265	85.94

[14]: `get_ct(df, 'purpose')`

loan_status	Charged Off	Fully Paid	Fully Paid(%)
purpose			
car	155	1344	89.66
credit_card	516	4491	89.69
debt_consolidation	2703	15572	85.21
educational	52	265	83.60
home_improvement	333	2554	88.47
house	59	310	84.01
major_purchase	211	1905	90.03
medical	101	569	84.93
moving	84	475	84.97
other	600	3238	84.37
renewable_energy	18	77	81.05
small_business	466	1320	73.91

vacation	50	302	85.80
wedding	92	843	90.16
All	5440	33265	85.94

```
[15]: get_ct( df, 'term' )
```

	loan_status	Charged Off	Fully Paid	Fully Paid(%)
term				
36 months	3073	25211	89.14	
60 months	2367	8054	77.29	
All	5440	33265	85.94	

```
[16]: get_ct( df, 'grade' )
```

	loan_status	Charged Off	Fully Paid	Fully Paid(%)
grade				
A	558	9136	94.24	
B	1362	10335	88.36	
C	1296	6642	83.67	
D	1097	4116	78.96	
E	710	2101	74.74	
F	317	718	69.37	
G	100	217	68.45	
All	5440	33265	85.94	

1.2.4 Exercise 3: (5 mts)

Based on the contingency tables above, which variables seem to have associations with `loan_status`? Use chi-square test to verify if the observed associations are statistically significant.

Answer. We compute the chi-square test on each contingency table from Exercise 2:

```
[17]: chi2, p, dof, ex = chi2_contingency(pd.crosstab(df['verification_status'],  
                                              df['loan_status']))  
print('verification_status: p-value of chisquare test =', p)
```

verification_status: p-value of chisquare test = 6.884998281535999e-17

```
[18]: chi2, p, dof, ex = chi2_contingency(pd.crosstab(df['emp_length'],  
                                              df['loan_status']))  
print('emp_length: p-value of chisquare test =', p)
```

emp_length: p-value of chisquare test = 0.06419956286286349

No es estadísticamente significativo

```
[19]: chi2, p, dof, ex = chi2_contingency(pd.crosstab(df['home_ownership'],  
                                              df['loan_status']))
```

```

print('home_ownership: p-value of chisquare test =', p)

home_ownership: p-value of chisquare test = 3.094754275730736e-05

[20]: chi2, p, dof, ex = chi2_contingency(pd.crosstab(df['purpose'], df['loan_status']))
      print('purpose: p-value of chisquare test =', p)

purpose: p-value of chisquare test = 6.53325963877868e-71

[21]: chi2, p, dof, ex = chi2_contingency(pd.crosstab(df['term'], df['loan_status']))
      print('term: p-value of chisquare test =', p)

term: p-value of chisquare test = 2.7851291583847273e-194

[22]: chi2, p, dof, ex = chi2_contingency(pd.crosstab(df['grade'], df['loan_status']))
      print('grade: p-value of chisquare test =', p)

grade: p-value of chisquare test = 0.0

```

It seems that all variables considered above are associated with `loan_status` since the proportions of fully-paid users are not the same across levels of these variables. Notably, the variable that we care about the most, `verification_status`, seems to have an counter-intuitive association with `loan_status`: the less reliable the income information is, the more likely a user would fully repay the debt, as reflected by the decreasing trend of fully paid proportions when `verification_status` increases from not verified to fully verified. The chi-square tests verify that all these observed associations, except for `emp_length`, are indeed statistically significant.

1.2.5 Exercise 4: (5 mts)

Based on the results above, which of the following statements are correct? Select all that apply.

1. There is a decrease of 2.5 on average in interest rate in users who fully paid their debts compared to users who didn't.
2. There is no statistically significant association between `emp_length` and `loan_status` since the proportion of fully paid users remains relatively the same across different levels of `emp_length`.
3. Since the chi-square test for the contingency table of `verification_status` vs. `loan_status` is significant, the decreasing trend of probability of pay-off is not likely due to randomness.
4. Since there is a significant difference between the probability of full payment in users with no verified income and users with verified income based on the chi-square test, we should always enforce income verification.

Answer. 1 is correct as indicated by the boxplots. 2 is not correct since the conclusion should be drawn based on the chi-square test (which is not significant). 3 is not correct. The chi-square test is not a trend test, it only suggests that there is significant difference in the probability of pay-off among the three levels of `verification_status`, but it cannot further verify that the differences form a specific trend. 4 is not correct since it implies a causal relationship between `verification_status` and `loan_status`. There is not enough evidence to deduce this, and we

should first assess for potential confounding variables like we did in previous cases before making such a strong conclusion.

1.2.6 Stratified contingency tables of verification_status vs. loan_status (15 mts)

As we saw in the previous exercise, we need to adjust for other variables in order to remove potential confounding impacts. If we want to know if `verification_status` really is associated with `loan_status` as indicated by the contingency table above, we should consider stratifying the contingency table between `loan_status` and `verification_status` by other variables that are also associated with `loan_status`.

1.2.7 Exercise 5: (10 mts)

Write a function to generate a contingency table for `verification_status` vs. `loan_status` stratified by `stra_var`. (Hint: If `index` includes two variables, the contingency table will be stratified by the first variable). Make sure to include a column for the proportion of fully-paid users in each level of `verification_status`. Use this function to examine the relationship between `verification_status` and `loan_status` when adjusting for `home_ownership` and `term`. Do we still see the same trend across different levels of `home_ownership` and `term`?

Answer. One possible solution is shown below:

```
[23]: def get_ct_stra(stra_var):
    ct_stra = pd.crosstab(index=[stra_var,df.verification_status], columns = df.
    ↪loan_status, margins = True)
    ct_stra["Fully Paid (%)"] = round(ct_stra["Fully Paid"]/
    ↪ct_stra["All"]*100,2)
    return ct_stra.drop(columns="All").drop("All", level=0)
```

```
[24]: get_ct_stra(df.home_ownership)
```

		Charged Off	Fully Paid	Fully Paid (%)
home_ownership	verification_status			
RENT	Not Verified	1072	7133	86.93
	Source Verified	815	4332	84.17
	Verified	865	4290	83.22
MORTGAGE	Not Verified	798	6064	88.37
	Source Verified	487	3443	87.61
	Verified	983	5480	84.79
OWN	Not Verified	170	1210	87.68
	Source Verified	110	622	84.97
	Verified	122	611	83.36
OTHER	Not Verified	10	42	80.77
	Source Verified	1	9	90.00
	Verified	7	29	80.56

```
[25]: get_ct_stra(df.term)
```

```
[25]: loan_status          Charged Off  Fully Paid  Fully Paid (%)  
      term    verification_status  
      36 months Not Verified      1580      12574      88.84  
                  Source Verified      688       6132      89.91  
                  Verified            805       6505      88.99  
      60 months Not Verified      470        1875      79.96  
                  Source Verified      725       2274      75.83  
                  Verified            1172      3905      76.92
```

We can see the trend is preserved when stratified by `home_ownership`. The proportion of fully paid users decreases as the verification status gets less reliable regardless of levels of `home_ownership`. The exception is level `OTHER`. But since we have very few users there, we should not put too much confidence on the results. The trend is no longer preserved in both levels of `term`. We can see the proportion of fully paid users with shorter loan term (36 months) remains the same across different levels of `verification status` at around 88%.

1.2.8 Exercise 6: (5 mts)

Comparing the stratified contingency tables to the unstratified counterparts, which of the following conclusions is correct? Select all that apply.

1. Based on the stratified contingency table, the difference in the probability of full payment between income-verified and not income-verified users might be attributable to the effect of loan term on the probability of full payment.
2. Based on the stratified contingency table, the difference in the probability of full payment between income-verified and not income-verified users is not likely attributable to the effect of home ownership on the probability of full payment.
3. Assume users with no income verified are more likely to borrow short term loans compared to users with verified income and users with short term loans are more likely to pay back their debts. If in fact income verification status has no effect on the probability of paying back one's debt, users with no verified income are on the margin still more likely to pay off their debt.

Answer. 1 is correct. As we see in the stratified table by `term`, the proportion of fully paid loans remains the same across different levels of verification status when `term` is fixed. On the other hand, the stratum-specific probability of full payment is quite different from stratum to stratum. 2 is correct since stratifying the contingency table by `home_ownership` does not significantly change the proportions of fully-paid loans. 3 is correct by definition, and in fact might be the reason why we observed the counter-intuitive trend. We can see the users with no verified income have many more short-term loans, while the stratified contingency table told us that users with longer loan terms are less likely to pay off their debts compared to short-term loan borrowers.

```
[41]: pd.crosstab(df['term'], df['verification_status'])
```

```
[41]: verification_status  Not Verified  Source Verified  Verified  
      term
```

36 months	14154	6820	7310
60 months	2345	2999	5077

1.3 Logistic regression: extending the linear model to binary outcomes (20 mts)

From the chi-square test, we know that verification status is marginally associated with the probability of pay-off. We observed from the contingency table that this probability increases when the information about income becomes less reliable. However, this observation cannot be verified with just the standard chi-square test.

We introduce the logistic regression model to examine if this trend is significant. The logistic regression model is a modified version of the standard linear regression model that deals with classification problems, where the outcomes are binary variables (in this case, whether the loan was paid off or charged off). To understand a logistic regression model, let's first introduce the idea of odds. **The odds of an event is the ratio of its success probability to its failure probability. In our case, the odds of fully paying the debt is defined as**

$$\text{Odds(Full payment)} = \frac{p}{1-p}.$$

where p is the probability of pay-off.

It turns out that it is easier to work with odds than probabilities. To investigate our question about trend in probability of pay-off p , we assume the odds is a function of income verification status:

$$\log(\text{Odds}) = \beta_0 + \beta_1 * \text{vs}. \quad (1)$$

In order to reflect that there is a trend over different verification statuses, we code different levels of `verification_status` using integers: “Not Verified” as 0, “Source Verified” as 1 and “Verified” as 2; i.e. we are renaming categories of a categorical variable to reflect a clear numerical trend. We assert that this coding scheme implies that if β_1 is positive, then there is an increasing trend of pay-off probability as income verification status becomes more reliable. On the other hand, if β_1 is negative, the opposite is true.

Once we fit the model in (1), we can perform a test to examine if β_1 is significantly different from zero. Note if $\beta_1 = 0$, then p is the same for different levels of `verification_status` and thus no trend is present. If β_1 is significantly different from zero, we can use the estimated sign of β_1 to determine the direction of the trend.

1.3.1 Numerical interpretation of logistic regression (5 mts)

Although the logistic regression model comes with more complexity than standard linear regression, the sign of the regression coefficient still represents the direction of influence of a specific variable: a positive coefficient means the probability of pay-off will increase if the associated variable increases and vice versa.

The other important part of information when interpreting a logistic regression model is the p -value for each regression coefficient. The p-value indicates the result of the following hypothesis test:

$$H_0 : \beta = 0 \text{ vs. } H_A : \beta \neq 0.$$

The test determines if the difference in the probability of pay-off is associated with changes in the corresponding variable.

Sometimes you might find that both the p -value and the estimated effect size (regression coefficient) for a variable is small. You should be cautious about this situation as p -values are not the only way to determine if variables are important.

$$\log\left(\frac{p}{1-p}\right) = \beta_0 + \beta_1 \cdot \text{var}$$

1.3.2 Geometric interpretation of logistic regression (5 mts)


The geometric perspective of logistic regression models is also crucial, as a logistic regression model naturally defines a “decision boundary”. If we specify that p is lower than a given decision boundary cut-off, where p_0 (e.g. 0.5) indicates defaulters, then (1) indicates that samples with $\beta_0 + \beta_1 * \text{verification_status} > \log(p_0)/(1 - \log(p_0))$ are non-defaulters and $\beta_0 + \beta_1 * \text{verification_status} < \log(p_0)/(1 - \log(p_0))$ are defaulters. The line which separates defaulters and non-defaulters, which is defined by the equation $\beta_0 + \beta_1 * \text{verification_status} = \log(p_0)/(1 - \log(p_0))$ is the decision boundary determined by model (1) and cut-off value p_0 .

Assume we know p is determined by the following logistic model:

$$\log(p/(1 - p)) = 2 * \text{loan_amnt} - \text{annual_inc}.$$

Here `annual_inc` is divided by 10^4 and `loan_amnt` is divided by 1000. The line $2 * \text{loan_amnt} = \text{annual_inc}$ in this case separates defaulters from non-defaulters and is the decision boundary with cut-off probability 0.5:

```
[44]: x = np.arange(0,100,1)
y = x/2
plt.plot(x, y, label = "Decision boundary");
plt.xlabel("annual_inc/1e4");
plt.ylabel("loan_amnt/1e3");
plt.text(10, 30, "Non-defaulter", fontsize = 18)
plt.text(60, 20, "Defaulter", fontsize = 18)
plt.legend();
```



The shape of the decision boundary in logistic regression is not always a line; oftentimes, it is a curve.

1.3.3 Visualizing the decision boundary (5 mts)

We consider the following logistic model for p :

$$\log(p/(1-p)) = \beta_0 + \beta_1 * \text{loan_amnt} + \beta_2 * \text{annual_inc} + \beta_3 * \text{loan_amnt}^2.$$

We know the corresponding decision boundary with cut-off value $p = 0.5$ is

$$\beta_0 + \beta_1 * \text{loan_amnt} + \beta_2 * \text{annual_inc} + \beta_3 * \text{loan_amnt}^2 = 0.$$

Use the following interactive plot to tweak the parameters and see how the decision boundary changes:

```
[45]: %matplotlib widget
import ipywidgets as widgets
from ipywidgets import interact

idx_plot = np.random.choice(df.shape[0], size=500, replace = False)
plt_data = df.iloc[idx_plot]

fig, ax = plt.subplots()
```

```

# plt.subplots_adjust(left=0.25, bottom=0.35)
plt.scatter(plt_data[plt_data.loan_status=="Fully Paid"].annual_inc/10000, \
            plt_data[plt_data.loan_status=="Fully Paid"].loan_amnt/1000, \
            c = 'blue', label = 'Fully paid')
plt.scatter(plt_data[plt_data.loan_status=="Charged Off"].annual_inc/10000, \
            plt_data[plt_data.loan_status=="Charged Off"].loan_amnt/1000, \
            c = 'red', label = 'Charged off')
plt.xlim(0,40);
plt.legend();


#plot boundary
beta0 = 1
beta1 = -1
beta2 = 1
beta3 = -1

x_plt = np.arange(0,40,1)
y_plt = (-beta0 - beta1*x_plt - beta3*x_plt**2)/(beta2)
plt.xlabel("annual_inc/1e4")
plt.ylabel("loan_amnt/1e3")
l, = plt.plot(x_plt, y_plt, lw = 2, color="k", label = "Current Decision\u2192boundary")
plt.ylim(0,60);
plt.legend();

def update(beta0=1, beta1=-1, beta2=1, beta3=-1):
    y_plt = (-beta0 - beta1*x_plt - beta3*x_plt**2)/(beta2)
    l.set_ydata(y_plt);
    fig.canvas.draw_idle();

interact(update, beta0 = widgets.FloatSlider(value=beta0,min=-10,max=10,step=0.5),\
         beta1 = widgets.FloatSlider(value=beta1,min=-10,max=10,step=0.5),
         beta2 = widgets.FloatSlider(value=beta2,min=0.5,max=10,step=0.5),
         beta3 = widgets.FloatSlider(value=beta3,min=-10,max=0,step=0.5));

Canvas(toolbar=Toolbar(toolitems=[('Home', 'Reset original view', 'home', 'home'), ('Back', 'B
interactive(children=(FloatSlider(value=1.0, description='beta0', max=10.0, min=-10.0, step=0.5),
```



1.4 Fitting the logistic regression model (25 mts)

We now fit the logistic model using the `statmodels` package with the specification we described above:

```
[50]: # code the discrete variable by the specification above
df_log1 = pd.DataFrame(columns=['verification_status','loan_status'])
df_log1['verification_status'] = df.verification_status.cat.codes
df_log1['loan_status'] = df.loan_status.cat.codes
df_log1['Intercept'] = 1

logit = sm.Logit(df_log1['loan_status'], df_log1[['Intercept','verification_status']])
logit_res = logit.fit()
print( logit_res.summary() )
```

Optimization terminated successfully.

Current function value: 0.405007

Iterations 6

Logit Regression Results

Dep. Variable:	loan_status	No. Observations:	38705
Model:	Logit	Df Residuals:	38703
Method:	MLE	Df Model:	1

```

Date: Fri, 29 May 2020 Pseudo R-squ.: 0.002350
Time: 19:31:55 Log-Likelihood: -15676.
converged: True LL-Null: -15713.
Covariance Type: nonrobust LLR p-value: 8.419e-18
=====
=====

            coef      std err          z      P>|z|      [0.025
0.975]
-----
Intercept      1.9469      0.022      88.290      0.000      1.904
1.990
verification_status -0.1460      0.017     -8.600      0.000     -0.179
-0.113
=====
=====
```

This indicates that when `verification_status` increases by 1, the odds of pay-off decreases by 14%. Since the p -value for the coefficient is smaller than 0.05, we can conclude that the decreasing trend we observed from the table is not likely due to randomness.

An important output above is `Pseudo R-squ.`. This metric is similar to R-squared for linear models. If this number is large, then the variables in the model explain a large portion of the drivers of people's tendency to pay off their debt. In our case, the metric is only 0.2%, which means that we have a lot of room to improve.

1.4.1 Evaluating the numerical coding scheme for `verification_status` (5 mts)

Note that our coding scheme assumes that the difference in the log-odds $p/(1-p)$ between people from any two consecutive levels of `verification_status` is the same. This assumption might be too strong and we ought to check if it is reasonable. We will fit the same model with a different coding scheme for `verification_status`. Instead of mapping levels to integers, we use dummy variables for different levels. With this specification, the difference in log-odds between two consecutive levels is not necessarily a constant. We set "Not Verified" as the reference level. The code and the results are shown below:

```
[52]: vs_dummy = pd.get_dummies(df.verification_status, drop_first=True)
vs_dummy.head()
```

```
[52]:   Source Verified Verified
0           0         1
1           1         0
2           0         0
3           1         0
4           1         0
```

```
[53]: #vs_dummy = pd.get_dummies(df.verification_status, drop_first=True )
df_log1 = pd.concat([df_log1,vs_dummy], axis = 1)
```

```

logit_dummy = sm.Logit(df_log1['loan_status'], df_log1[['Intercept', 'Source_U
→Verified', 'Verified']])
logit_dummy_res = logit_dummy.fit()
print( logit_dummy_res.summary() )

```

```

Optimization terminated successfully.
    Current function value: 0.405001
    Iterations 6
                    Logit Regression Results
=====
Dep. Variable:          loan_status    No. Observations:                 38705
Model:                  Logit         Df Residuals:                  38702
Method:                 MLE          Df Model:                      2
Date:      Fri, 29 May 2020   Pseudo R-squ.:            0.002366
Time:           19:32:57     Log-Likelihood:        -15676.
converged:            True       LL-Null:             -15713.
Covariance Type:    nonrobust    LLR p-value:      7.140e-17
=====
===
                   coef      std err      z      P>|z|      [0.025
0.975]
-----
Intercept          1.9528      0.024    82.741      0.000      1.907
1.999
Source Verified   -0.1696      0.037    -4.558      0.000     -0.242
-0.097
Verified           -0.2916      0.034    -8.566      0.000     -0.358
-0.225
=====
===

```

Based on the results, we can see the difference in log-odds between people from level “Source Verified” and level “Not Verified” is -0.16 while the difference between “Verified” and “Not Verified” is -0.28. The results confirm that the probability of pay-off is indeed decreasing when information of income becomes more reliable. But this also means that the differences in log-odds are not constant and the previous coding scheme might not be reasonable. We can use the Pseudo R-squ. as a metric to examine if one coding scheme is better than the other.

1.4.2 Evaluating the logistic regression model (15 mts)

A major drawback of the pseudo R^2 metric is that it cannot be directly translated into the prediction accuracy of a model. Remember that the goal of our whole analysis is to predict whether or not a user is going to pay off his debt. But how do we evaluate if the predication is accurate or not? This question is not as simple as it seems and the reason is because logistic models only produce probabilities while what we observed are binary labels.

A straightforward solution is that you can select a cut-off value for the predicted probability such that people with probabilities smaller than the cut-off value are deemed as defaulters. This converts the probabilities into binary labels, after which we can see if the predicted labels are the same as the observed labels to measure prediction accuracy.

However, this only takes into account a single cut-off value. We ought to summarize the prediction accuracy across all different cut-off values to produce a unified evalution of our model. For logistic models, this is usually implemented via the **receiver operating characteristic (ROC) curve**.

In order to understand a ROC curve, we need to first know the **true positive rate** (TPR) and **true negative rate** (TNR). In our context, TPR measures the proportion of people who paid off their debt that are correctly identified by our model, whereas TNR measures the proportion of defaulters that are identified as defaulters. Each point in the plot indicates the TPR and 1-TNR at a specific cut-off value:

There is a constant trade-off between TPR and TNR. It is usually hard to find a rule with high TPR and TNR at the same time. People will choose to prioritize one of them based on the particular scenario. For instance, in our case, TNR is usually more important than TPR since we would rather avoid lending to defaulters than miss lending to one good candidate.

The ROC curve plots TPR against 1-TNR and illustrates the behavior of a logistic model when the cut-off point changes. We can further summarize the whole ROC curve by the **area under the curve (AUC)**. Roughly speaking, AUC indicates the probability that a randomly selected defaulter will have smaller p than a randomly selected non-defaulter.

If AUC is close to one, the rule is close to a perfect rule. The ROC curve for a random guess rule (set p to 0.5 for all applicants) is the diagonal line in the plot and usually serves as a baseline to check if the logistic model learns any information from the data above and beyond pure randomness. The ROC curves of the above two logistic regression models are shown below:

```
[54]: predict_p = logit_res.predict(df_log1[['Intercept', 'verification_status']])
predict_dummy_p = logit_dummy_res.predict(df_log1[['Intercept', 'Source_→Verified', 'Verified']])
roc_p = roc_curve(df_log1["loan_status"], predict_p)
roc_dummy_p = roc_curve(df_log1["loan_status"], predict_dummy_p)
auc_p = auc(roc_p[0], roc_p[1])
auc_dummy_p = auc(roc_dummy_p[0], roc_dummy_p[1])


plt.figure()
lw = 2
plt.plot(roc_p[0], roc_p[1], color='darkorange',
         lw=lw, label='ROC curve for coding scheme 1 (AUC = %0.2f)' % auc_p)
plt.plot(roc_dummy_p[0], roc_dummy_p[1], color='darkgreen',
         lw=lw, label='ROC curve for coding scheme 2 (AUC = %0.2f)' % auc_dummy_p)
plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--', label='Random→guess')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
```

```

plt.ylabel('True Positive Rate')
plt.title('Classification of Defaulters')
plt.legend(loc="lower right")
plt.show()

```

Canvas(toolbar=Toolbar(toolitems=[('Home', 'Reset original view', 'home', 'home'), ('Back', 'Bac



From the figure we find that both ROCs have AUCs quite close to 0.5, confirming that `verification_status` alone is not very predictive of loan status above and beyond random guessing. This also confirms that the first coding scheme has identical performance to the second coding scheme. Since the first coding scheme is more parsimonious (i.e. contains fewer parameters) than the second scheme, we will use the first coding scheme for the rest of this analysis.

1.4.3 Exercise 7: (5 mts)

Based on the results above (pseudo R^2 and AUCs), what can we conclude about the appropriate code scheme for `verification_status`?

Answer. There is no difference in the two metrics when comparing the two different coding schemes. In statistics, we prefer simple models and therefore the first coding scheme is appropriate. Note that in the second coding scheme, we need to estimate one more parameter than in the first coding scheme.

1.5 Adding additional variables to the model (20 mts)

In the previous section, we verified that there is indeed a decreasing trend in pay-off probability associated with more reliable income information. We also performed two model-checking approaches and determined that the integer coding for `verification_status` is adequate. Yet the discovered decreasing trend is counterintuitive. This might be because many other factors are also associated with the probability of pay-off as we have seen in the exploratory analyses. Similar to linear regression, not accounting for these factors will result in completely different answers for the effect of `verification_status`, which is already hinted at in the stratified contingency tables.

We find that except for `annual_inc`, all other variables in the dataset are correlated with `verification_status`. For example, there seems to be a strong correlation between `verification_status` and `loan_amnt`. People with no verified income source tend to borrow less money than people with verified income source. This might be the reason for the decreasing trend associated with `verification_status`, since smaller loans are more likely to be paid off. Let's now incorporate all these variables (including `annual_inc` for now) into the logistic model along with `verification_status` and see if the decreasing trend associated with `verification_status` is still present:

```
[55]: # Preprocessing the variables
df_log2 = pd.concat([(df.loan_amnt - df.loan_amnt.mean())/df.loan_amnt.std(), \
                     (df.int_rate - df.int_rate.mean())/df.int_rate.std(), \
                     (df.annual_inc - df.annual_inc.mean())/df.annual_inc, \
                     std(), \
                     pd.get_dummies(df.home_ownership, prefix="home", \
                     drop_first=True), \
                     pd.get_dummies(df.purpose, prefix="purpose", \
                     drop_first=True), \
                     pd.get_dummies(df.grade, prefix="grade", drop_first=True)], \
                     axis=1)
df_log2["verification_status"] = df.verification_status.cat.codes
df_log2["emp_length"] = df.emp_length.cat.codes
df_log2["term"] = df.term.cat.codes
df_log2["Intercept"] = 1
```

```
[56]: logit_full1 = sm.Logit(df.loan_status.cat.codes, df_log2)
logit_full1_res = logit_full1.fit()
print(logit_full1_res.summary())
```

Optimization terminated successfully.

Current function value: 0.378075

Iterations 7

Logit Regression Results

Dep. Variable:	y	No. Observations:	38705
Model:	Logit	Df Residuals:	38676
Method:	MLE	Df Model:	28
Date:	Fri, 29 May 2020	Pseudo R-squ.:	0.06869

Time: 19:36:42 Log-Likelihood: -14633.
 converged: True LL-Null: -15713.
 Covariance Type: nonrobust LLR p-value: 0.000
 =====

		coef	std err	z	P> z
[0.025	0.975]				
-----	-----	-----	-----	-----	-----
loan_amnt		-0.0075	0.019	-0.389	0.697
-0.045	0.030				
int_rate		-0.4050	0.054	-7.454	0.000
-0.511	-0.299				
annual_inc		0.3615	0.031	11.641	0.000
0.301	0.422				
home_MORTGAGE		0.0382	0.035	1.095	0.273
-0.030	0.107				
home_OWN		0.0209	0.060	0.348	0.728
-0.097	0.138				
home_OTHER		-0.4454	0.268	-1.660	0.097
-0.971	0.080				
purpose_credit_card		0.0364	0.100	0.363	0.717
-0.160	0.233				
purpose_debt_consolidation		-0.2190	0.091	-2.405	0.016
-0.398	-0.041				
purpose_educational		-0.5663	0.178	-3.186	0.001
-0.915	-0.218				
purpose_home_improvement		-0.1624	0.107	-1.523	0.128
-0.371	0.047				
purpose_house		-0.3556	0.171	-2.077	0.038
-0.691	-0.020				
purpose_major_purchase		-0.0024	0.114	-0.021	0.983
-0.227	0.222				
purpose_medical		-0.3945	0.141	-2.794	0.005
-0.671	-0.118				
purpose_moving		-0.4147	0.150	-2.766	0.006
-0.708	-0.121				
purpose_other		-0.4050	0.099	-4.105	0.000
-0.598	-0.212				
purpose_renewable_energy		-0.7242	0.286	-2.530	0.011
-1.285	-0.163				
purpose_small_business		-0.9414	0.105	-8.961	0.000
-1.147	-0.735				
purpose_vacation		-0.4371	0.179	-2.439	0.015
-0.788	-0.086				
purpose_wedding		0.1348	0.142	0.948	0.343
-0.144	0.414				
grade_B		-0.2733	0.074	-3.709	0.000

-0.418	-0.129				
grade_C		-0.3549	0.103	-3.448	0.001
-0.557	-0.153				
grade_D		-0.3775	0.132	-2.860	0.004
-0.636	-0.119				
grade_E		-0.2973	0.160	-1.859	0.063
-0.611	0.016				
grade_F		-0.3238	0.194	-1.669	0.095
-0.704	0.057				
grade_G		-0.1802	0.238	-0.757	0.449
-0.647	0.286				
verification_status		0.0082	0.020	0.411	0.681
-0.031	0.047				
emp_length		-0.0147	0.004	-3.301	0.001
-0.023	-0.006				
term		-0.4520	0.037	-12.081	0.000
-0.525	-0.379				
Intercept		2.6163	0.123	21.216	0.000
2.375	2.858				
<hr/>					
<hr/>					

1.5.1 Exercise 8: (5 mts)

Based on the results of the logistic model with all variables included, which of the following conclusions is correct? Select all that apply.

1. There is no evidence that supports the association between probability of pay-off and verification status after accounting for other variables.
2. Loans borrowed for small business have the smallest pay-off probability.
3. The odds of pay-off in users with LendingClub grade B decreases by 24% compared to users with grade A after accounting for other variables.

Answer. 1 is correct as the p - value for the variable `verification_status` is not smaller than 0.05. 2 is not correct as the statement does not account for the effects of other variables. 3 is correct as we can calculate it using the estimated coefficient.

1.5.2 Why do we observe a marginal trend associated with `verification_status`? (10 mts)

Interestingly, we find that the regression coefficient for `verification_status` is now positive and the p - value for it is very large (0.68). This indicates that after adjusting for all other variables in the dataset, `verification_status` is no longer significantly associated with the probability of pay-off. Let's now find out what confounding variables were responsible for introducing the initial decreasing trend that we observed.

We can investigate this problem using the regression results along with the correlation between `verification_status` and all other variables. First let's take a look at the pairwise relationship

between `verification_status` and all other variables:

```
[57]: def get_rctable(var1, var2):
    res = pd.crosstab(df[var2], df[var1])
    chi2, p, dof, ex = chi2_contingency(res)
    output = round(res.div(res.sum(axis = 1), axis = 0)*100, 2)
    return output.style.set_caption("%s vs. %s: Chi-square p-value=% .2f" %_
    ↪(var1, var2, p))
```

```
[58]: display(get_rctable("emp_length", "verification_status"))
display(get_rctable("home_ownership", "verification_status"))
display(get_rctable("purpose", "verification_status"))
display(get_rctable("term", "verification_status"))
display(get_rctable("grade", "verification_status"))
```

<pandas.io.formats.style.Styler at 0x7fb998cd4da0>

<pandas.io.formats.style.Styler at 0x7fb99e23a5f8>

<pandas.io.formats.style.Styler at 0x7fb9d96f7898>


<pandas.io.formats.style.Styler at 0x7fb99b74fc50>

<pandas.io.formats.style.Styler at 0x7fb9af5e12e8>

```
[59]: sns.boxplot(y="loan_amnt", x='verification_status', data = df)
```


Canvas(toolbar=Toolbar(toolitems=[('Home', 'Reset original view', 'home', 'home'), ('Back', 'B

```
[59]: <matplotlib.axes._subplots.AxesSubplot at 0x7fb99b5620b8>
```



```
[60]: sns.boxplot(y="annual_inc", x='verification_status', data = df).  
      set_yscale('log')
```


```
Canvas(toolbar=Toolbar(toolitems=[('Home', 'Reset original view', 'home', 'home'), ('Back', 'B
```



```
[61]: sns.boxplot(y="int_rate", x='verification_status', data = df)
```

```
Canvas(toolbar=Toolbar(toolitems=[('Home', 'Reset original view', 'home', 'home'), ('Back', 'B
```

```
[61]: <matplotlib.axes._subplots.AxesSubplot at 0x7fb998e0cd68>
```



1.5.3 Exercise 9: (5 mts)

Based on the relationships we observed in the contingency tables and the boxplots above, which of the following statements are correct? Select all that apply.

- A. The marginal trend is that people with verified income tend to borrow loans with higher interest rates
- B. The marginal trend is that people with verified income tend to have longer employment length
- C. The marginal trend is that people with verified income tend to have mortgages

Answer. A and B are both correct since high interest rate and long employment length are both associated with lower probability of pay-off. C is not correct since home ownership is not significantly associated with probability of pay-off after controlling for other variables.

1.6 Do we need to include verification_status? (10 mts)

From the above, we might reach the conclusion that `verification_status` does not matter when we want to predict if a borrower is going to default on his/her debt. But does this mean that in the future, we can just ignore the verification status of an applicant when he/she is filing for a loan application through LC?

1.6.1 Exercise 10: (5 mts)

Which of following statements is true regarding what we should do with `verification_status`, based on the results we have so far?

- A. Since the p - value for `verification_status` is large, we can remove it from our classification model and in the future, we do not need to collect this part of information from applicants.
- B. Although the p - value for `verification_status` is large, the estimated regression coefficient for it is non-zero. We should keep the variable in the model and continue to collect the information in the future.
- C. We should carry out an additional evaluation to compare the prediction accuracies of the model with and without `verification_status` before we can make any further decision.
- D. None of the above.

Answer. A and B are both incorrect. We should not only base our decision based on the p - value alone. C is correct.

Let's evaluate the prediction accuracy using ROC. The results are shown below:

```
[62]: predict_withvs = logit_full1_res.predict(df_log2)
logit_full_novs = sm.Logit(df.loan_status.cat.codes, df_log2.loc[:, df_log2.
    ↪columns != 'verification_status'])
predict_novs = logit_full_novs.fit(disp=0).predict(df_log2.loc[:, df_log2.
    ↪columns != 'verification_status'])

roc_vs = roc_curve(df.loan_status.cat.codes, predict_withvs)
roc_novs = roc_curve(df.loan_status.cat.codes, predict_novs)
auc_vs = auc(roc_vs[0], roc_vs[1])
auc_novs = auc(roc_novs[0], roc_novs[1])


plt.figure()
lw = 2
plt.plot(roc_vs[0], roc_vs[1], color='darkorange',
          lw=lw, label='With verification_status (AUC = %0.2f)' % auc_vs)
plt.plot(roc_novs[0], roc_novs[1], color='darkgreen',
          lw=lw, label='Without verification_status (AUC = %0.2f)' % auc_novs)
plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--', label='Random
    ↪guess')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
```

```

plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Classification of Defaulters')
plt.legend(loc="lower right")
plt.show()

```

Canvas(toolbar=Toolbar(toolitems=[('Home', 'Reset original view', 'home', 'home'), ('Back', 'Bac



1.6.2 Exercise 11: (5 mts)

Based on all the results above, what can we conclude about the verification process?

Answer. We find that the AUC for the model with `verification_status` included is the same as the model without `verification_status` included (0.69). Since we always want our model to be as simple as possible, this means that we should not include `verification_status` as a variable when predicting if a user is likely to pay off his debt.

1.7 Conclusions (5 mts)

In this case, we investigated whether the verification status of applicants' income source is important when predicting if an applicant will pay off his/her debt. We constructed a logistic model to first examine the marginal relationship between verification status and probability of pay-off.

We discovered a counterintuitive trend: less reliable income source information was correlated with higher pay-off probability. However, after fitting a larger logistic model with all independent variables that were available, we found that this trend was introduced via confounding effects of interest rate, annual income, and loan term.

After accounting for these variables, verification status was no longer significantly associated with pay-off probability. Based on p - values and the model ROC, we concluded that the verification process is irrelevant and we could potentially remove it from the required items and simplify the lending process.

1.8 Takeaways (5 mts)

Logistic regression naturally extends the concept of linear regression and is the benchmark model for performing classification tasks. The coefficients have similar interpretation as that of linear regression. ROC curves are then useful for comparing different classification models.

[]: